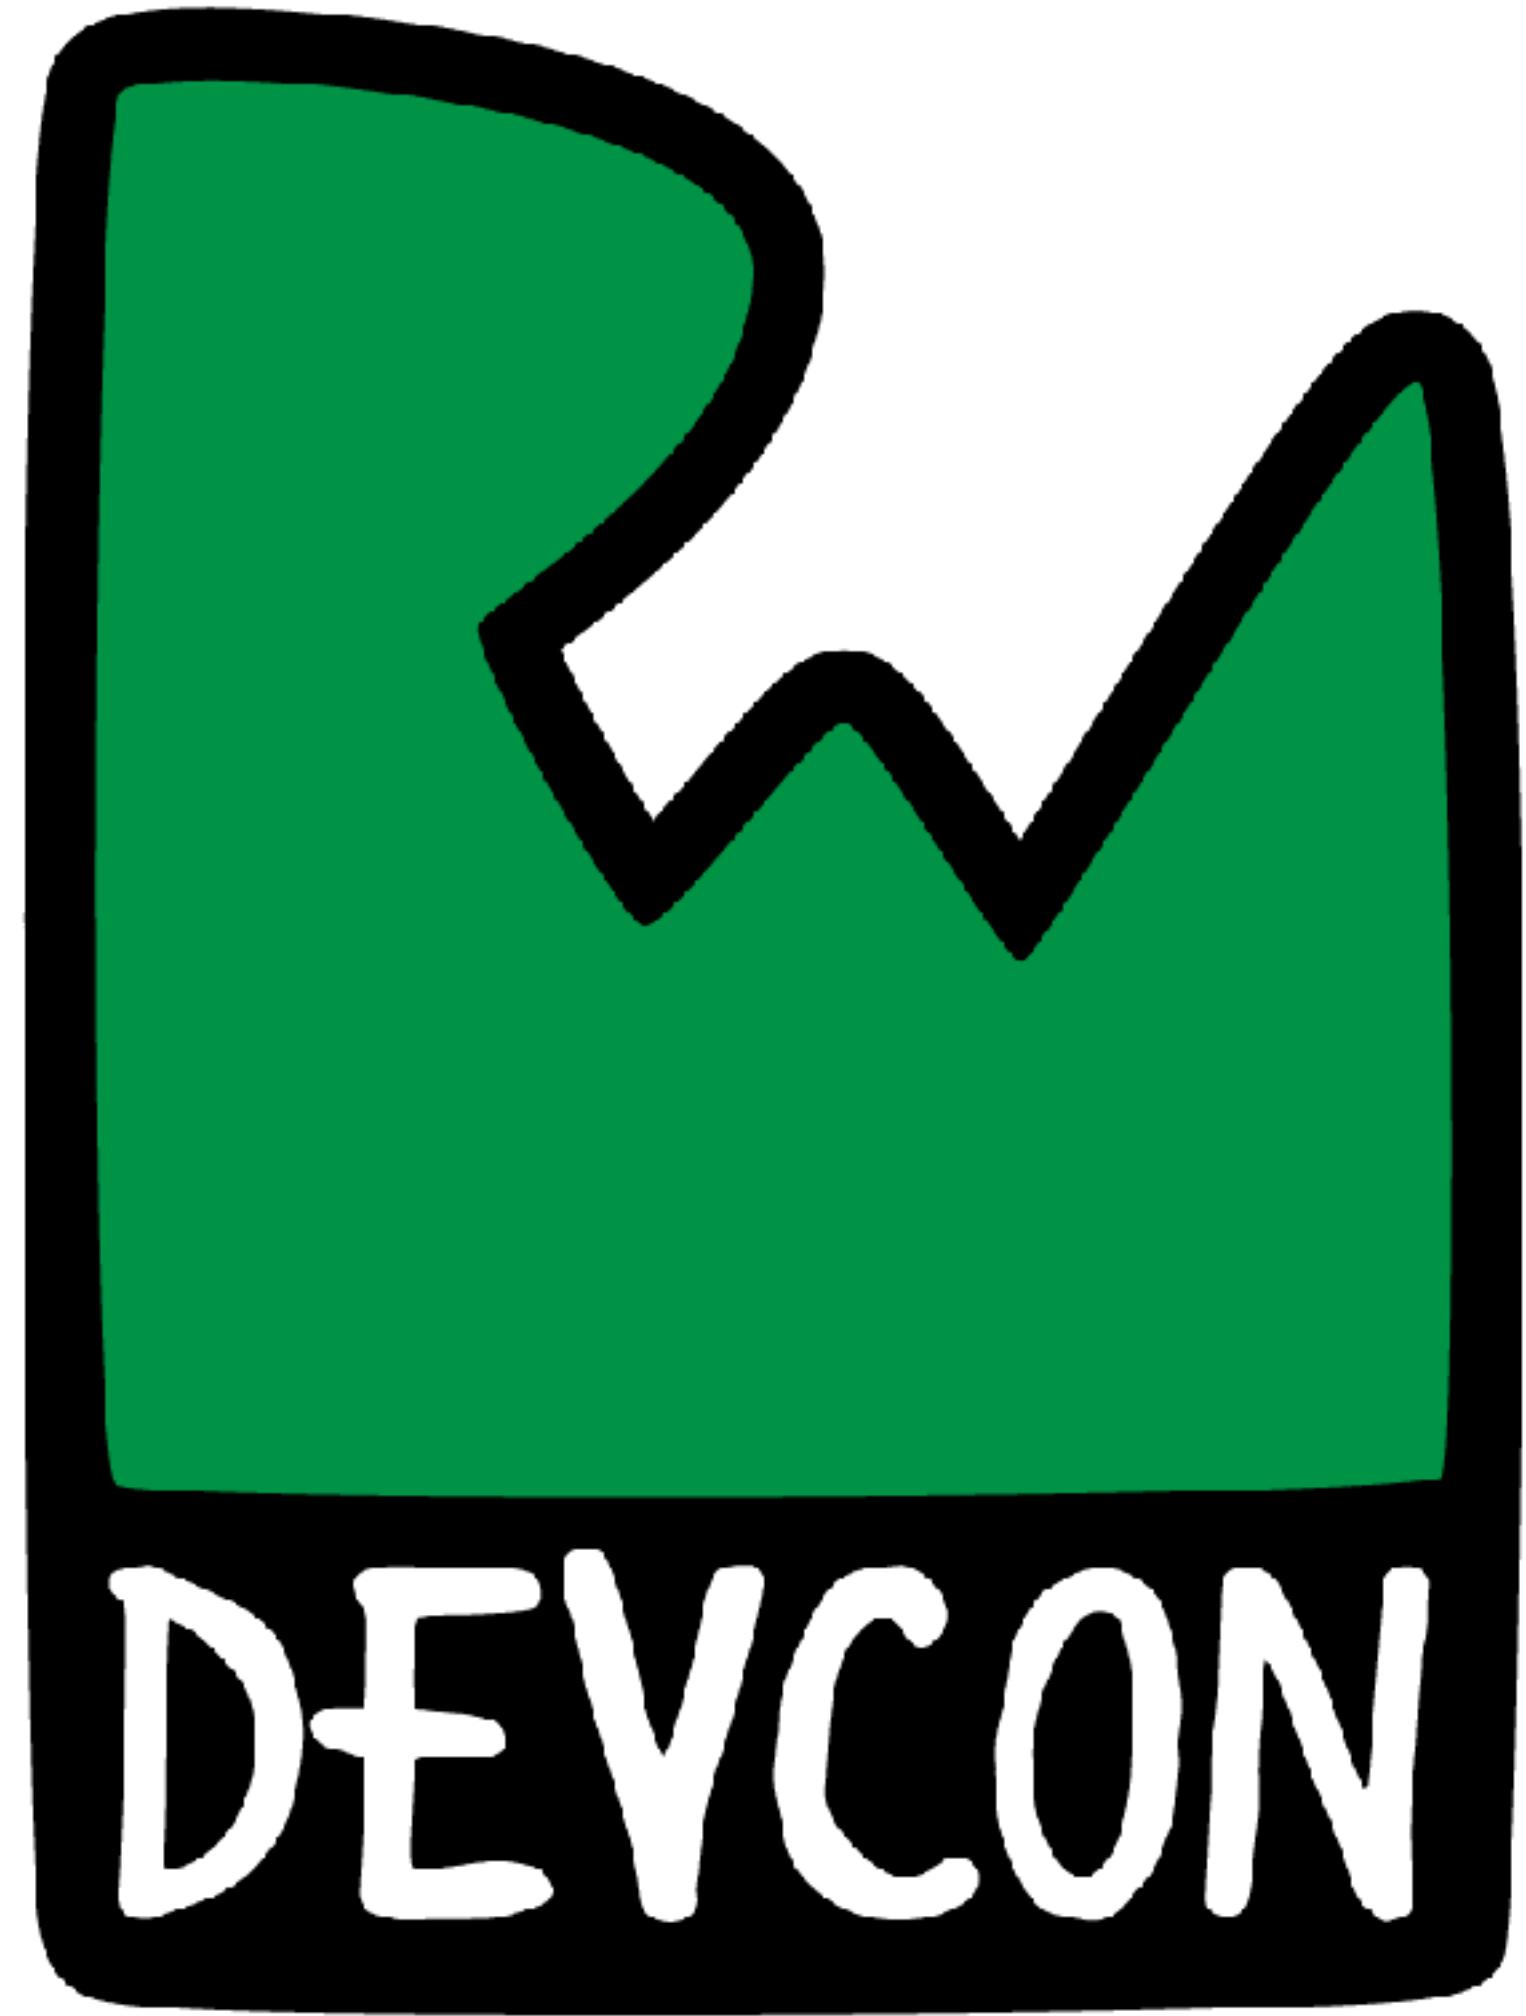


# Workshop: ARKit in Depth

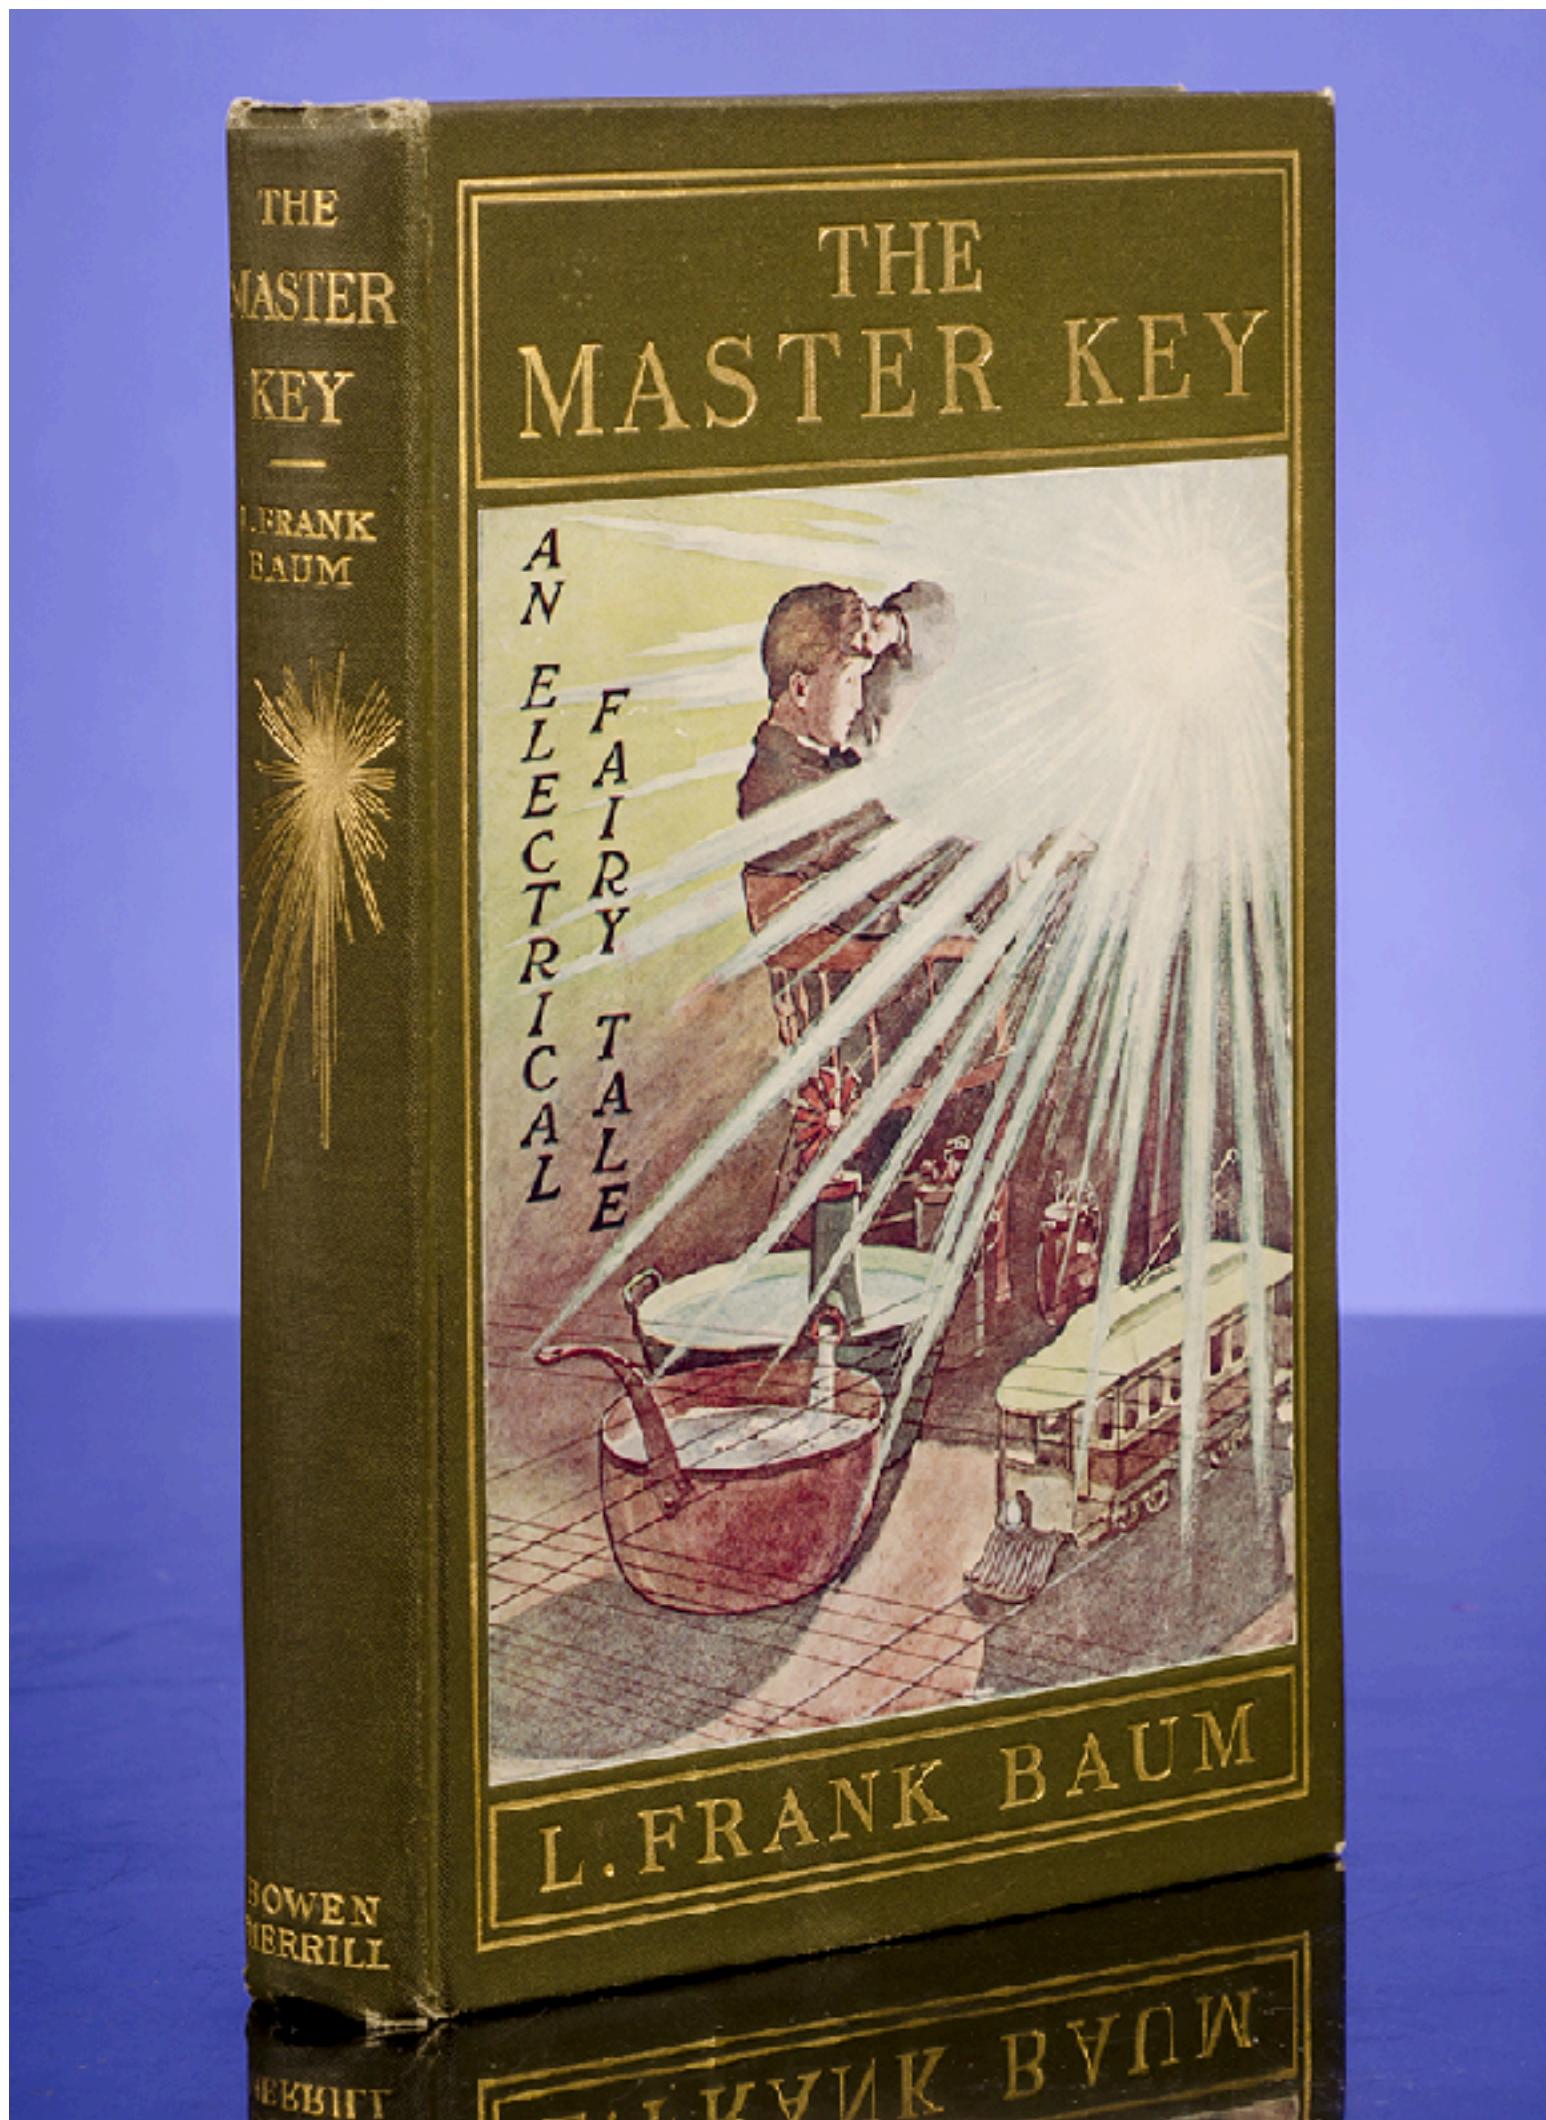


RW  
OVERVIEW



R  
W

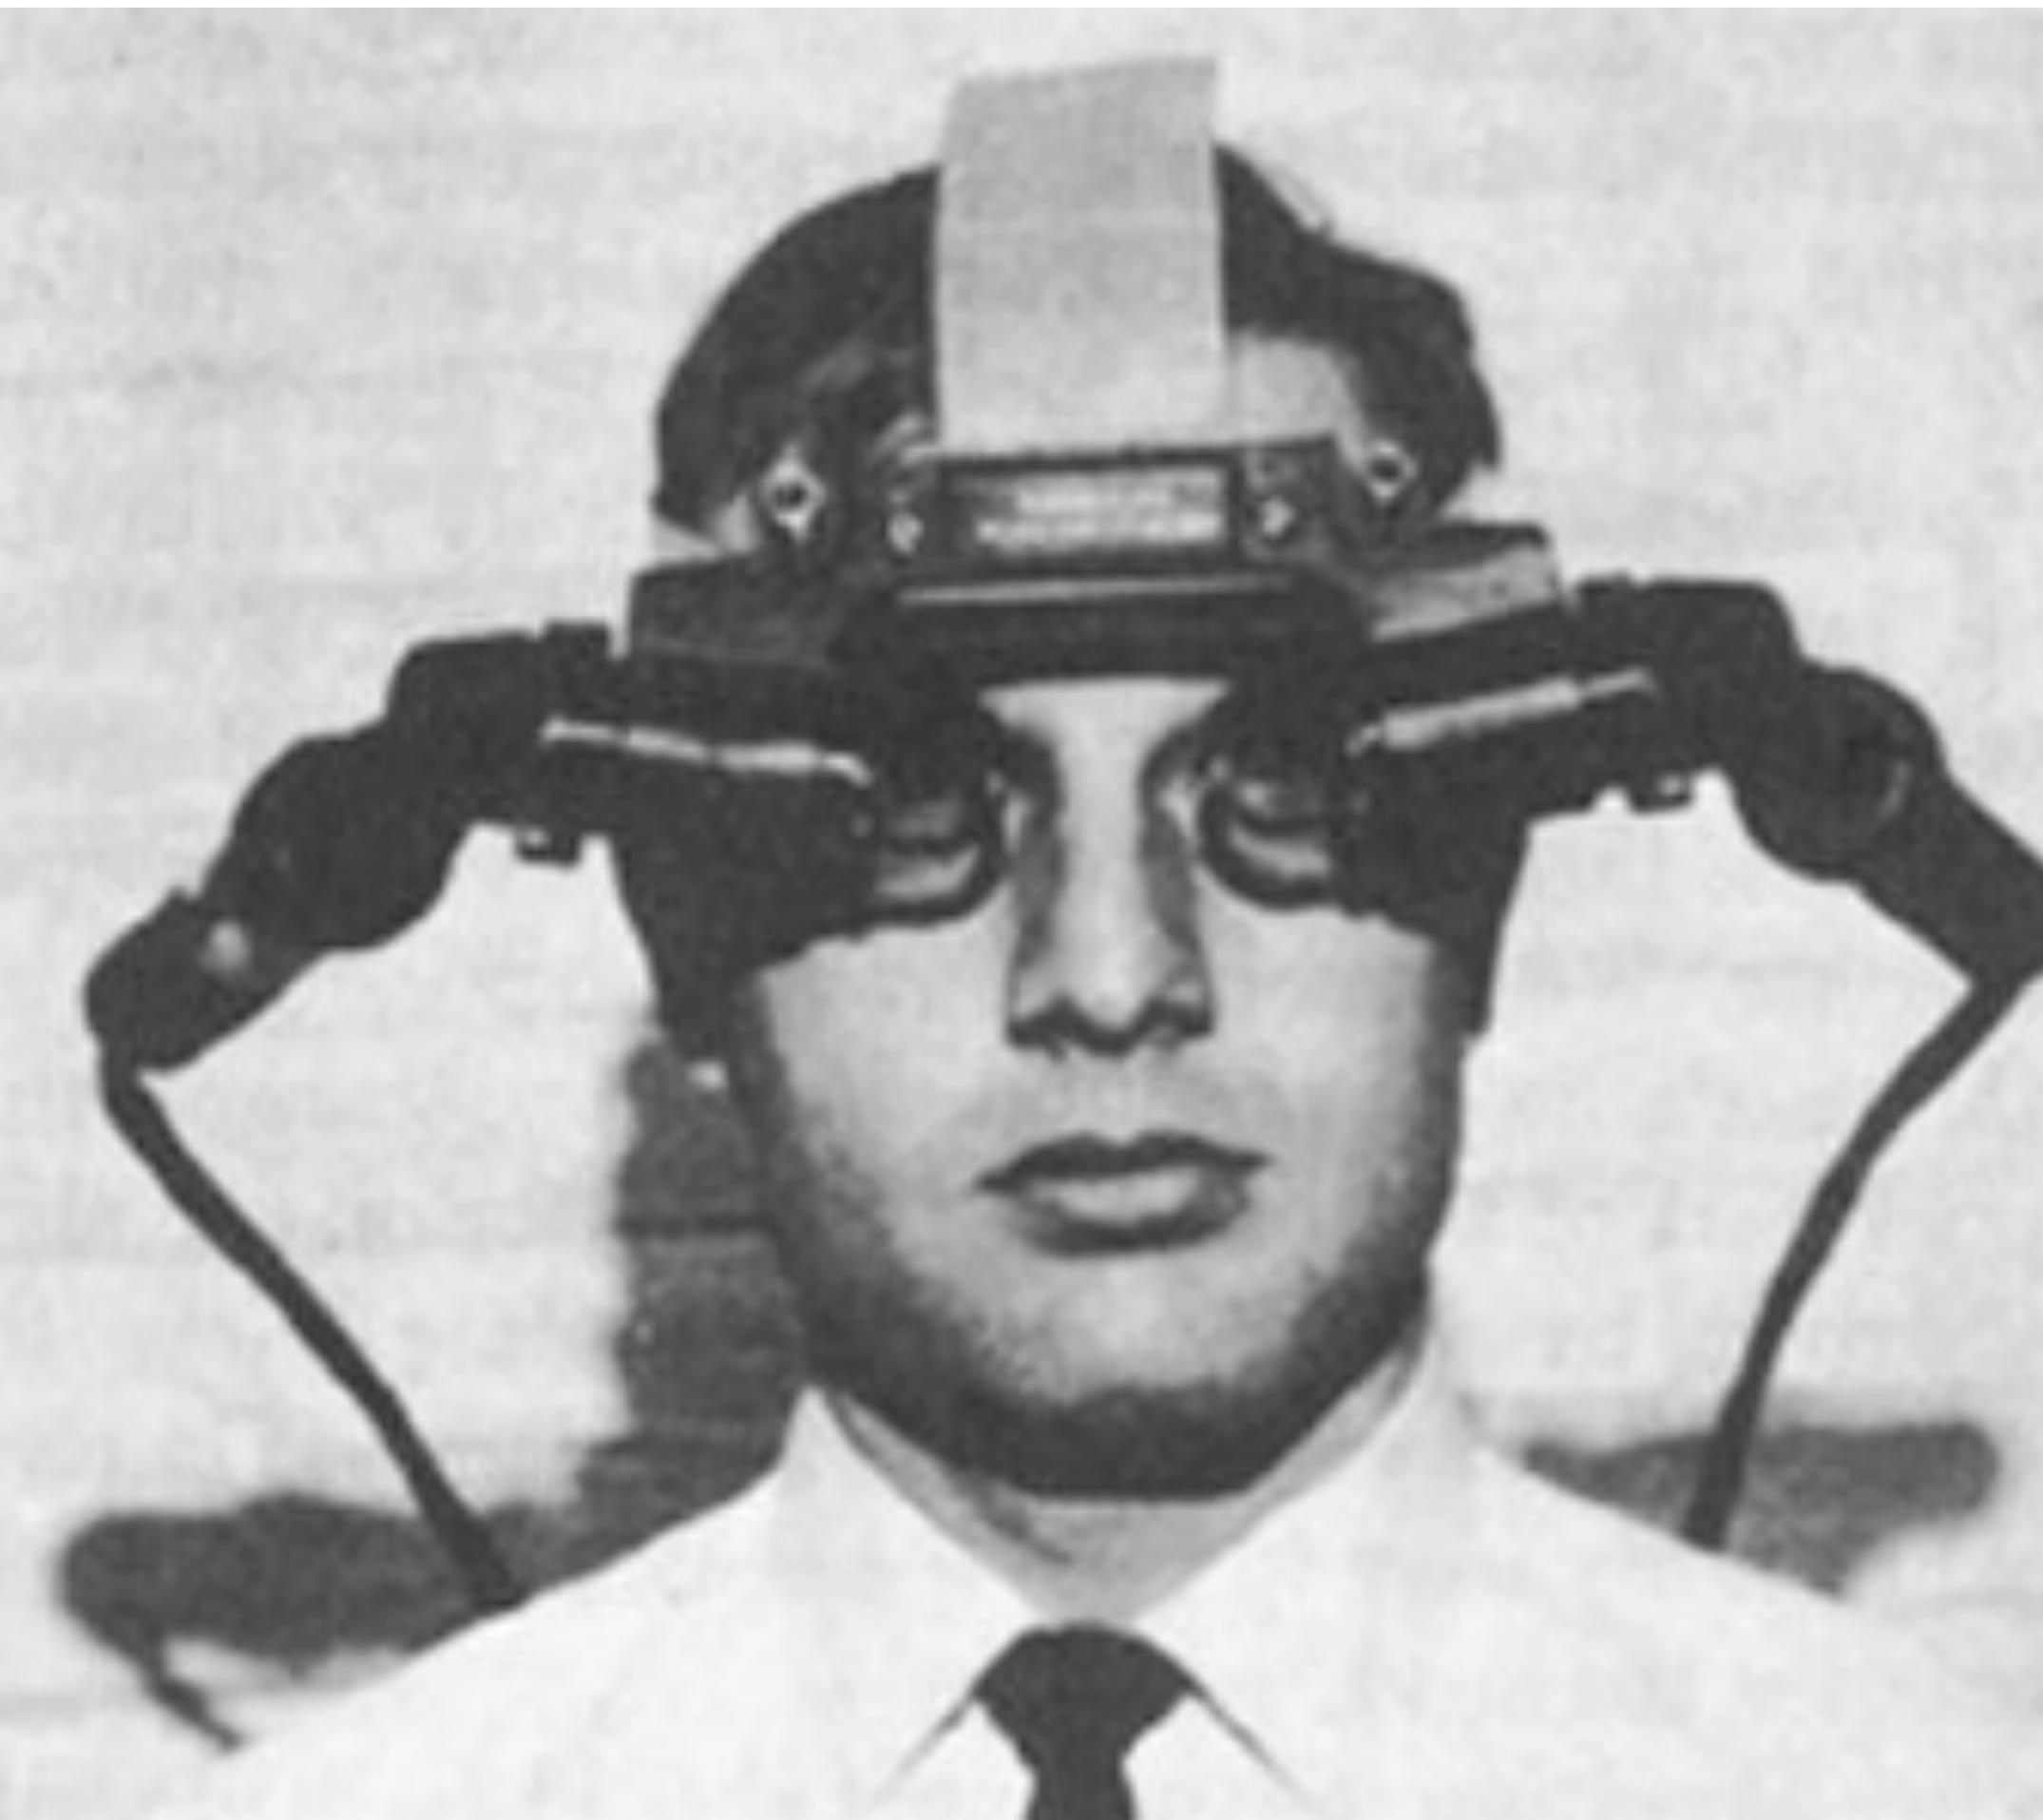
# 1901: L. FRANK BAUM



"It consists of this pair of spectacles.  
While you wear them every one you meet  
will be marked upon the forehead with a  
letter indicating his or her character..."

- Good: G
- Evil: E
- Wise: W
- Foolish: F
- Kind: K
- Cruel: C

# 1968: IVAN SOUTHERLAND



R  
W

# 1990: TOM CAUDELL

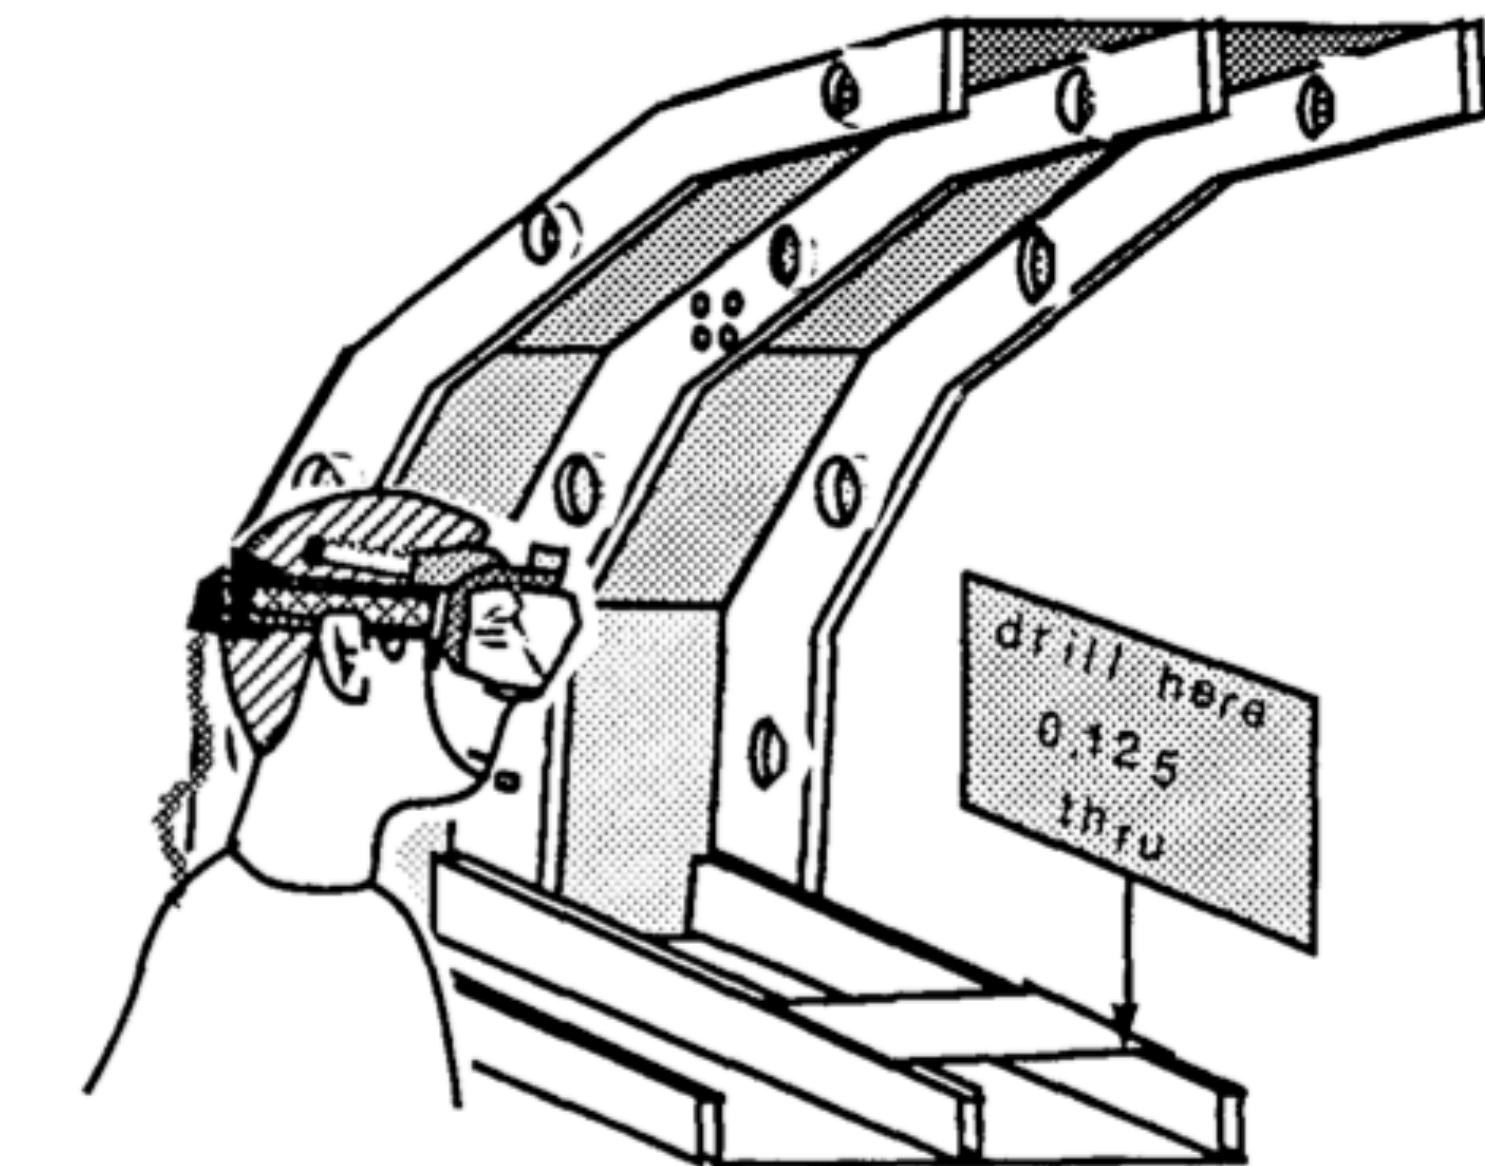


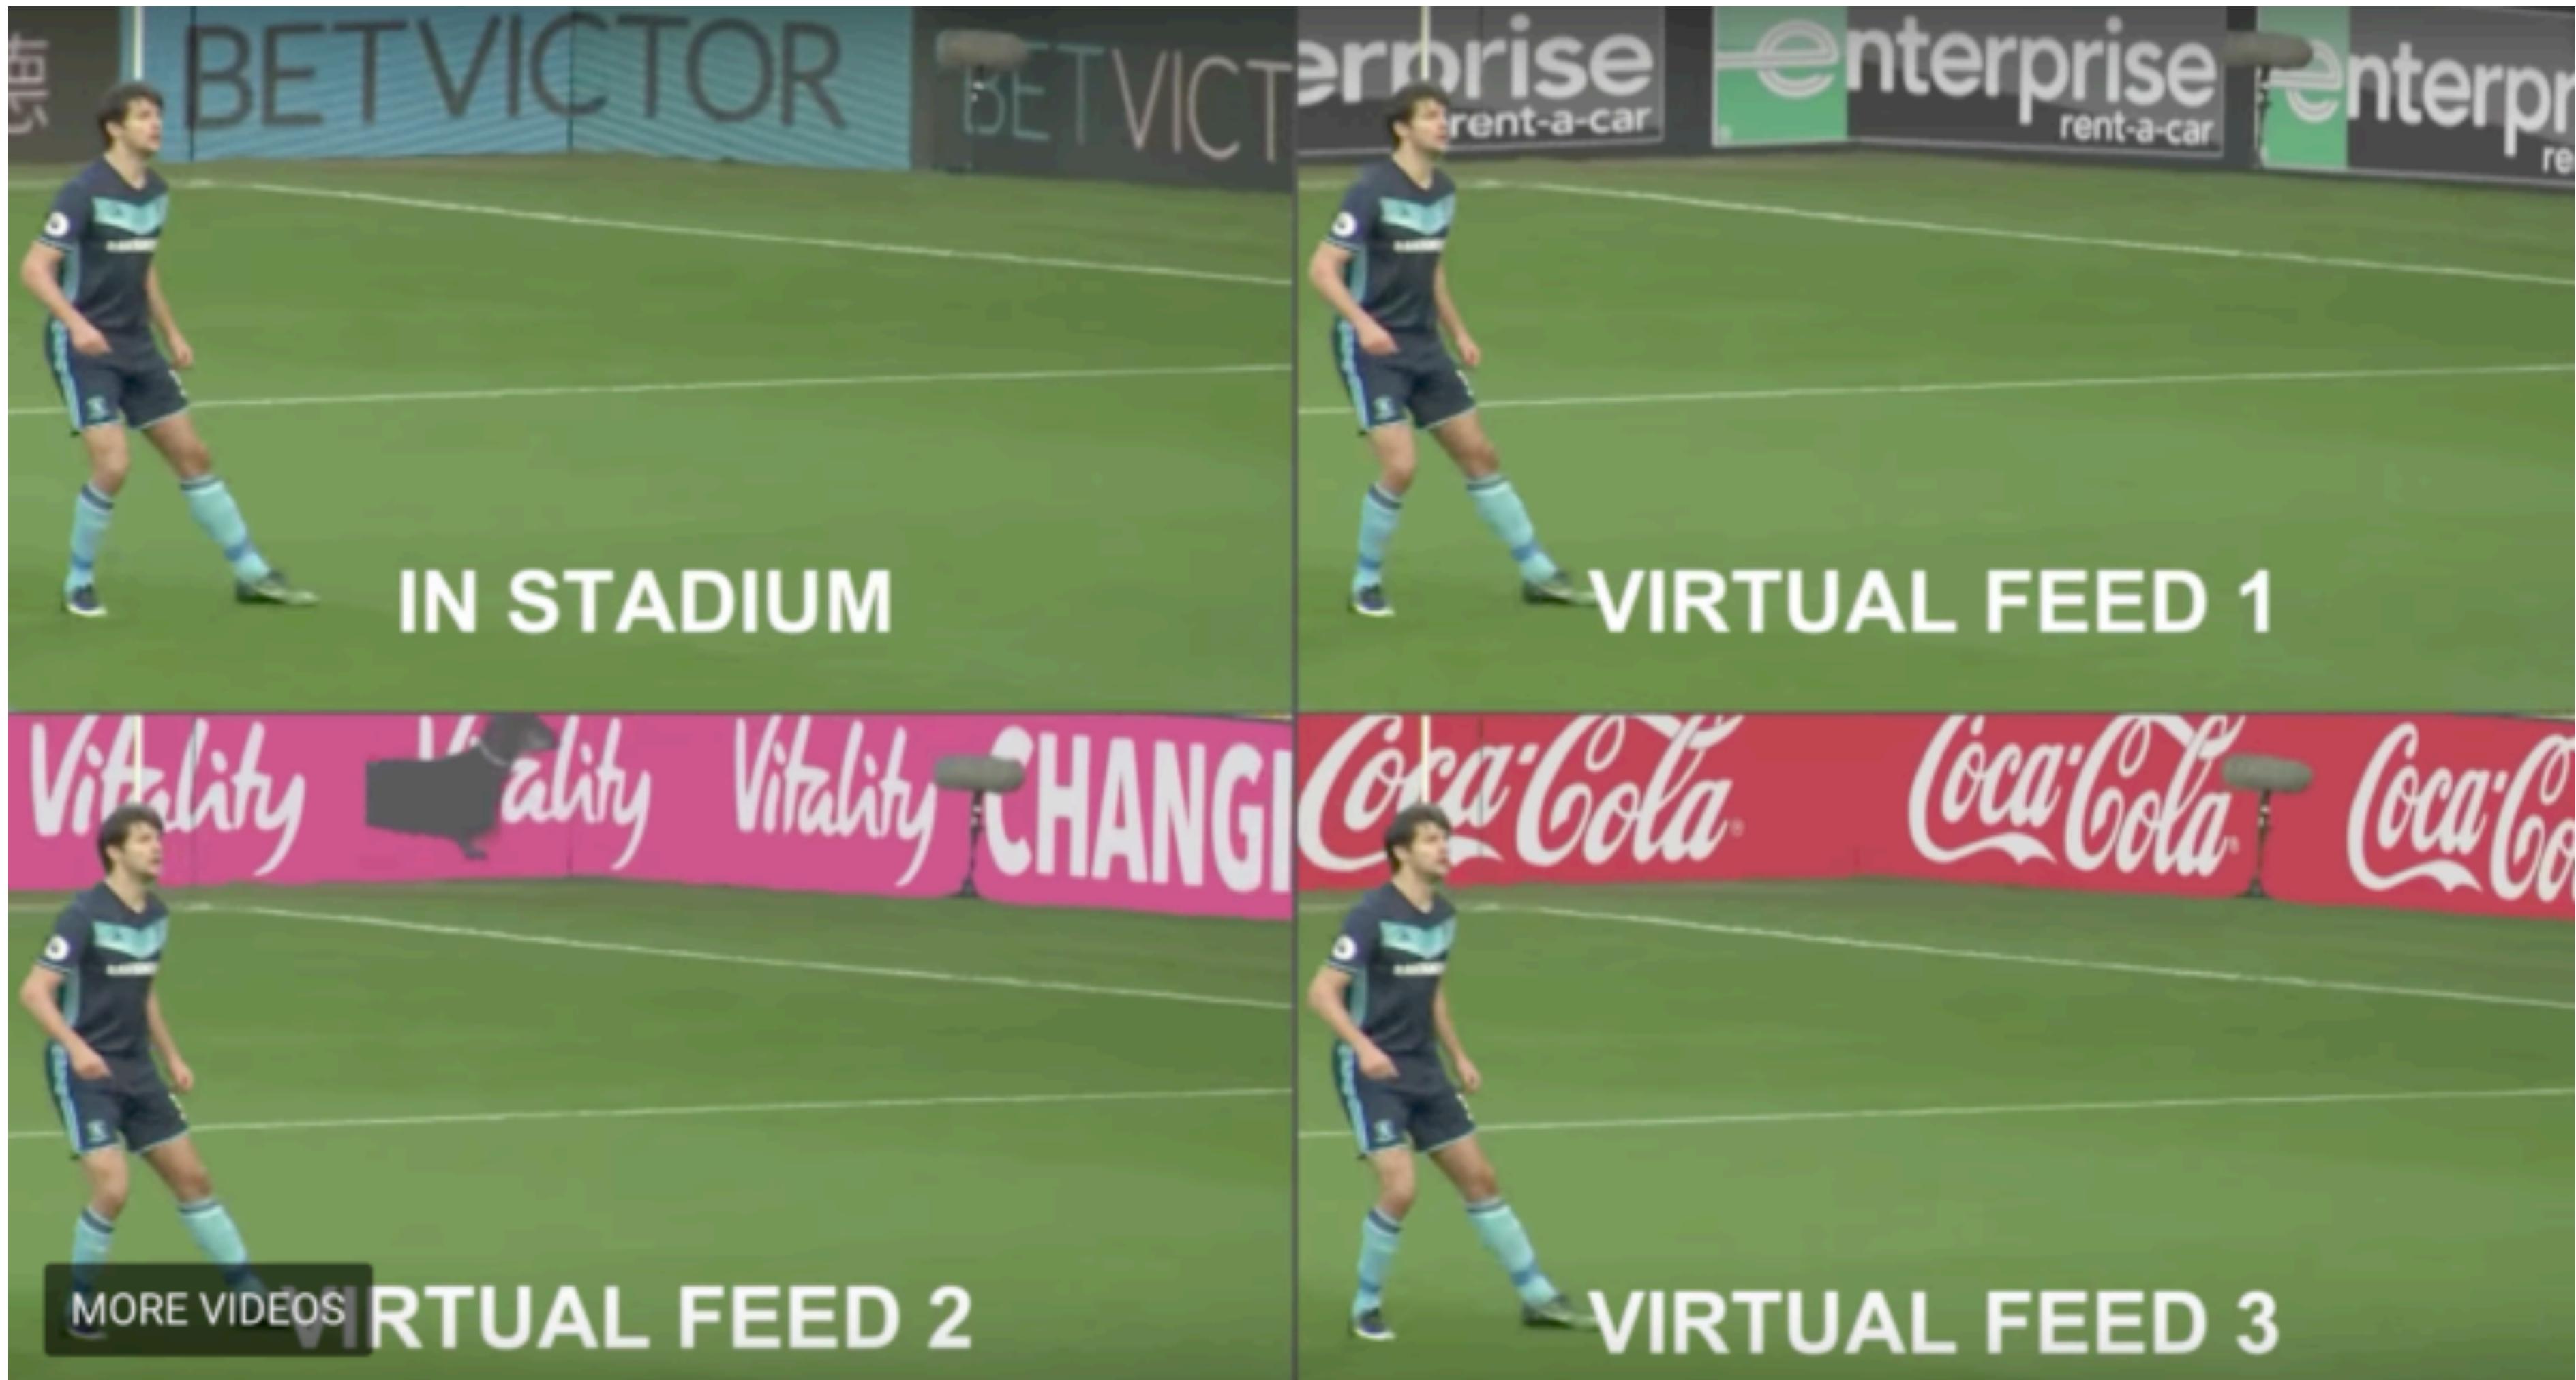
Figure 1. An application where the HUDset is used to dynamically mark the position of a drill/rivet hole inside an aircraft fuselage.

# 1980s, 1990s AND ON: STEVE MANN



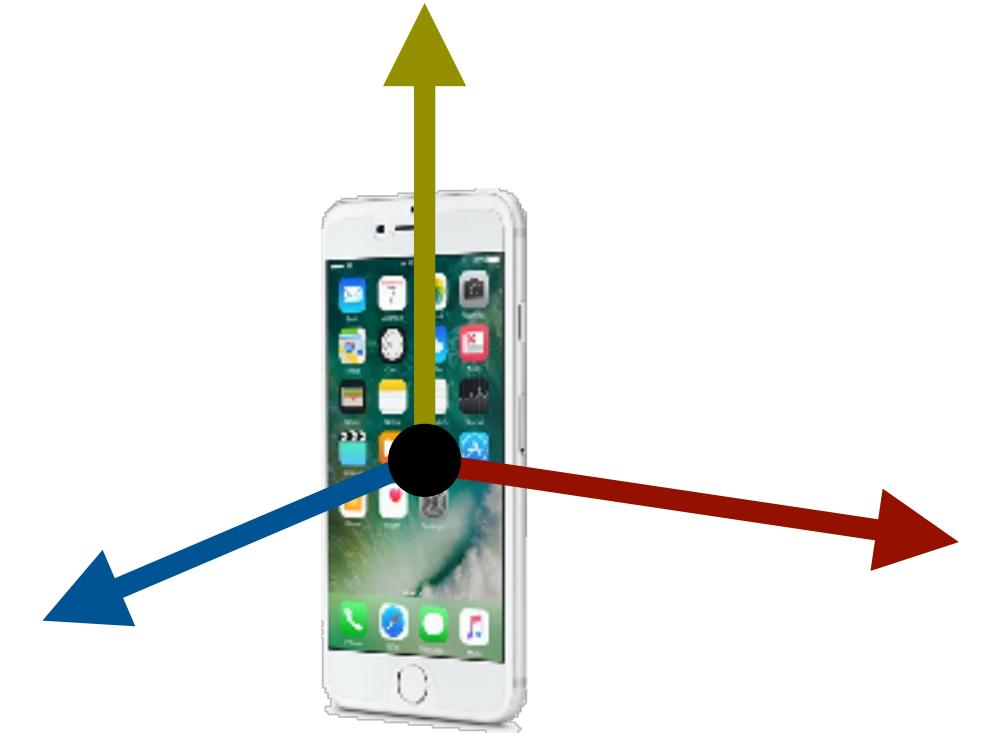
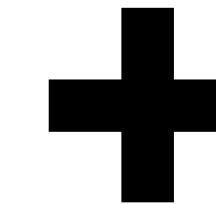
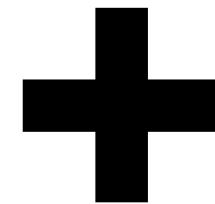
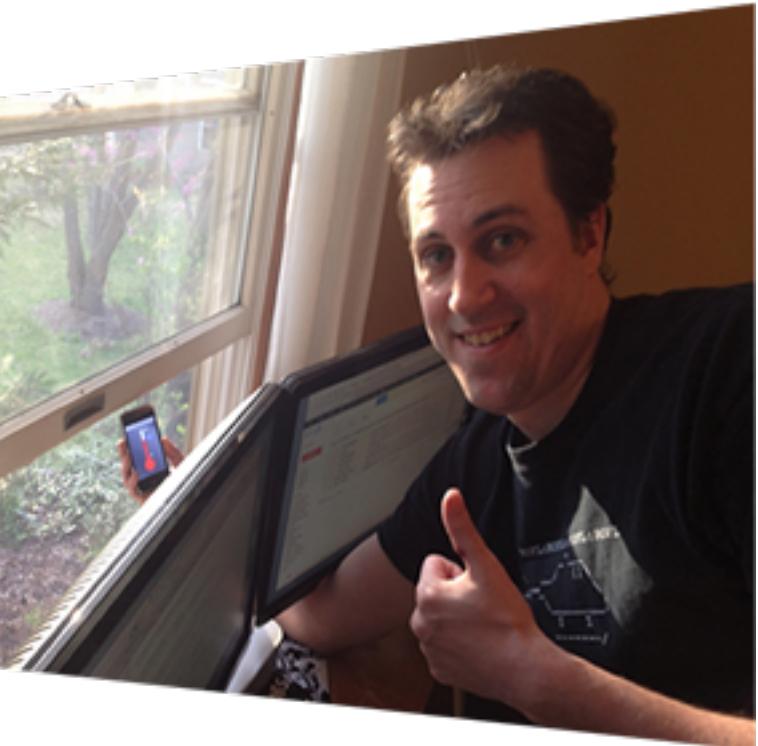
R  
W

# 1998: SPORTS BROADCASTING



R  
W

# ARKIT IN A NUTSHELL



## Real-world images

These come from a camera, and can simply be a backdrop for your app (think Pokémon Go), or can provide input for it (think IKEA Place).

## Virtual images

These are 2D or 3D images drawn on top of the real-world images from the camera.

## Sensor smarts

This is the AR application's ability to detect its position and orientation, as well as objects and conditions in the real world.



# ARKIT USER REQUIREMENTS



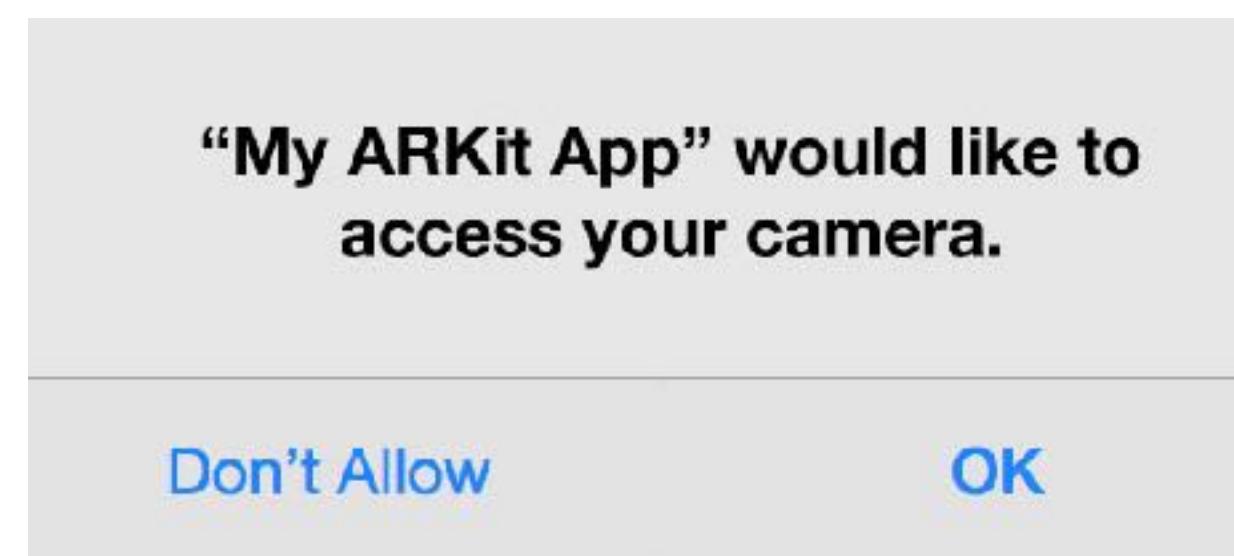
## iDevice with A9 processor or later

- iPhone: 6S / 6S Plus / 7S / 7S Plus / 8 / 8 Plus / X
- iPad Pro: 9.7" / 10.5" / 12.9"
- iPad: 2017 model

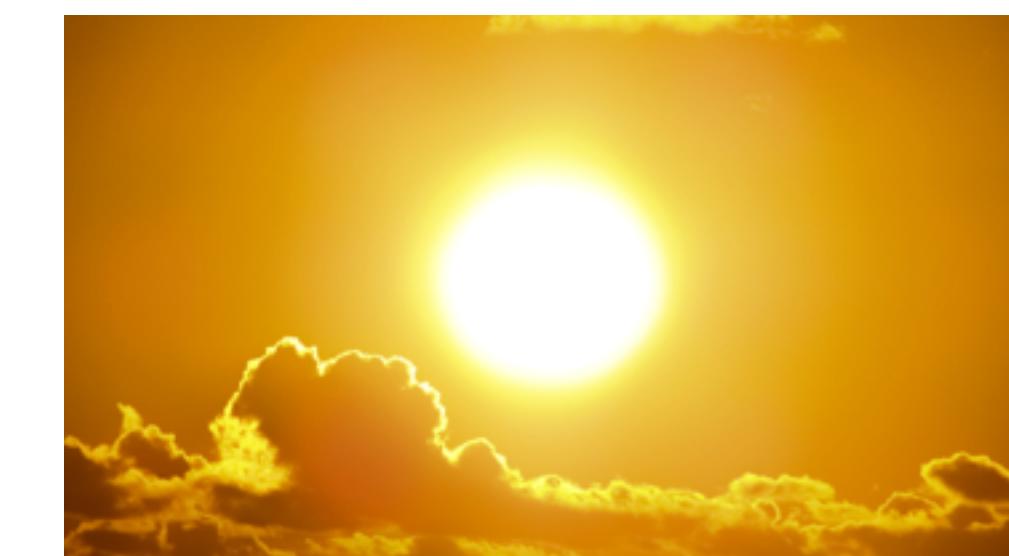


## iOS 11

The first iOS version that supports ARKit



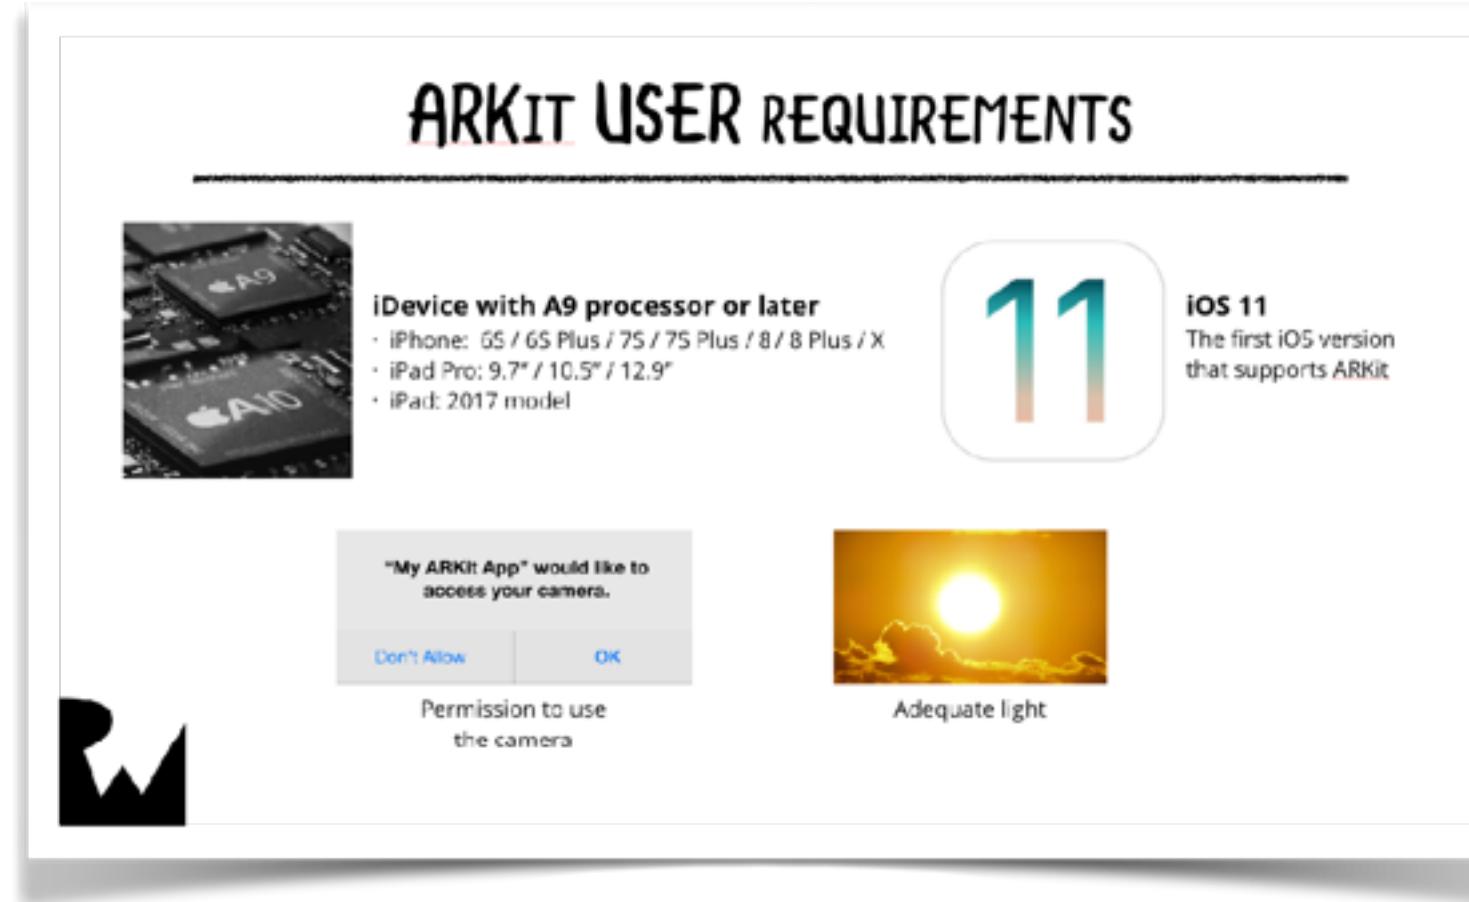
Permission to use  
the camera



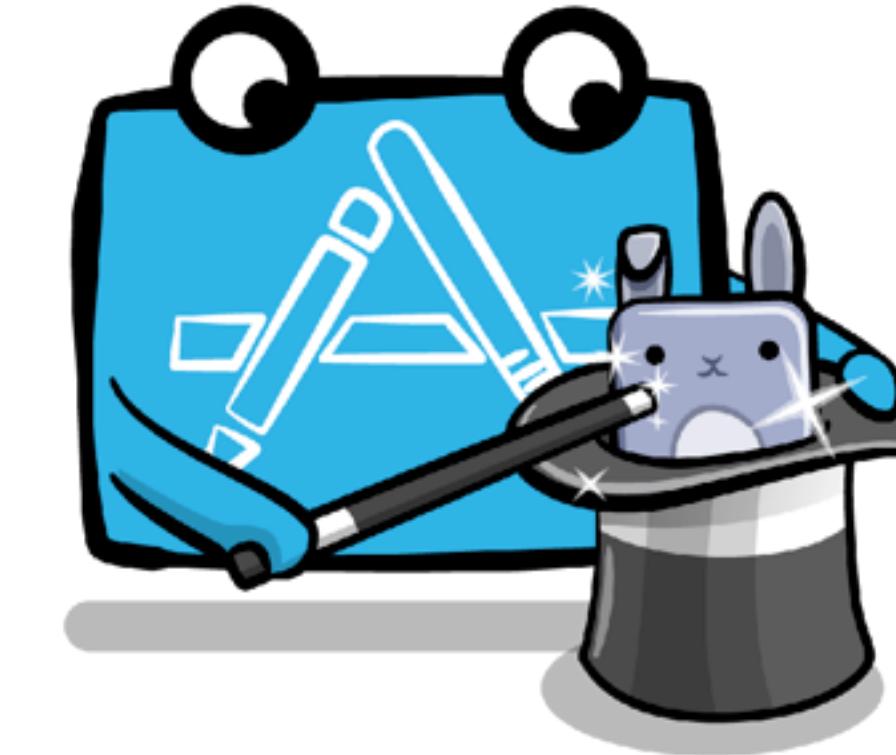
Adequate light



# ARKIT DEVELOPER REQUIREMENTS



The user requirements



Xcode 9.3 or later



SpriteKit / SceneKit basics



Readiness to walk about  
and wave your iDevice

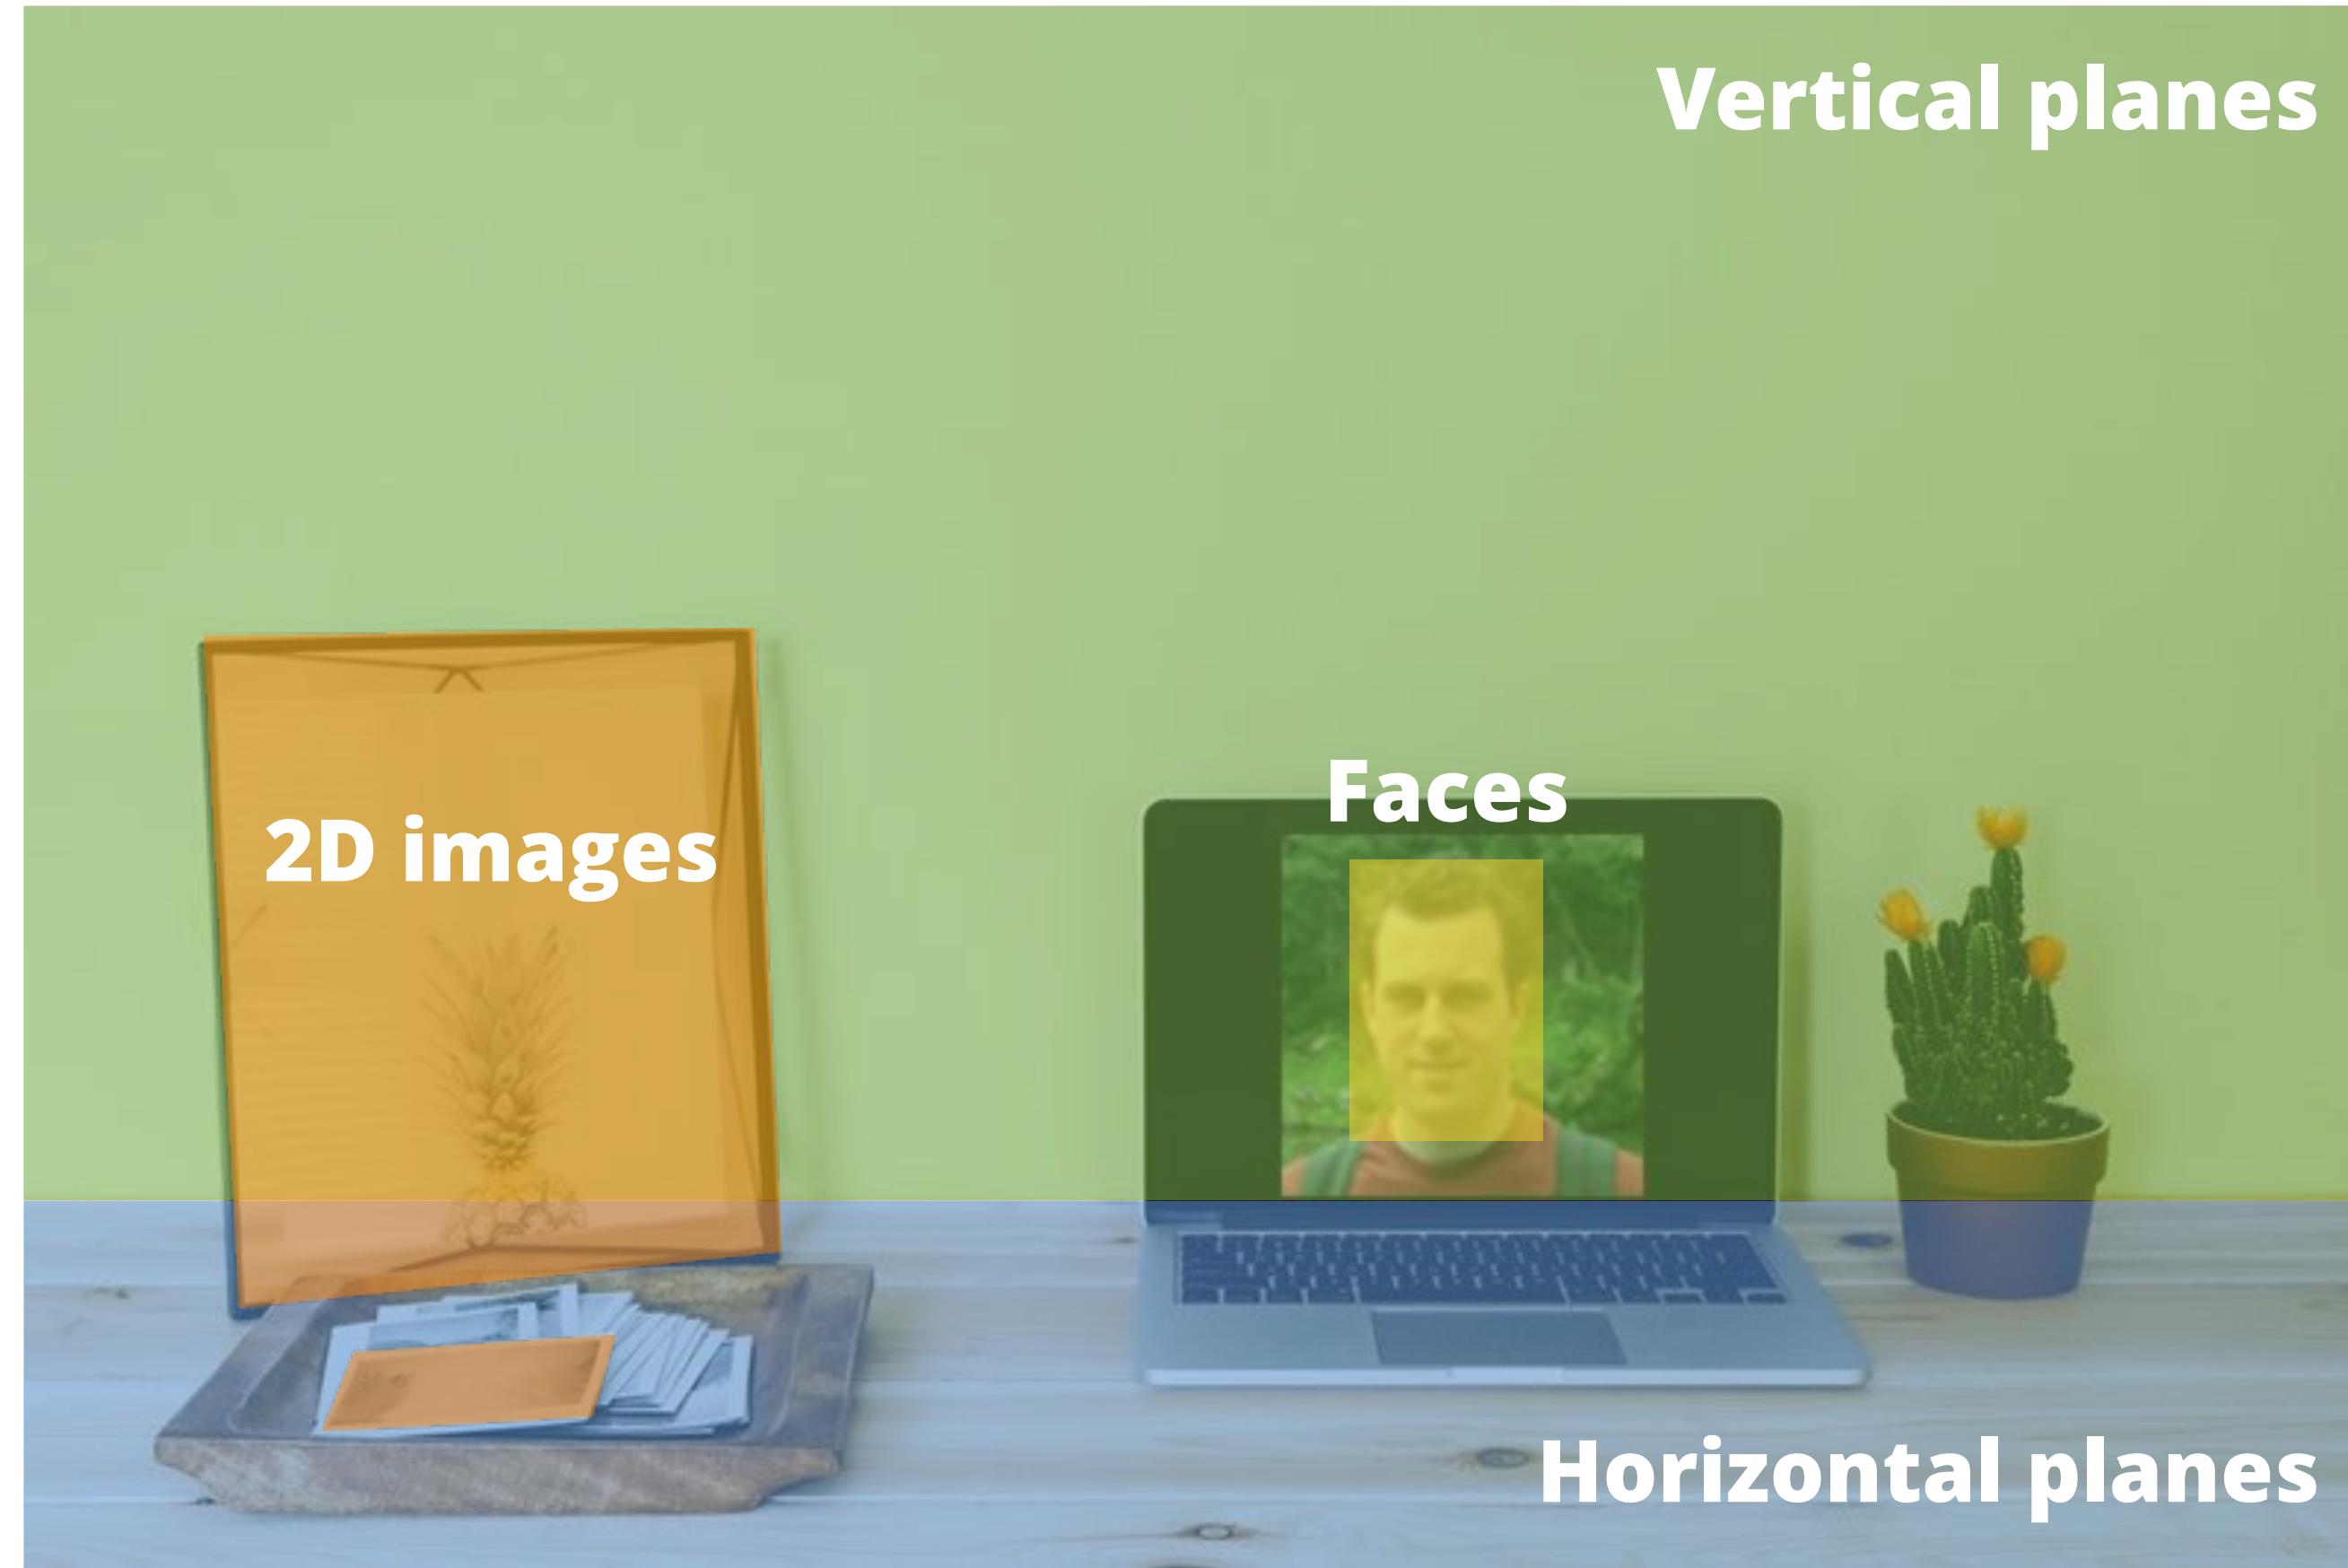


Willingness to do  
at least a little 3D  
math...

...and deal with  
changes and upgrades.



# WHAT REAL-WORLD THINGS CAN ARKIT IDENTIFY?



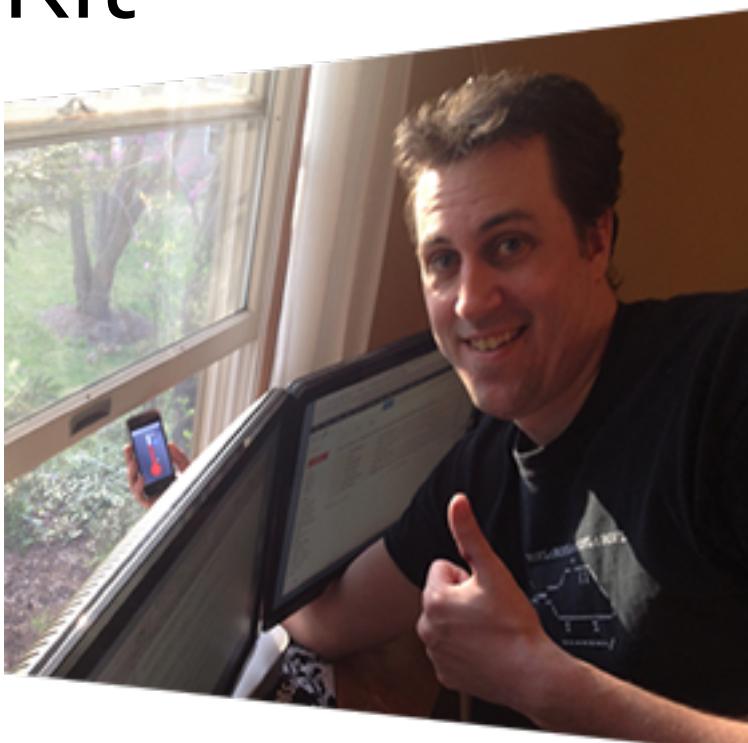
# SPRITEKIT AR VS. SCENEKIT AR

## SpriteKit AR:

Overlaying camera images with  
2D graphics using SpriteKit



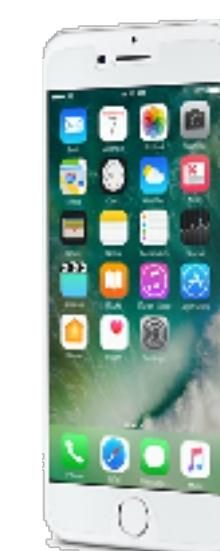
2D graphics drawn to  
an **AR SpriteKit view (ARSKView)**



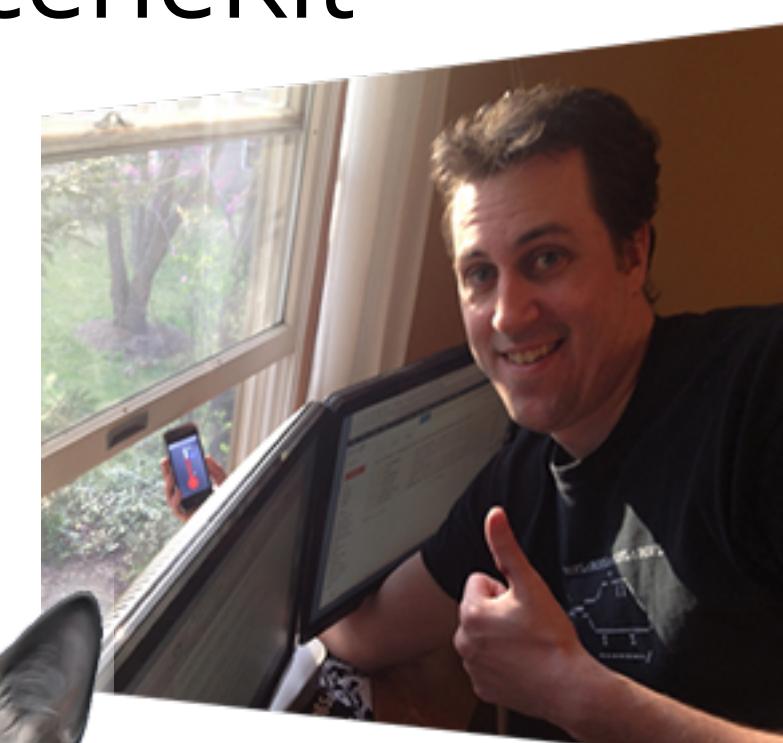
Camera image drawn to  
the same AR SpriteKit  
view

## SceneKit AR:

Overlaying camera images with  
3D graphics using SceneKit



3D graphics drawn to  
an **AR SceneKit view (ARSCNView)**



Camera image drawn to  
the same AR SceneKit  
view

# THE DEMOS



## Demo 1: *Happy AR Painter*

In which you begin your ARKit journey by paying homage to the great Bob Ross by building an app that paints in real-world 3D.



## Demo 2: *Raykea*

Let's make our own version of the most popular ARKit app. Why should makers of semi-disposable furniture have all the fun?



## Demo 3: *BaedekAR*

One of my clients is a cruise line company, and they're interested in this app, which is an AR museum guide.

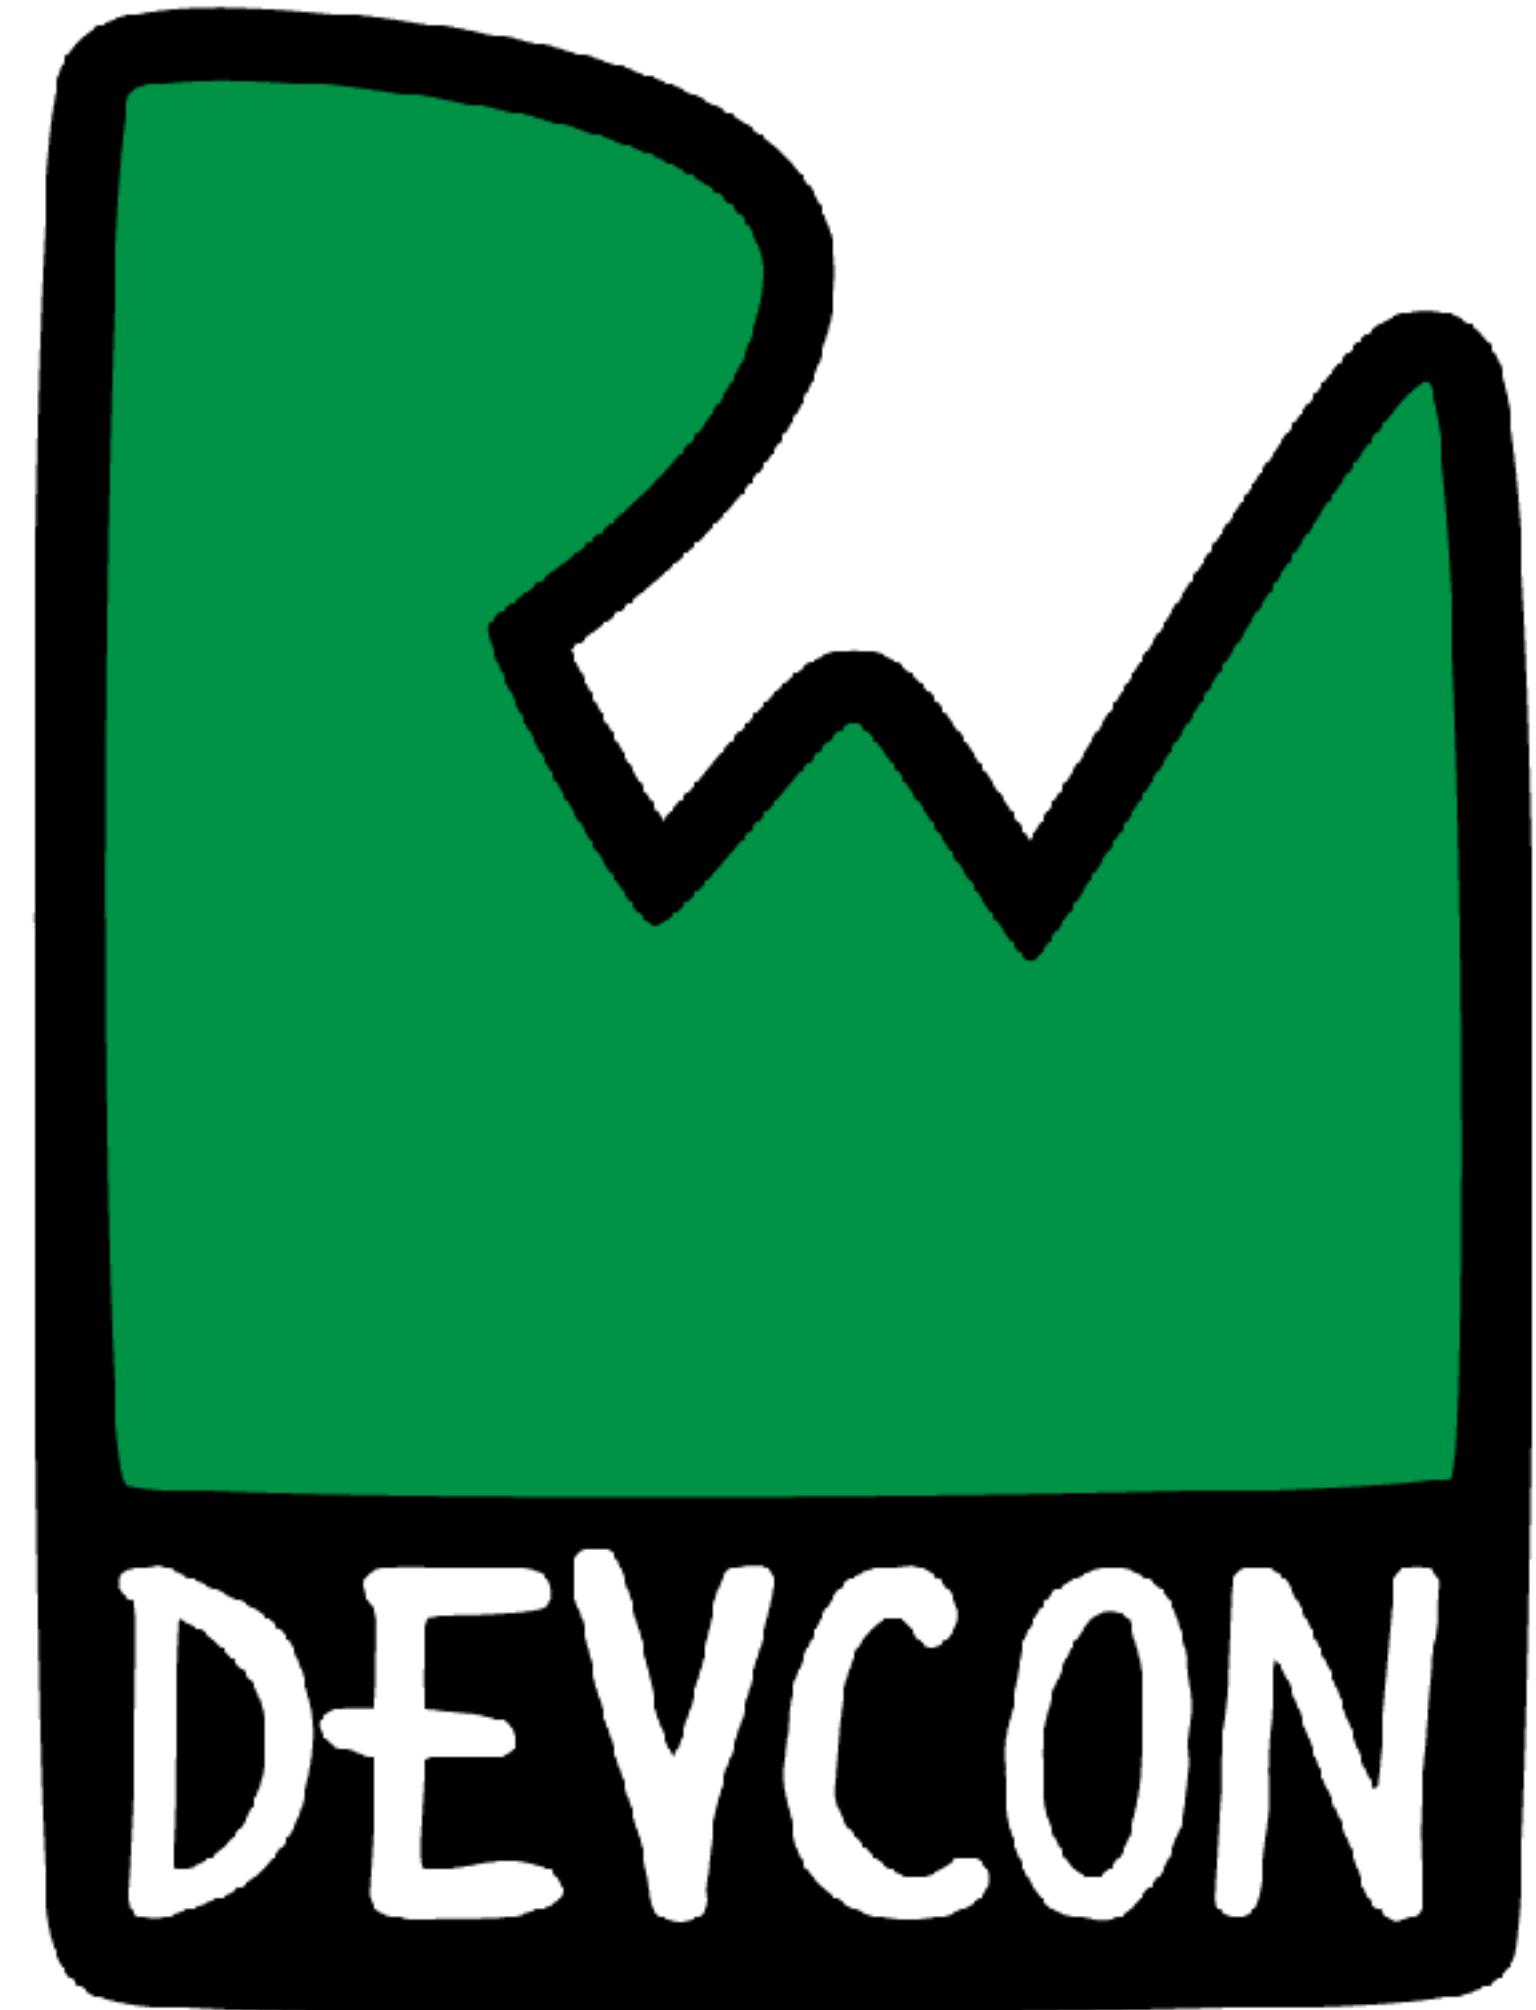


## Demo 4: *Vision Quest*

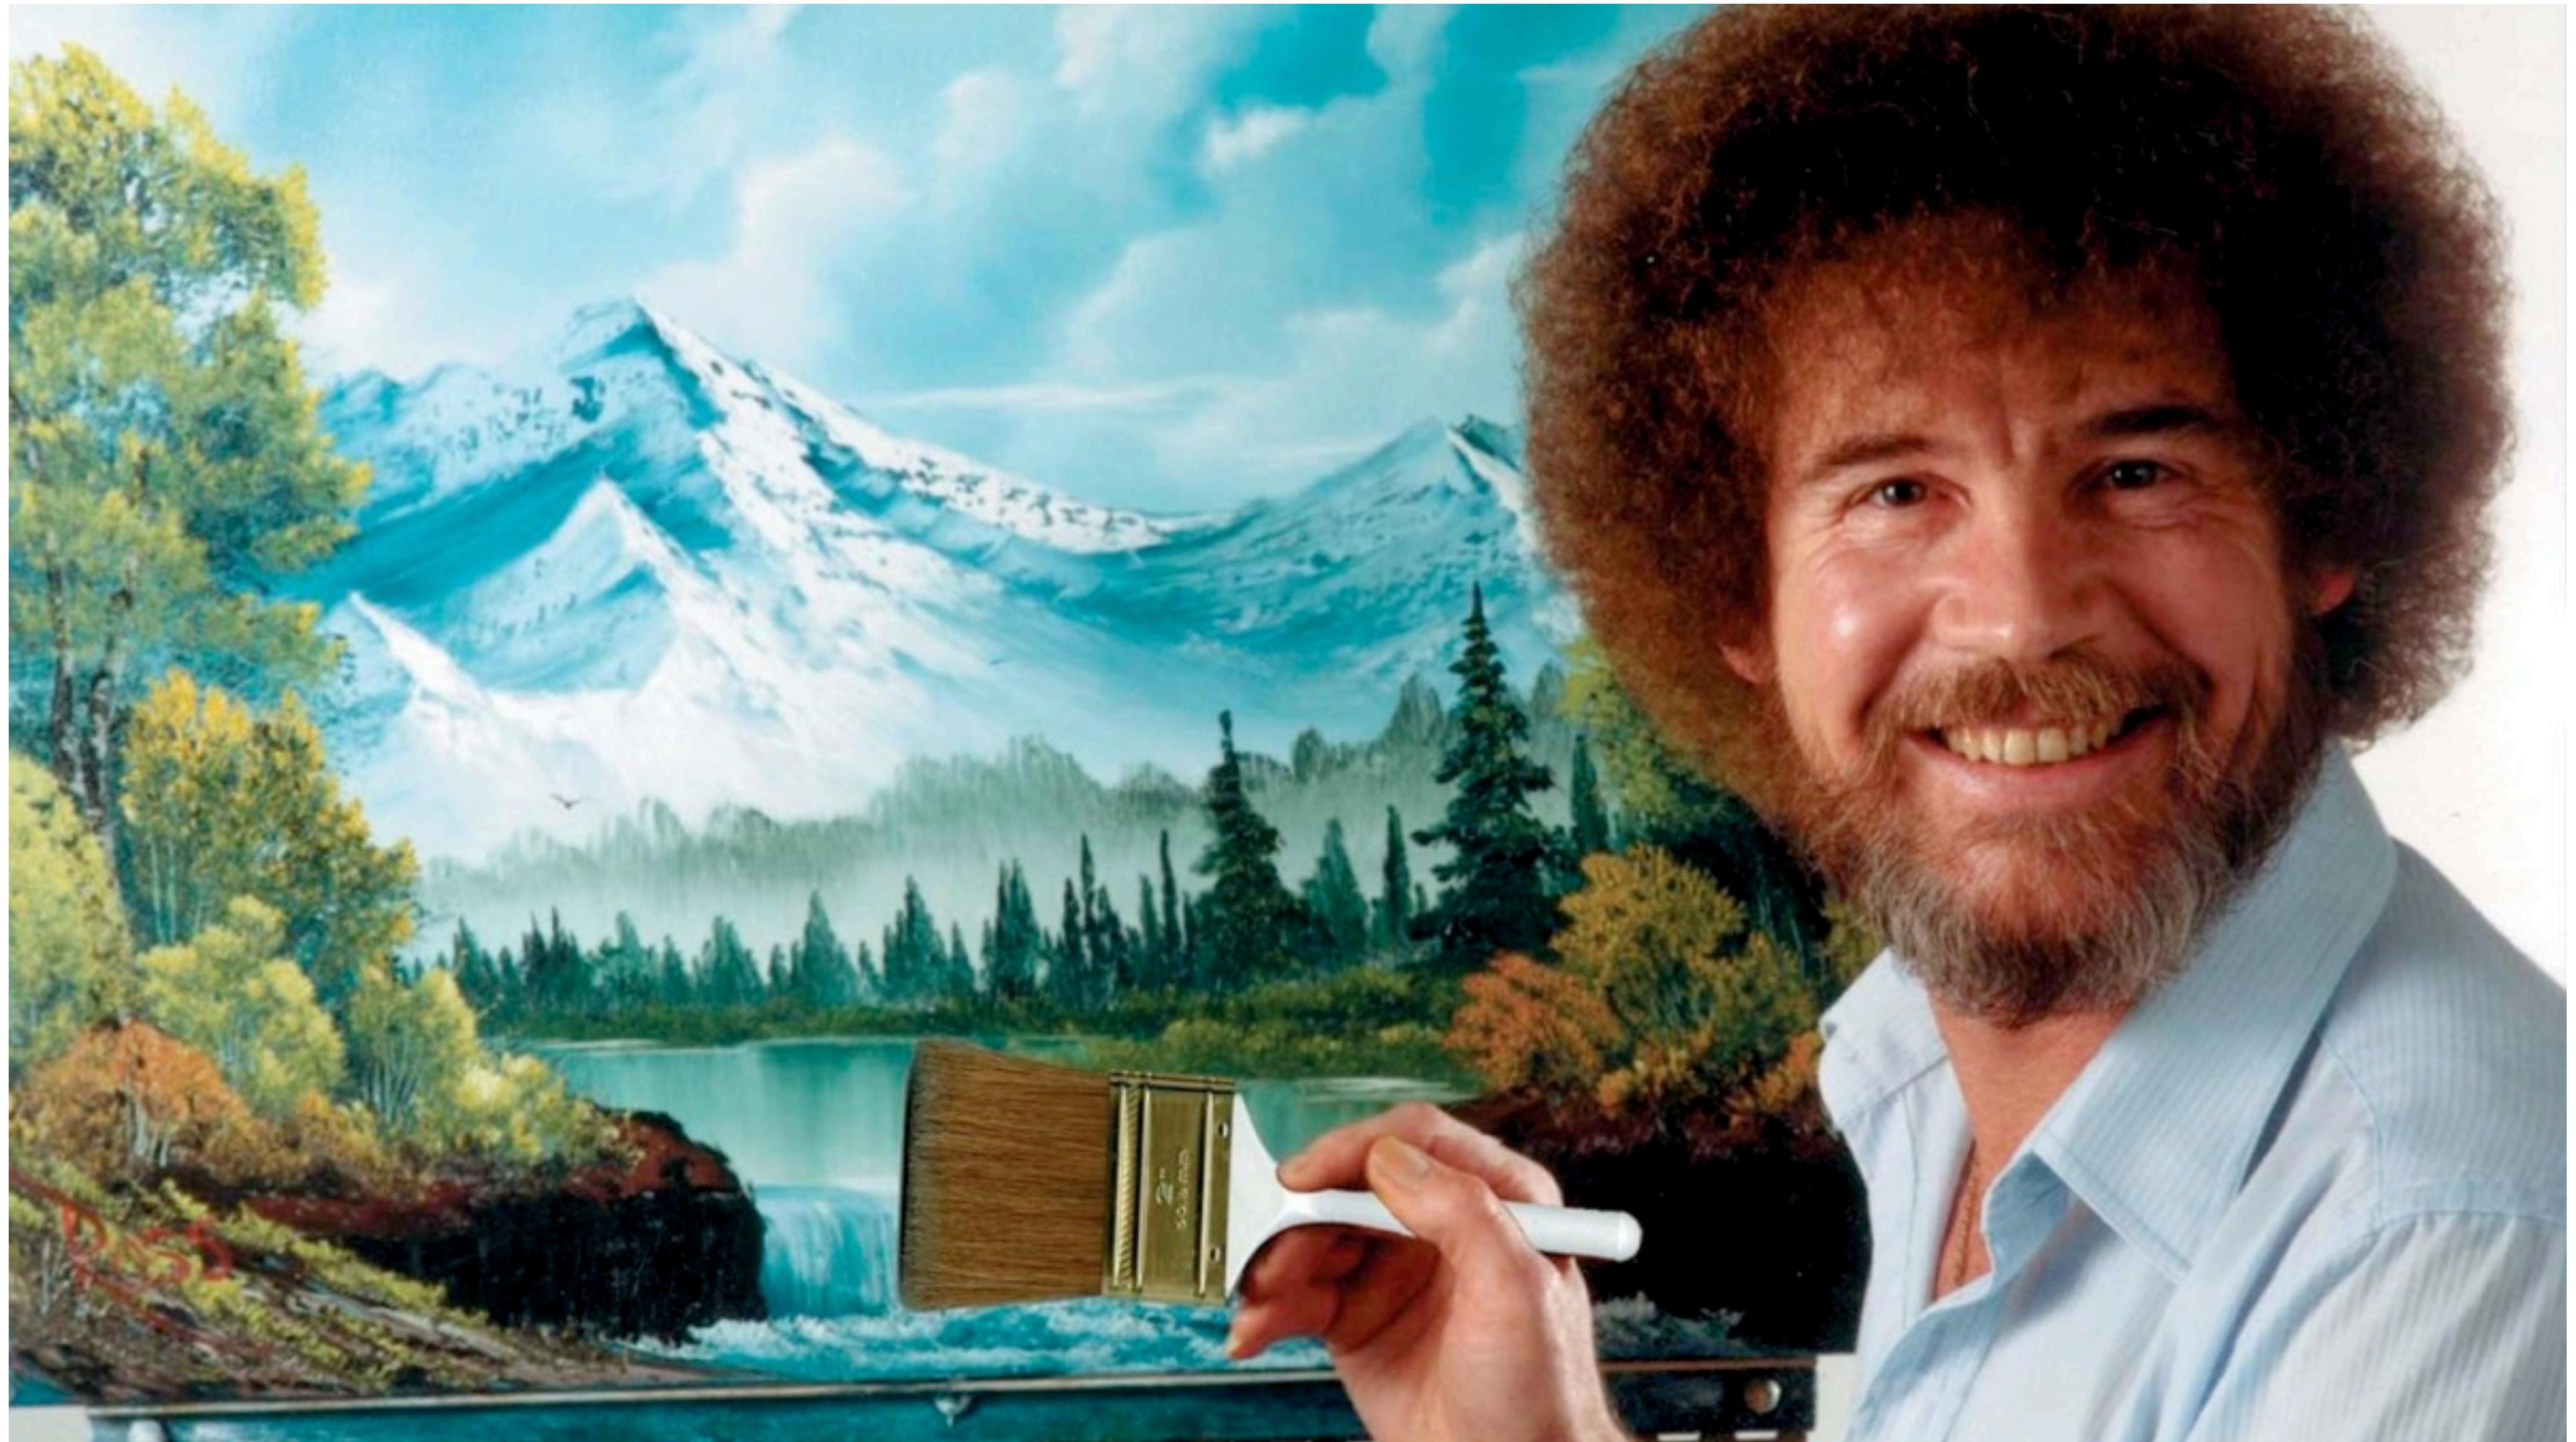
We'll close the workshop either by accidentally creating SkyNet or combine machine learning with ARKit to make an object recognizer.



# Workshop: ARKit in Depth

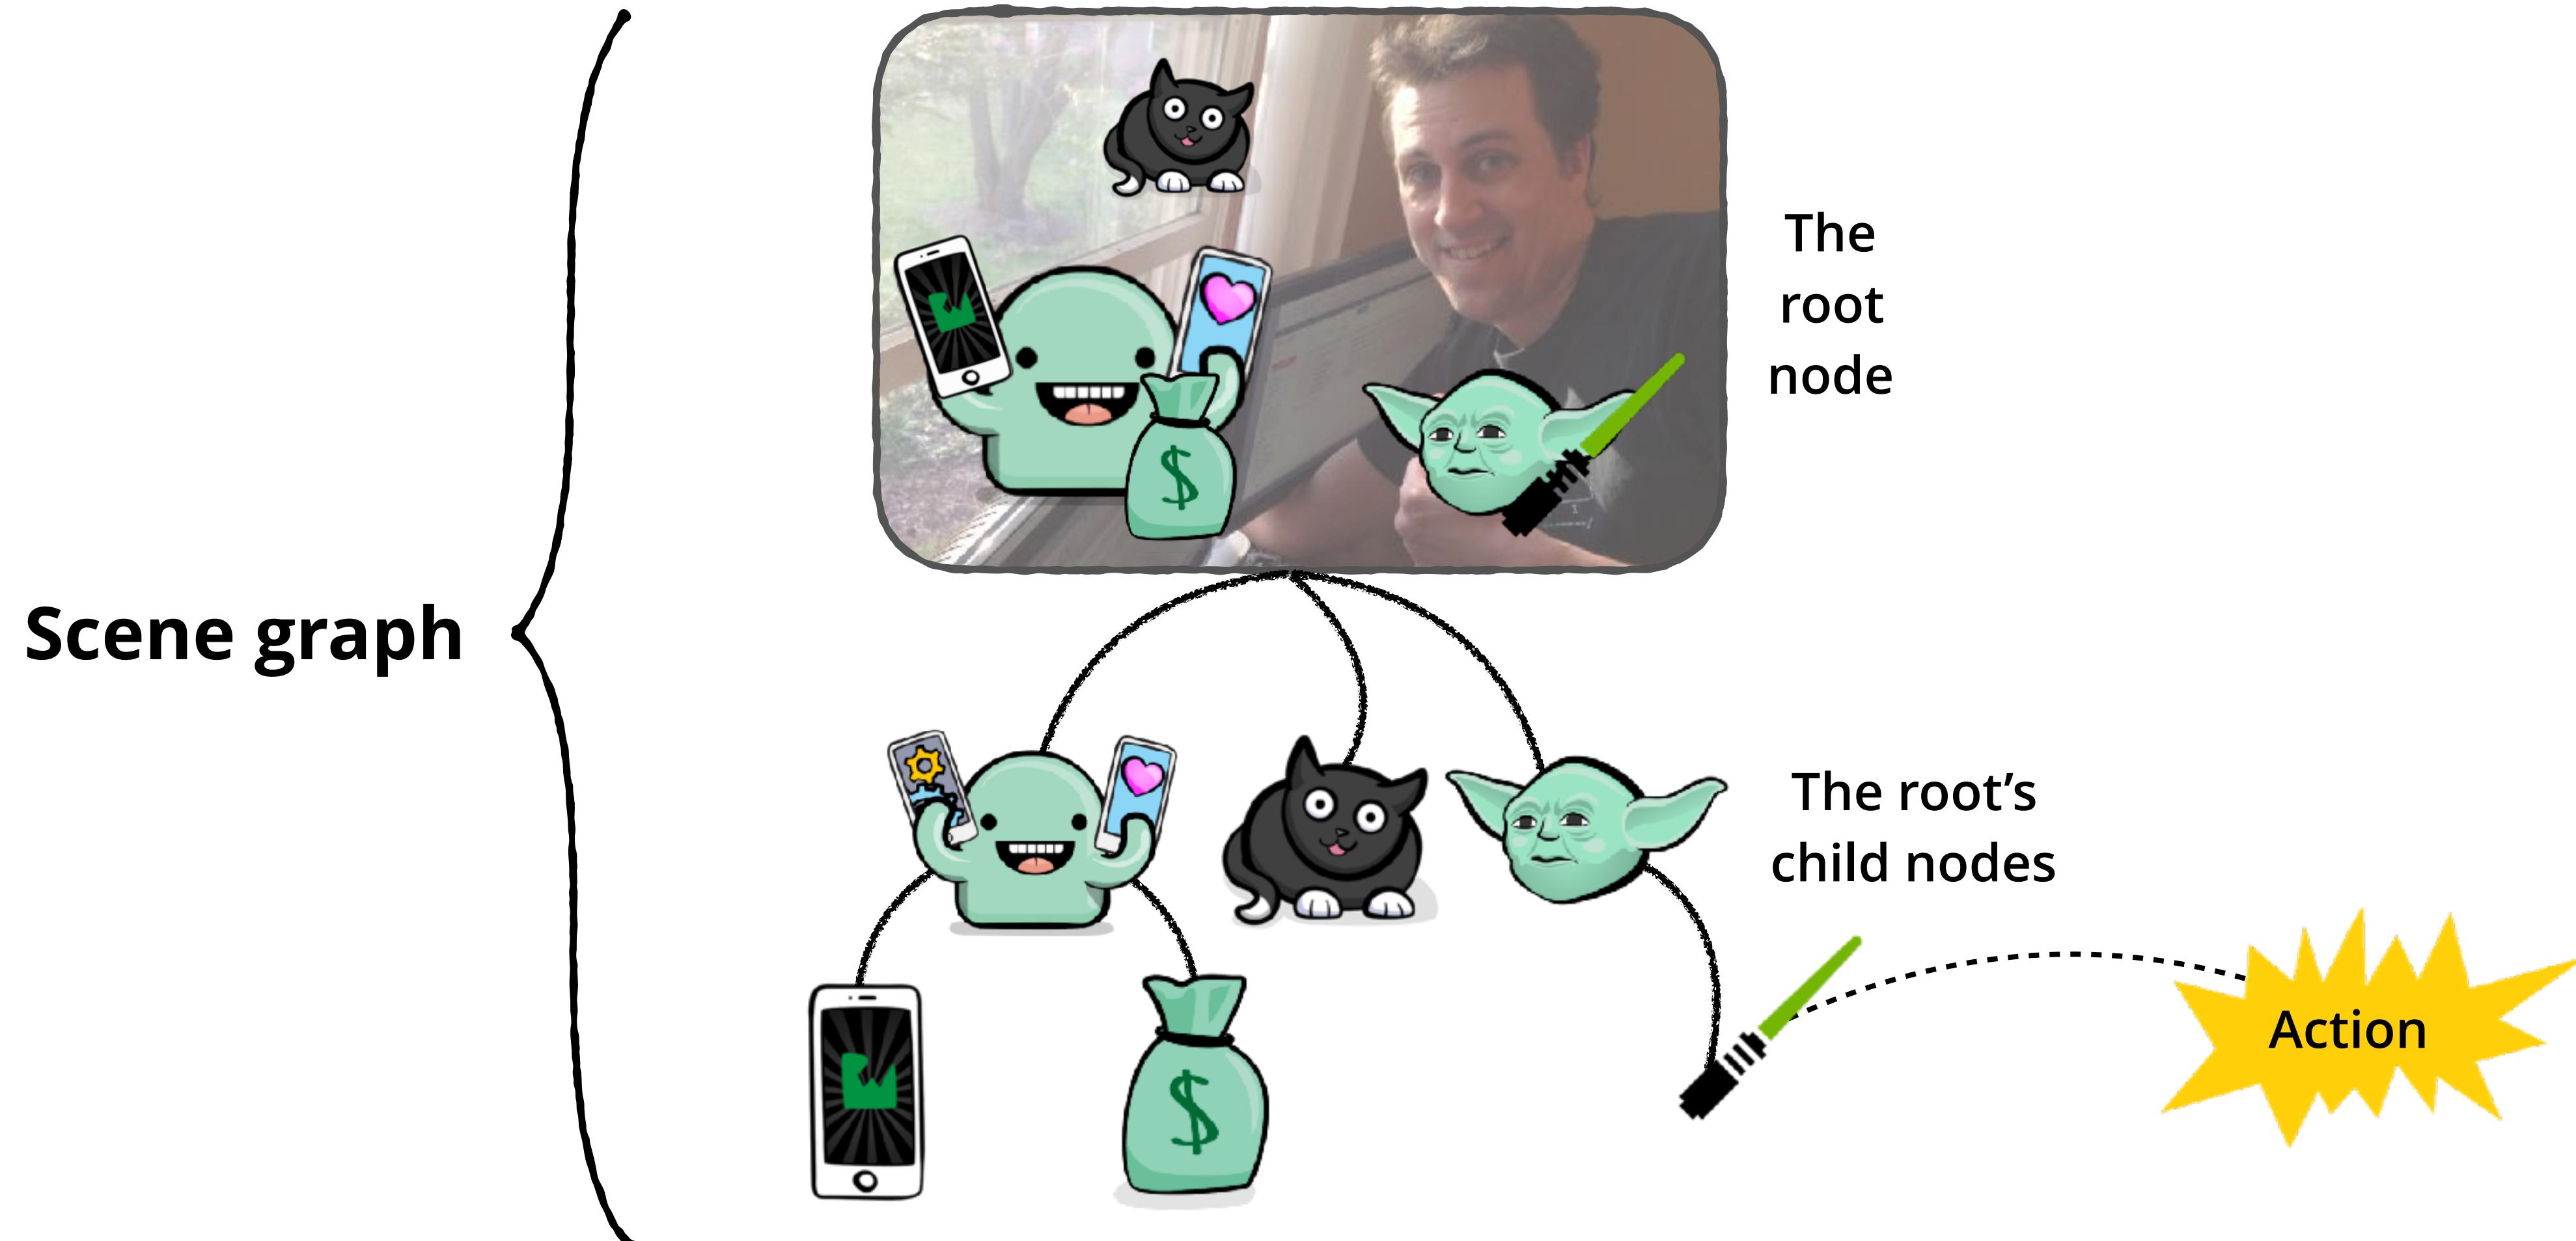


HAPPY AR PAINTER



R  
W

# QUICK SCENEKIT OVERVIEW



# SCENEKIT'S 3D COORDINATE SYSTEM

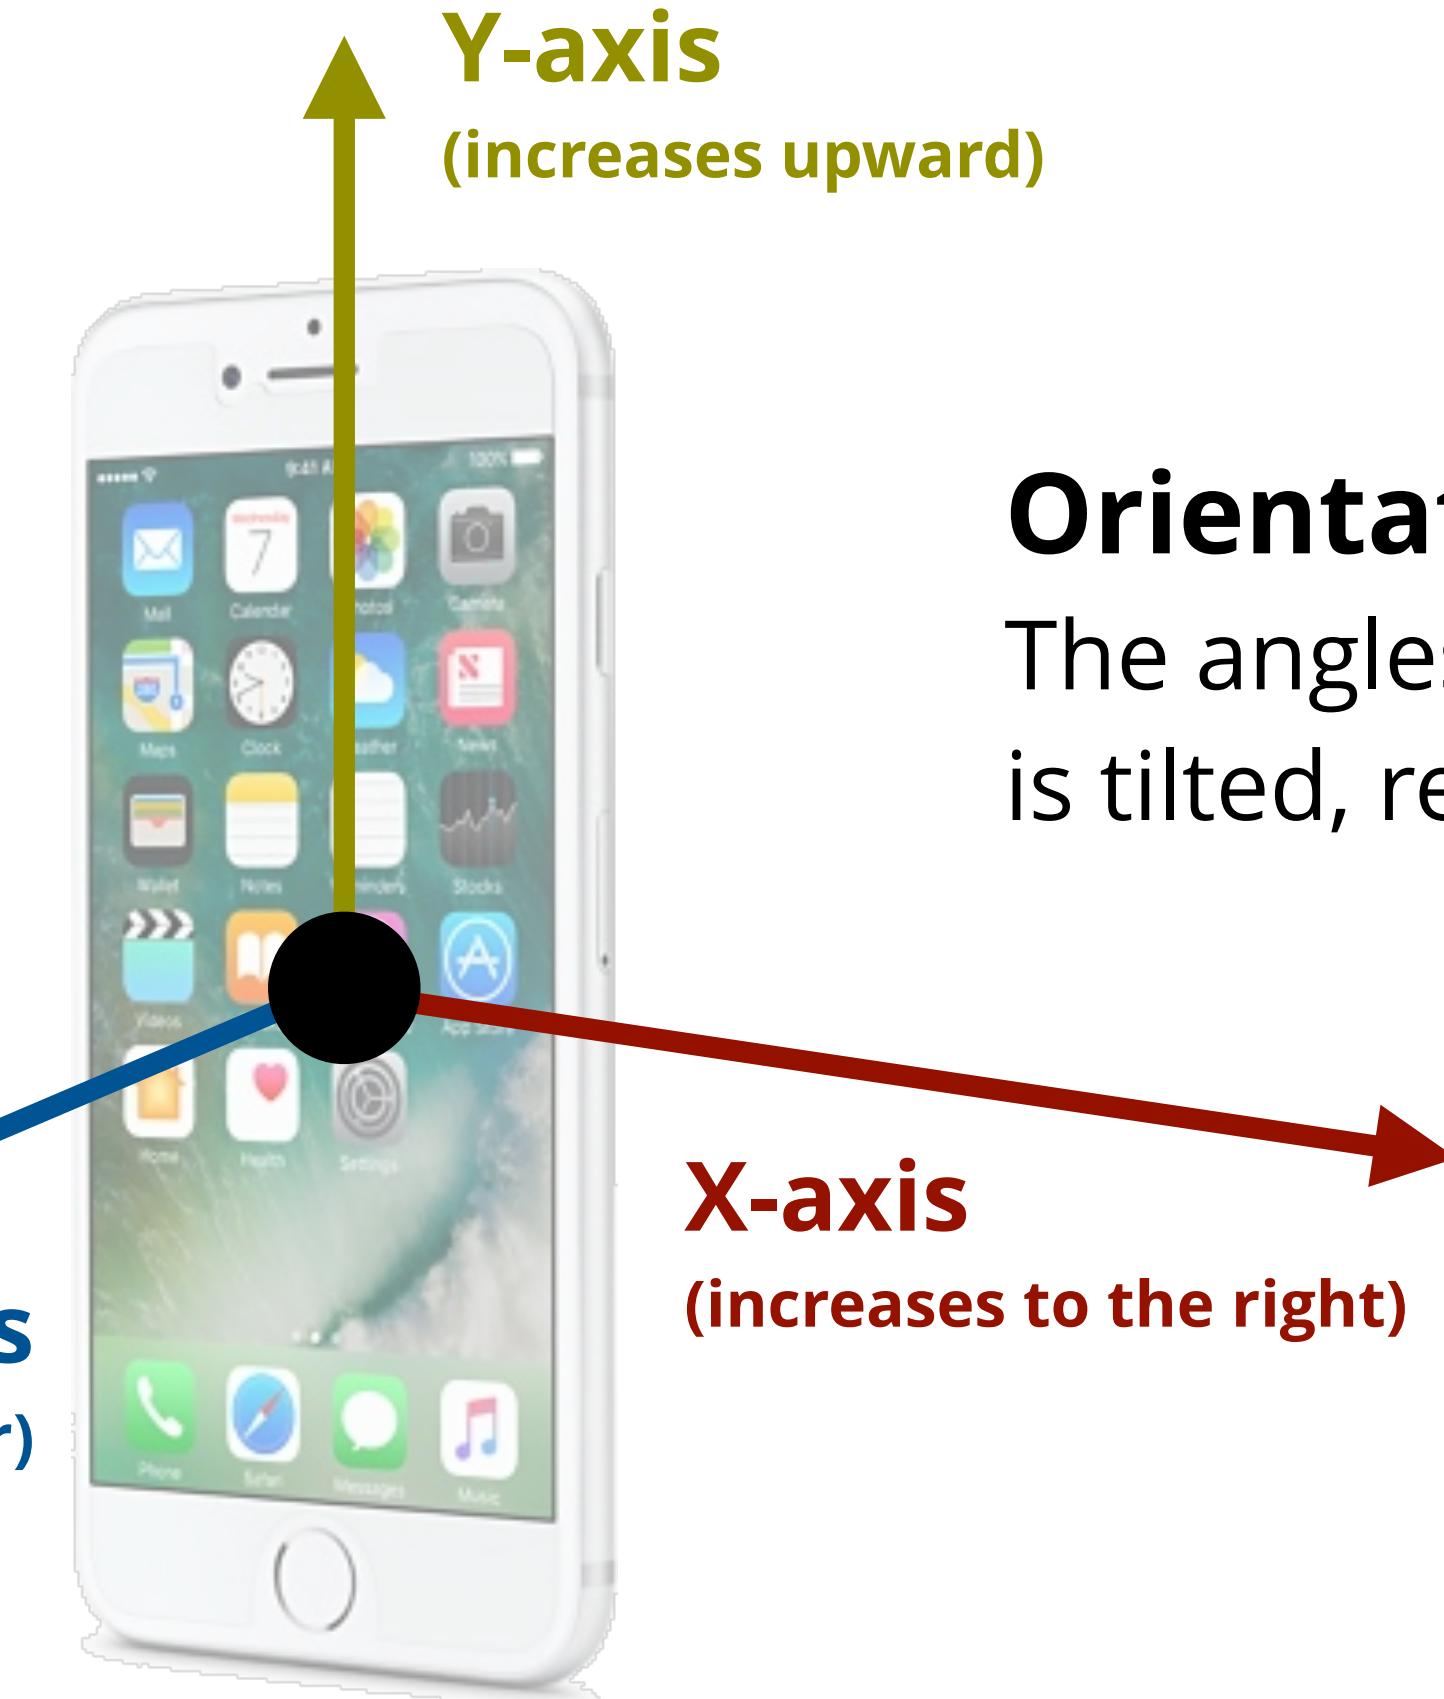
## Origin:

The device's position in space when the AR session begins.

## Position:

A position in space relative to the origin.

(increases towards the user)



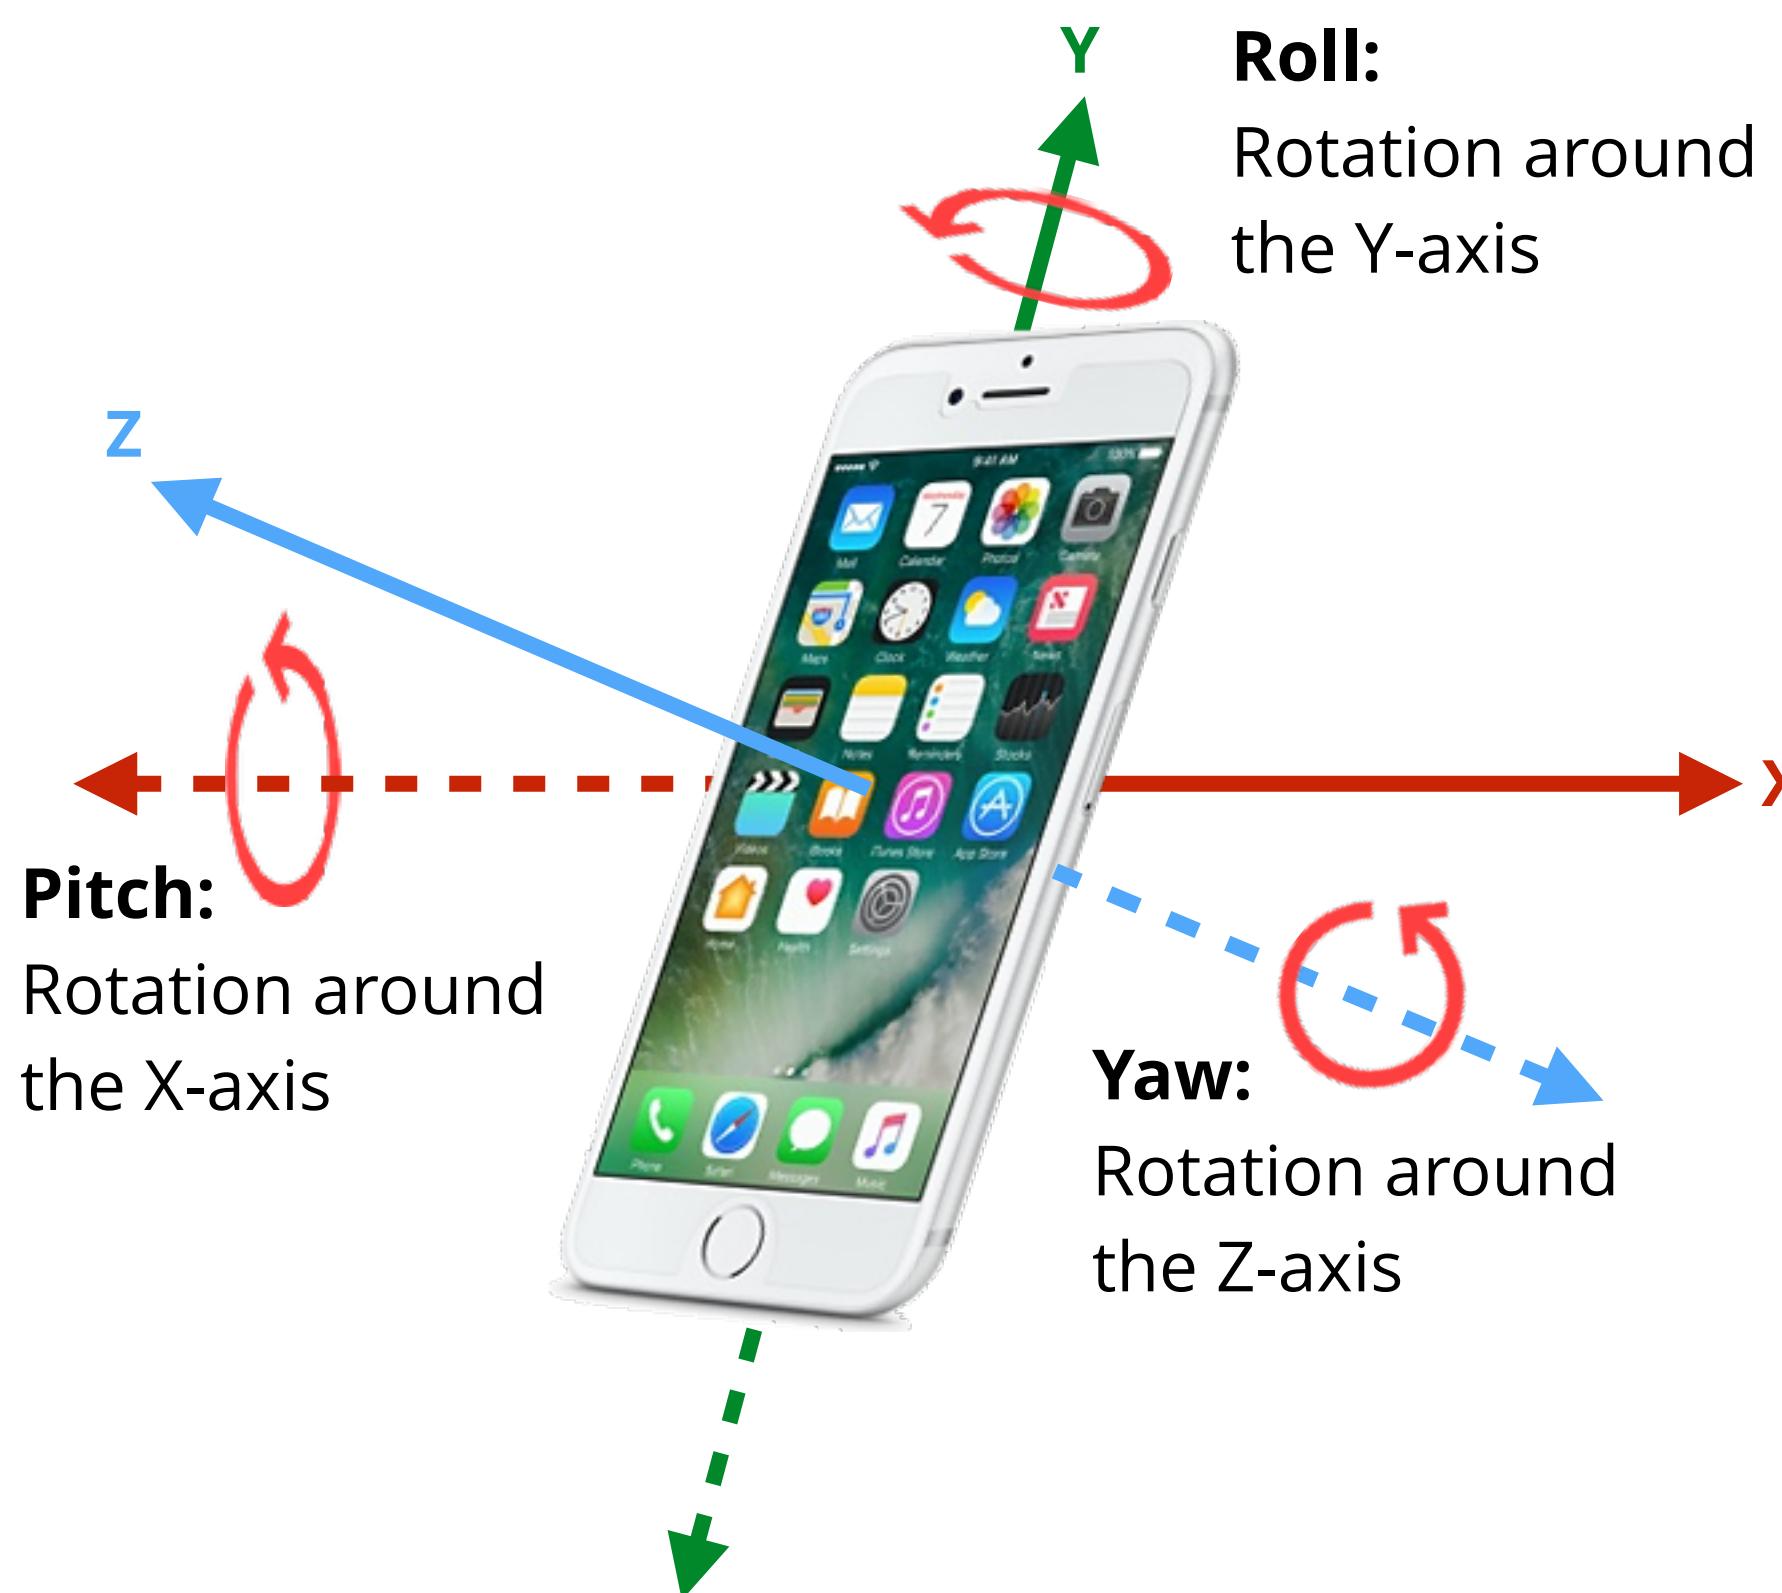
## Orientation:

The angles at which something is tilted, relative to x-, y-, and z-axes.

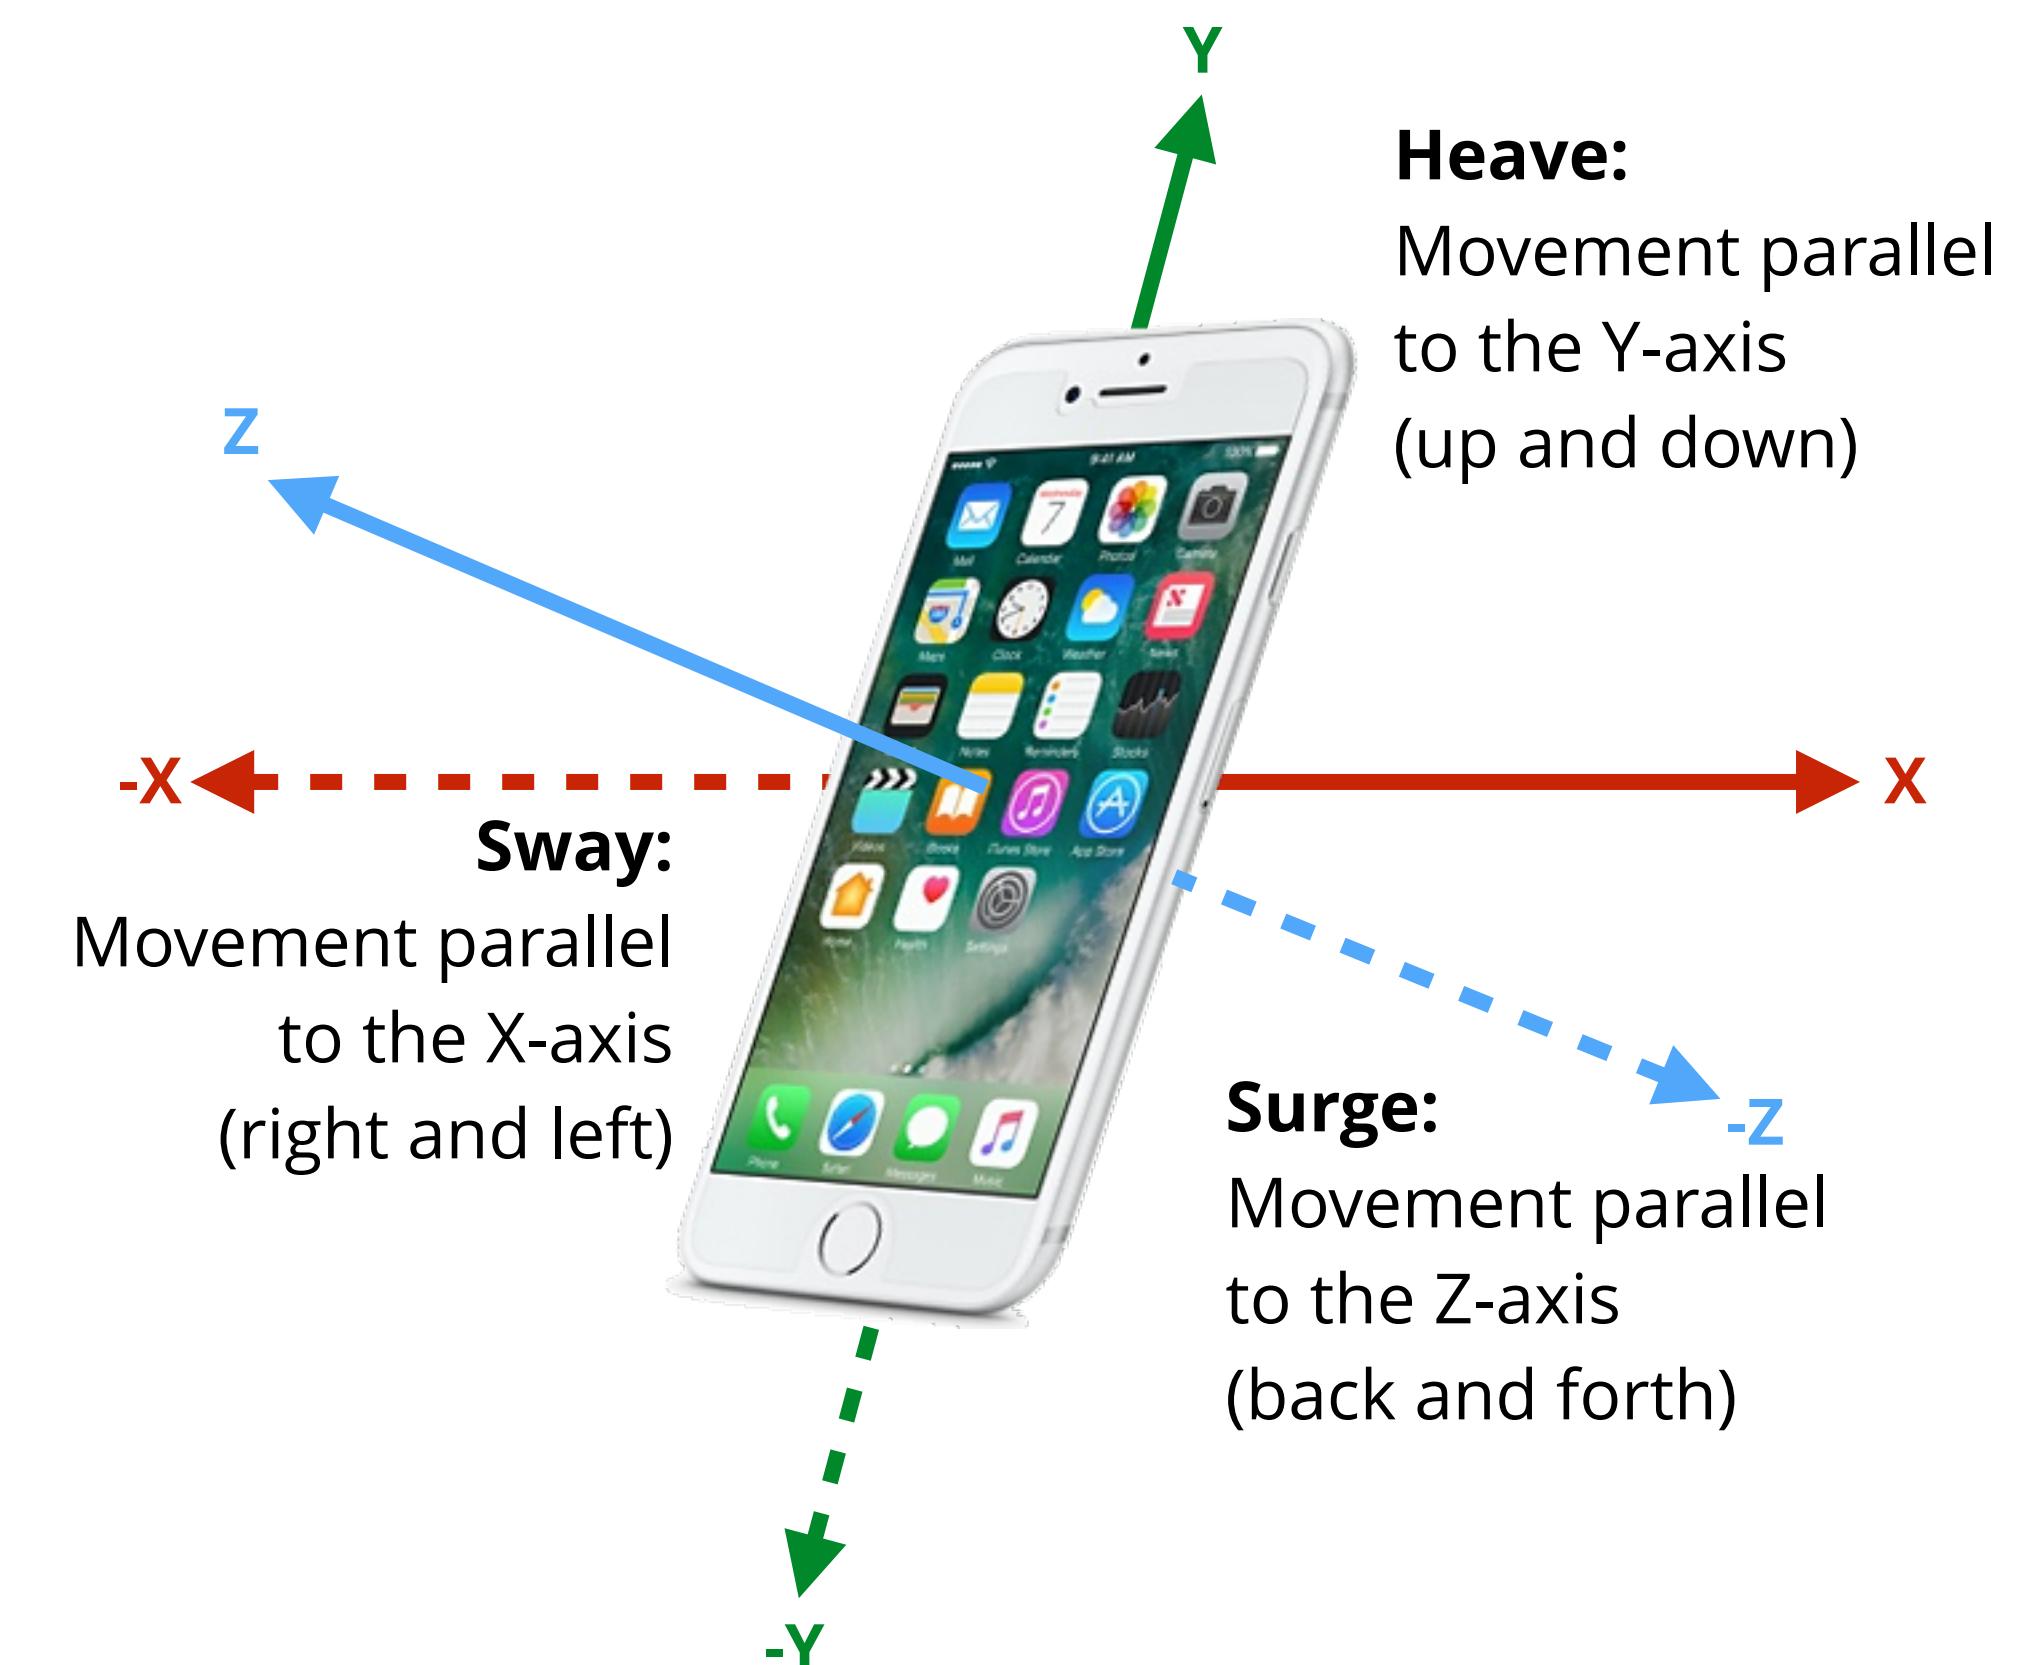


# CONFIGURATION: 3 OR 6 DEGREES OF FREEDOM?

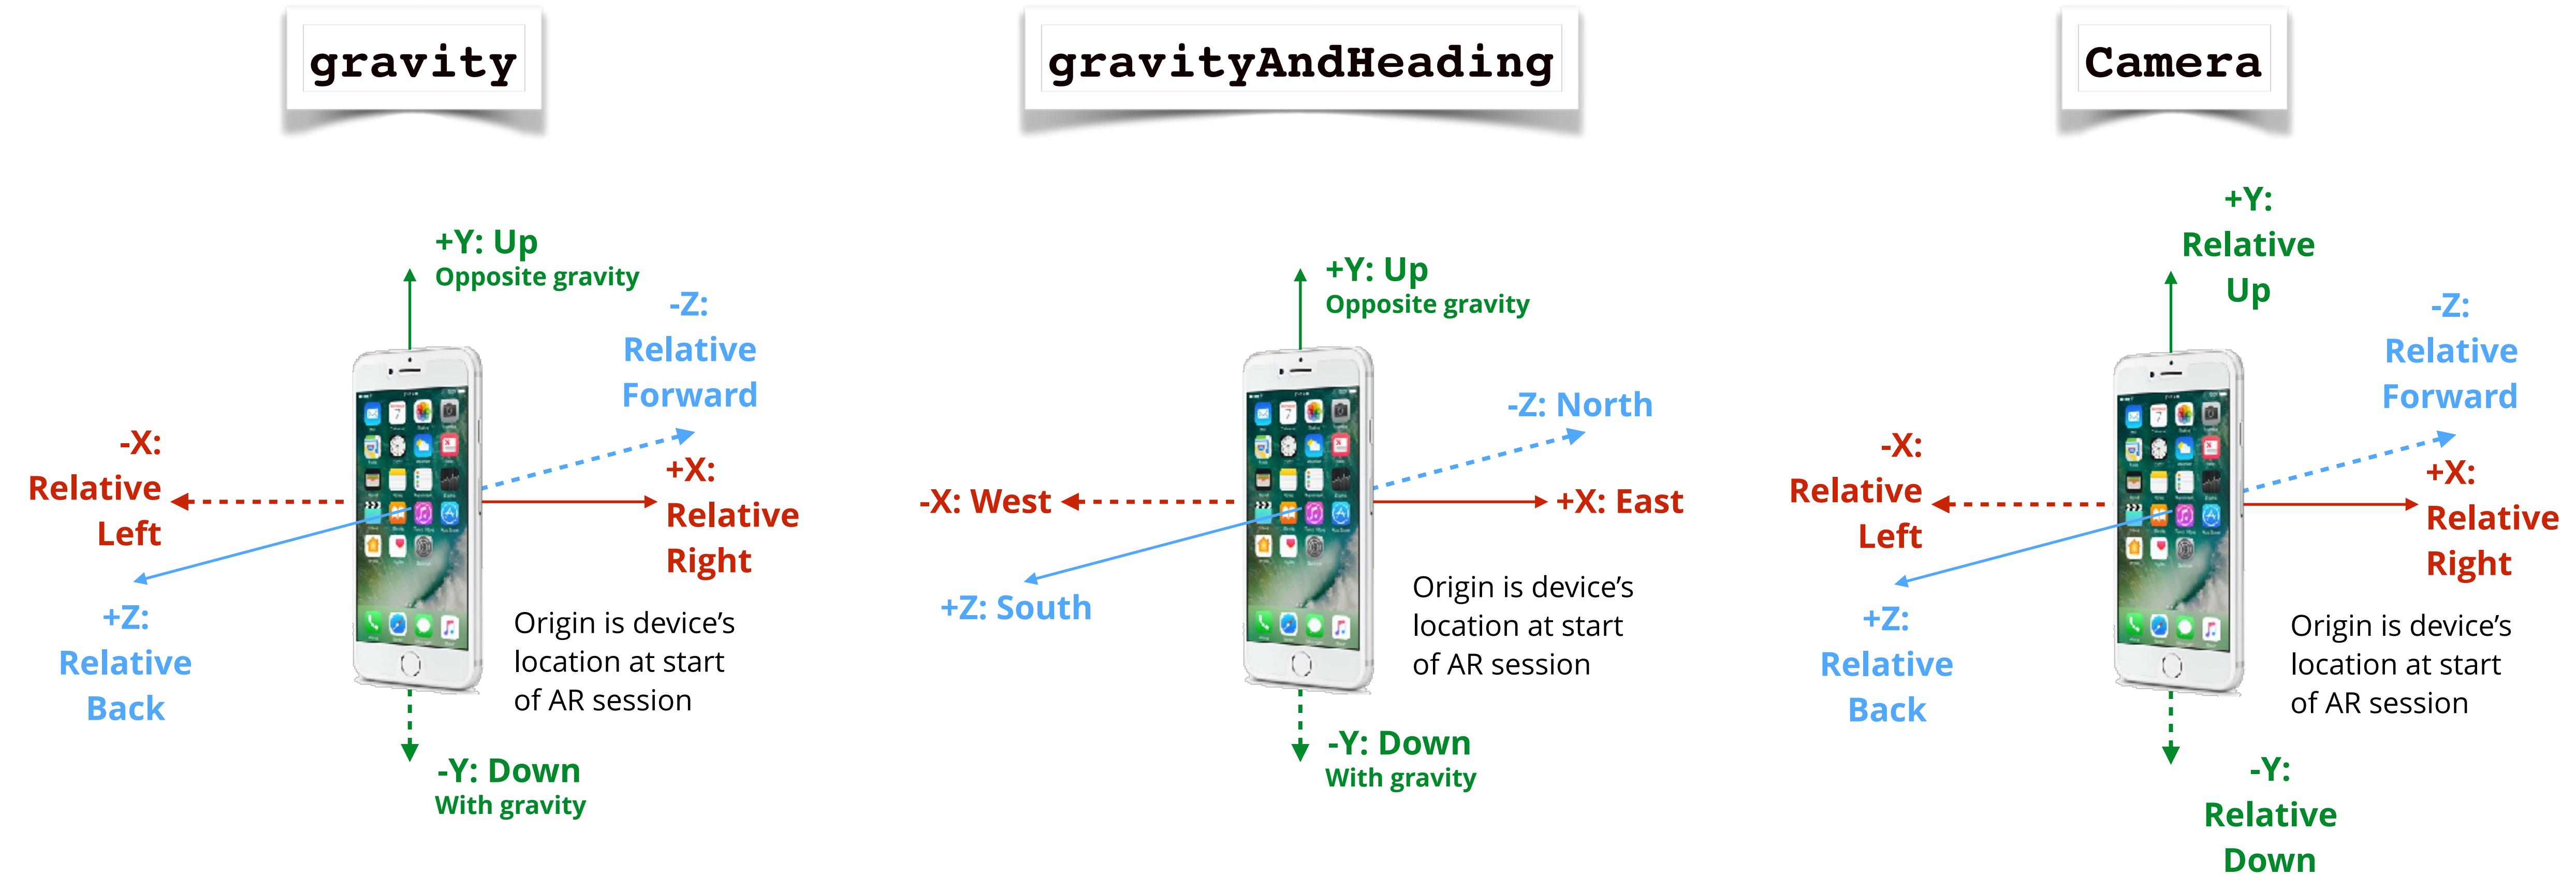
**AROrientationTrackingConfiguration**  
**ARWorldTrackingConfiguration**



**ARWorldTrackingConfiguration**

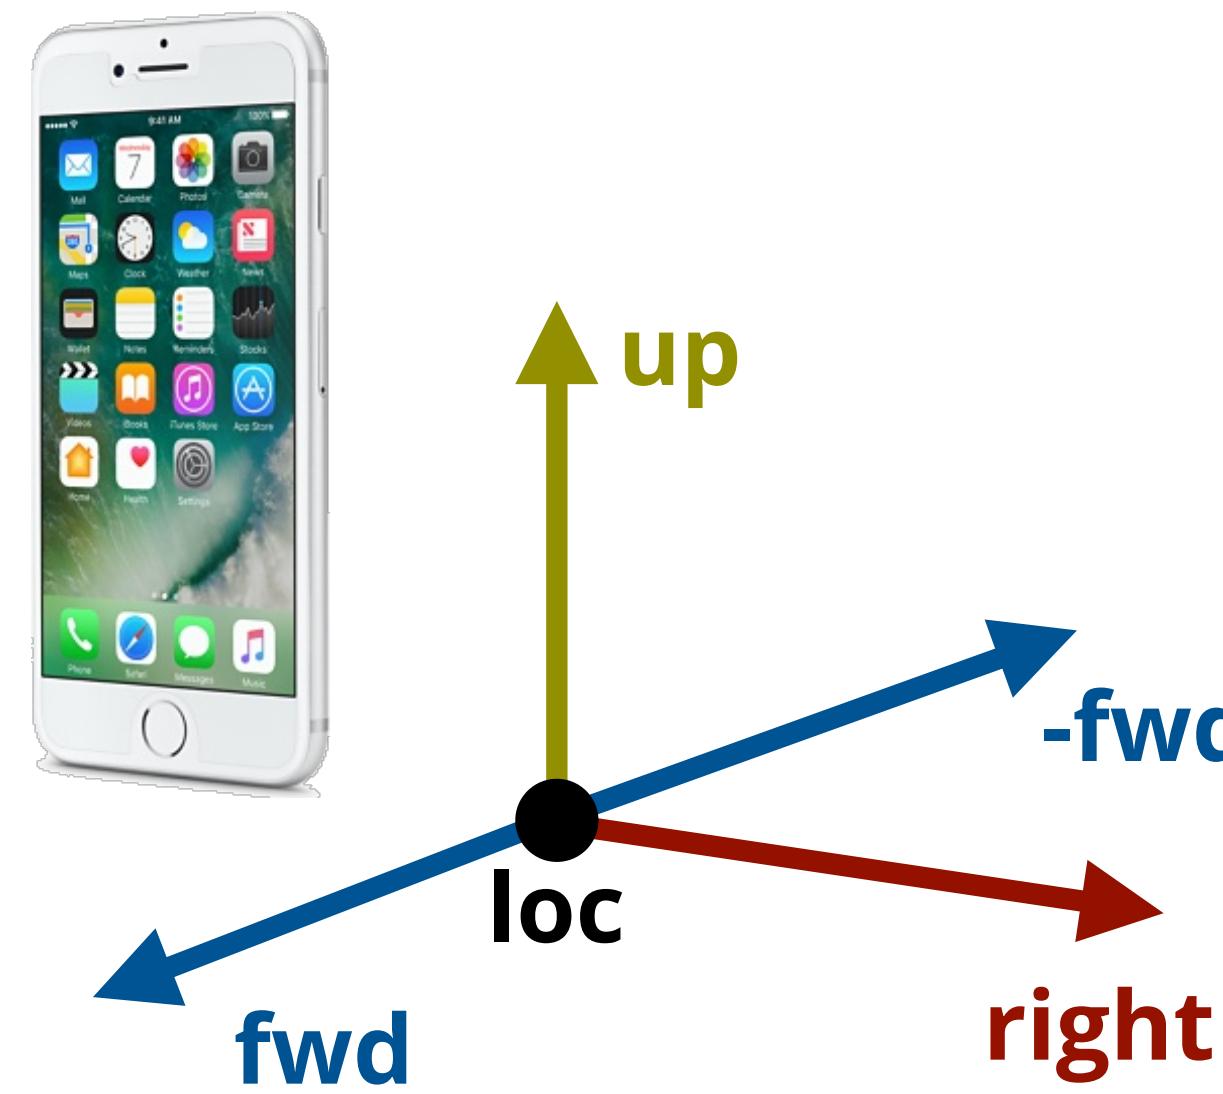


# ALIGNMENT: RELATING TO THE REAL WORLD



# THE SCENEVIEW'S TRANSFORMATION MATRIX

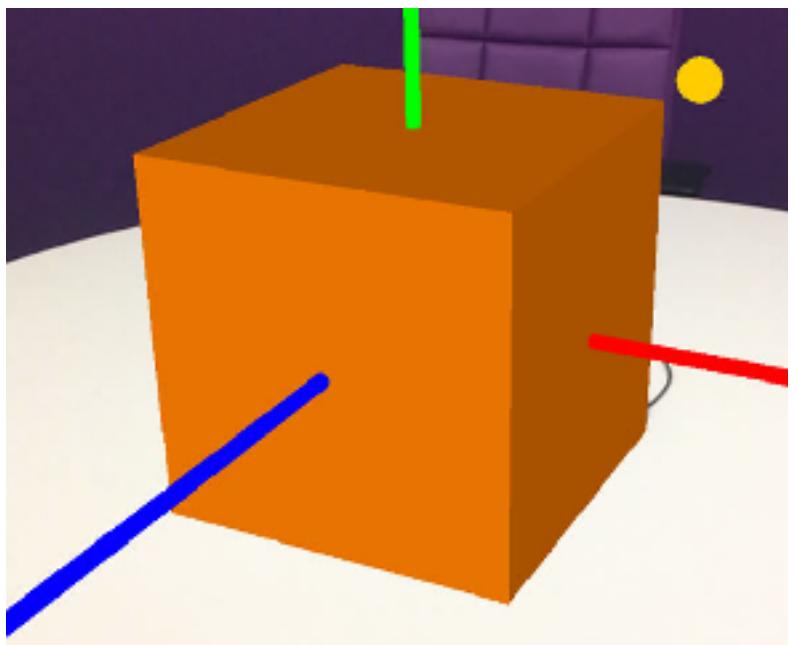
A 4-by-4 matrix representing the scene's *location* and *orientation*:



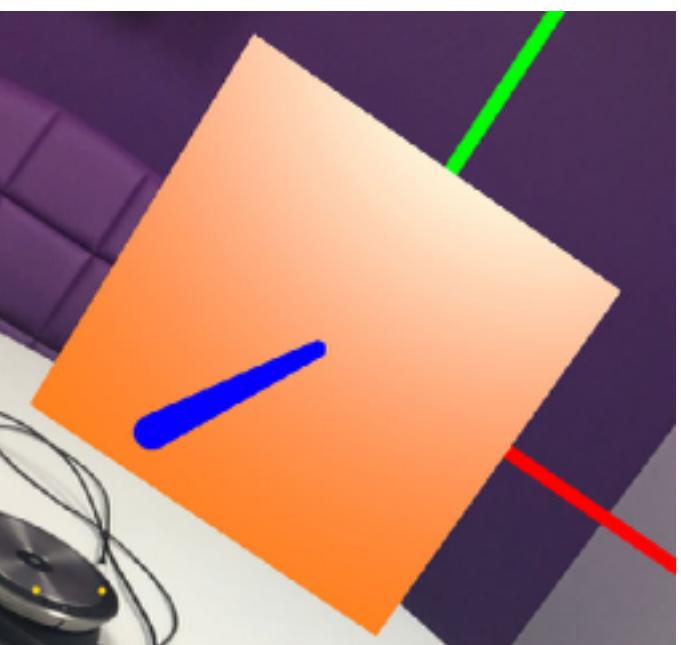
		<b>Orientation:</b> Which way the camera's facing		<b>Location:</b> Where the camera is	
	$\text{right}_x$	$\text{up}_x$	$-\text{fwd}_x$	$\text{loc}_x$	
	$\text{right}_y$	$\text{up}_y$	$-\text{fwd}_y$	$\text{loc}_y$	
	$\text{right}_z$	$\text{up}_z$	$-\text{fwd}_z$	$\text{loc}_z$	
0		0	0	1	

# DIFFUSE VS. SPECULAR REFLECTION

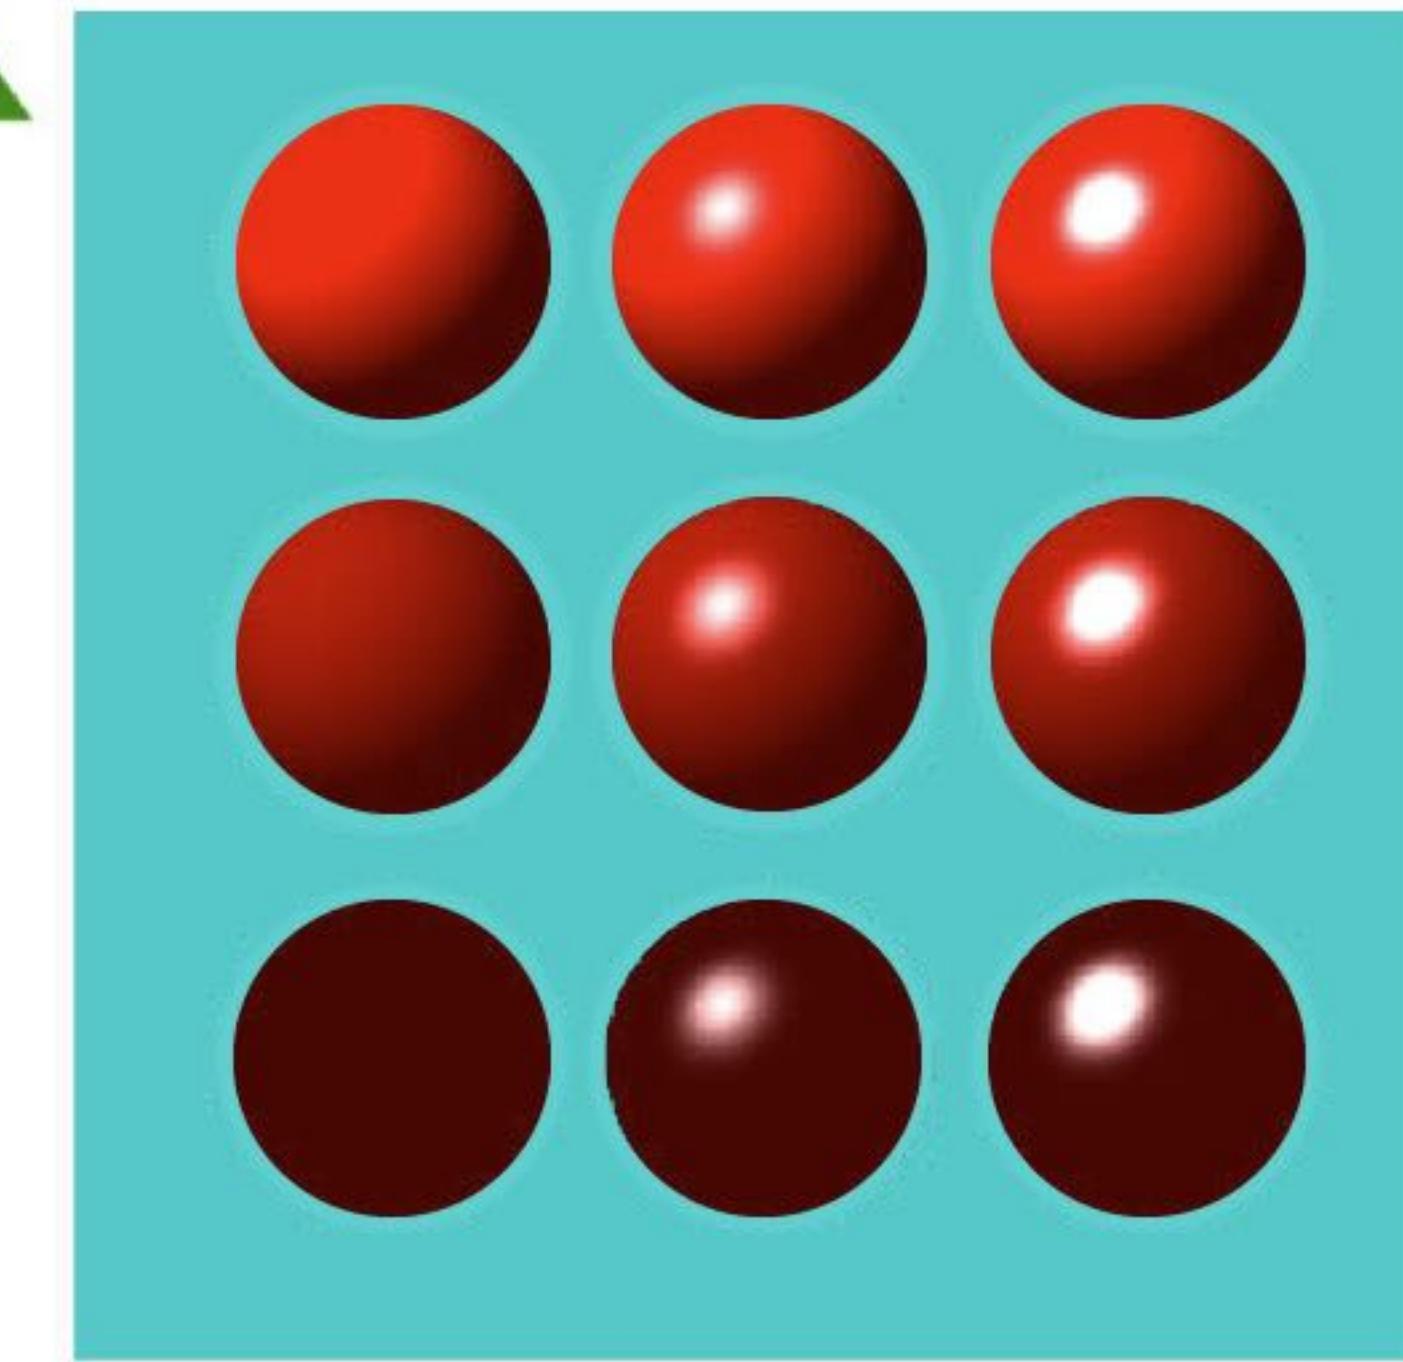
It might be easier to *show* you what these are than to *tell* you about them.



Diffuse reflection



Specular reflection



More specular reflection

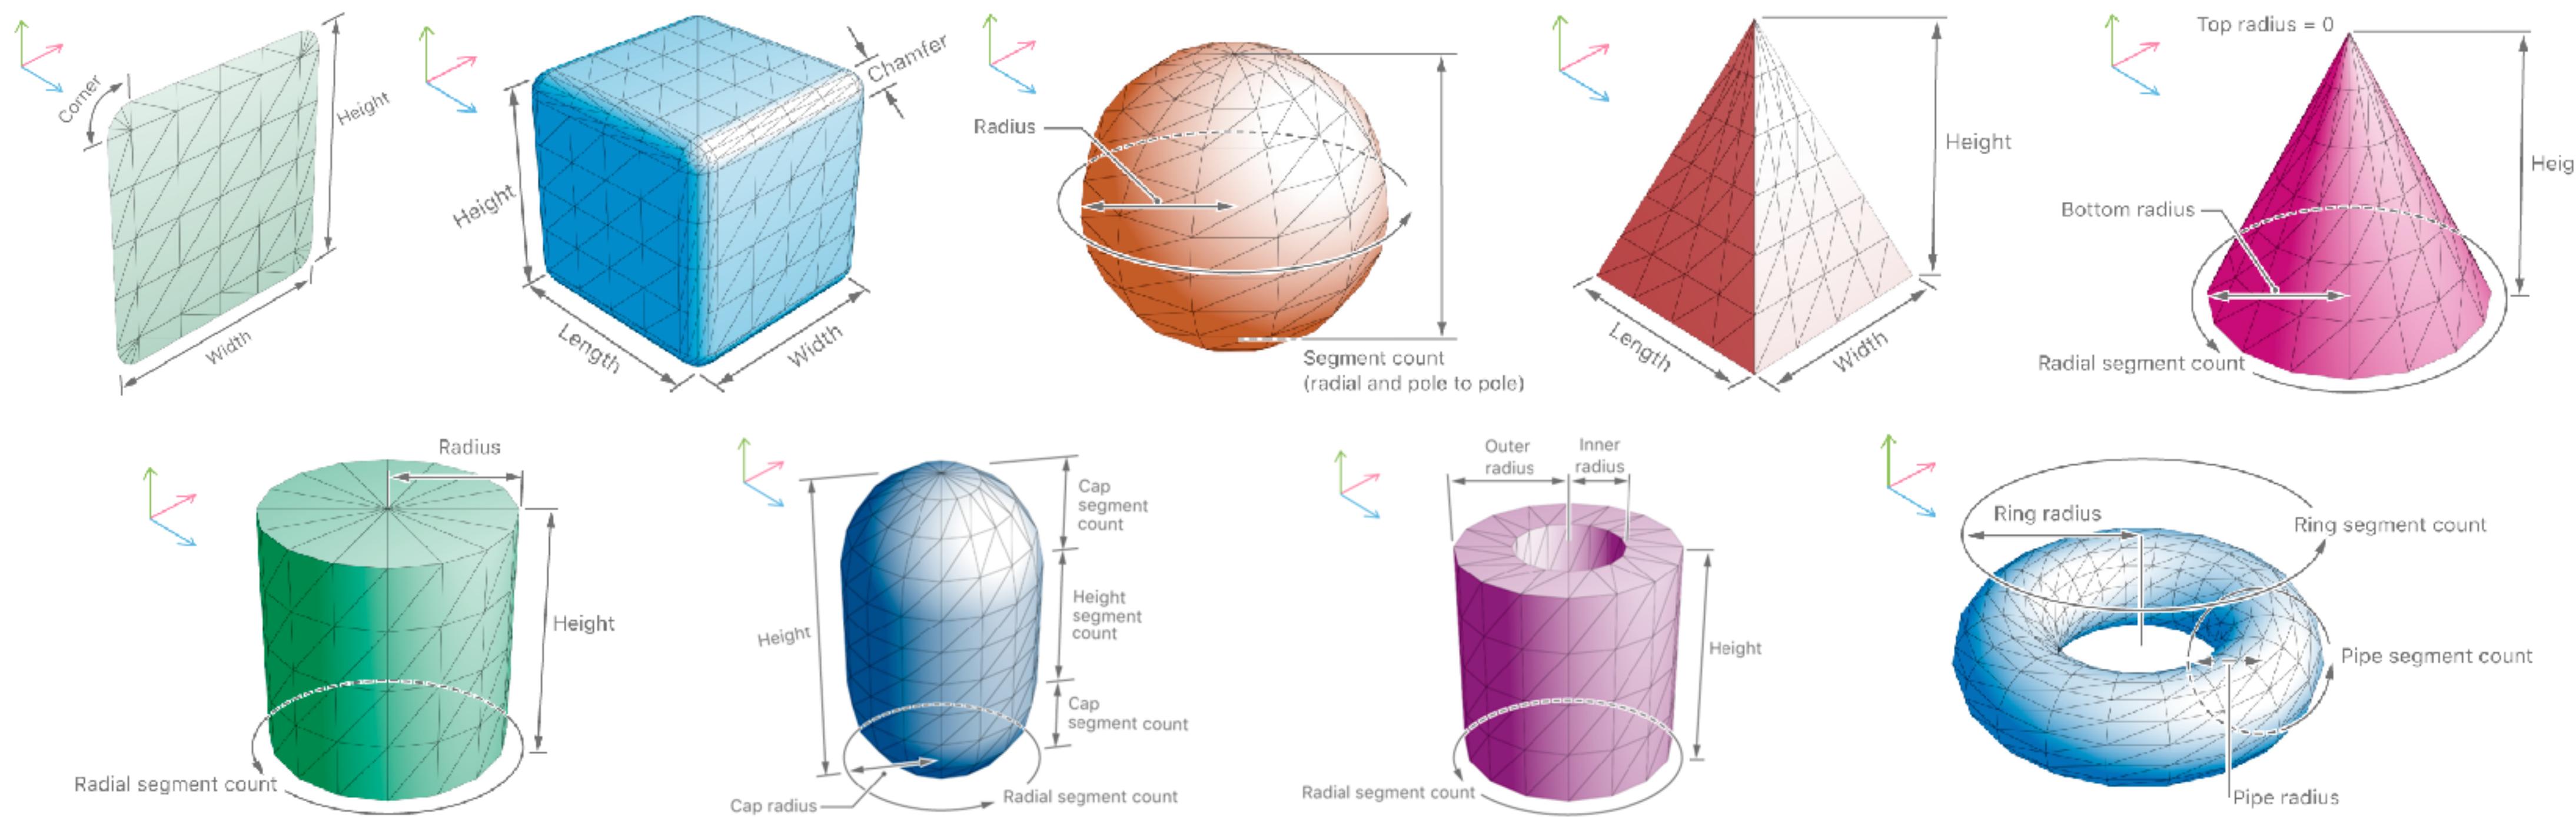


# DEMO 1: HAPPY AR PAINTER



# How HAPPY AR PAINTER WORKS, PART 1

## 1. Define a SceneKit geometry:



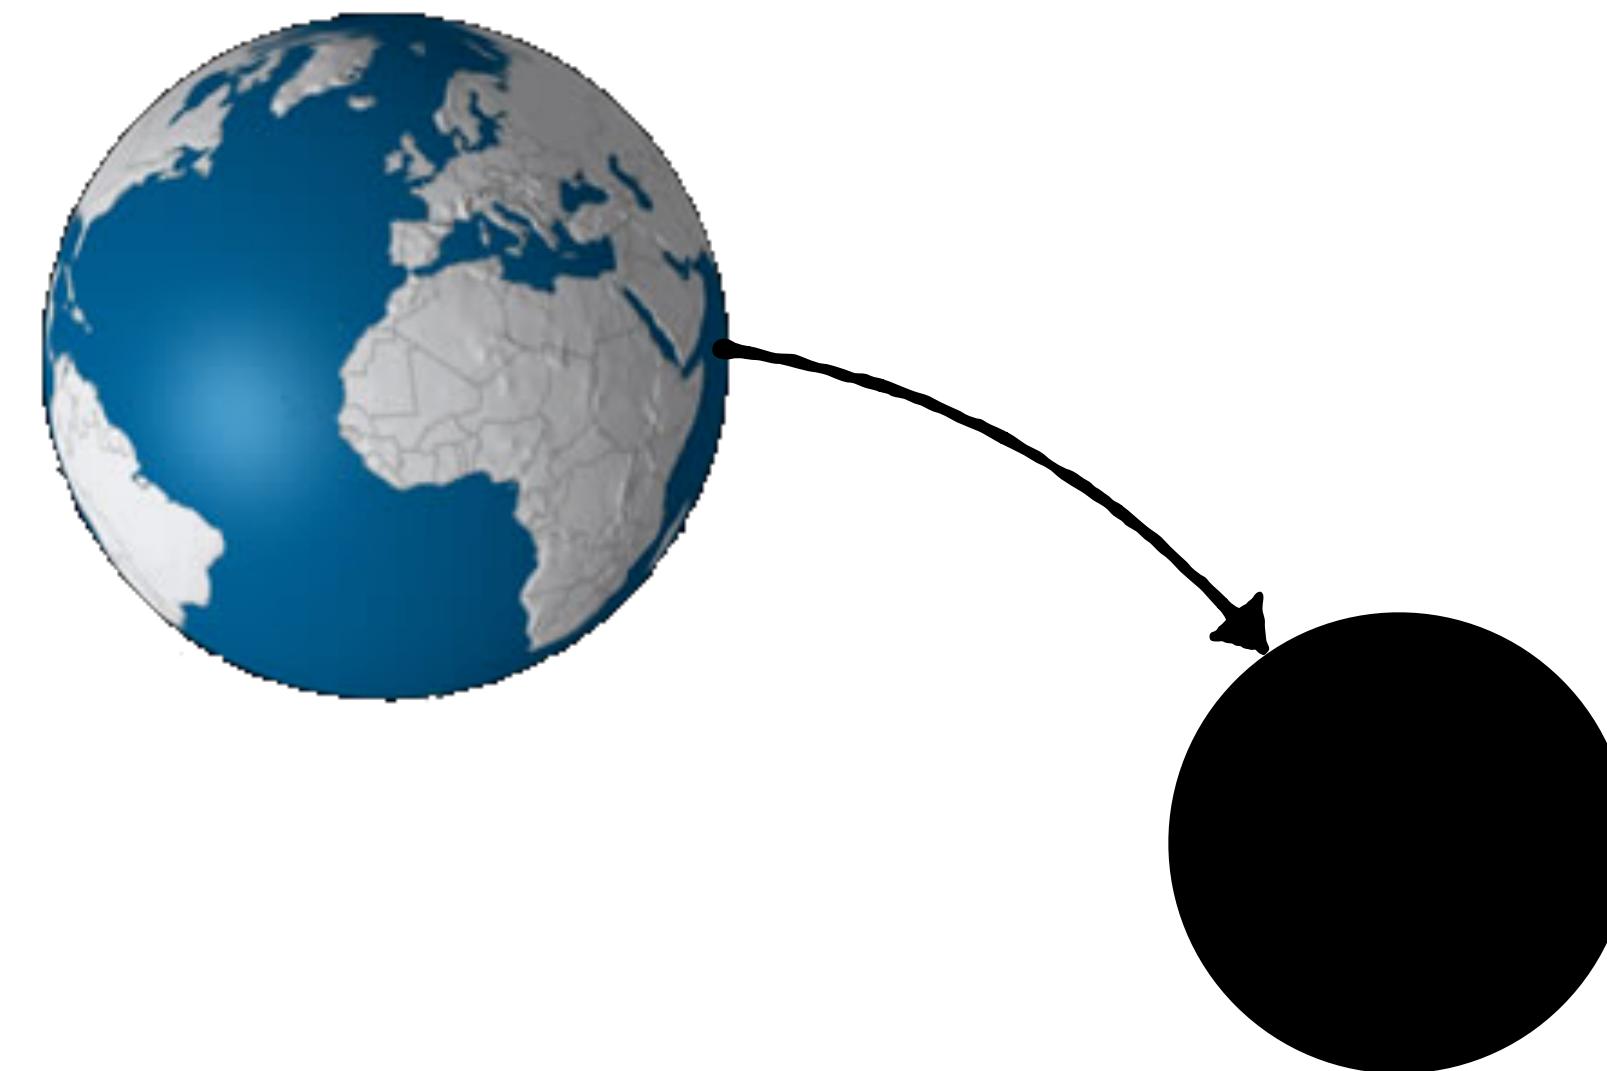
# How HAPPY AR PAINTER WORKS, PART 2

---

2. Apply reflective properties  
to the geometry.

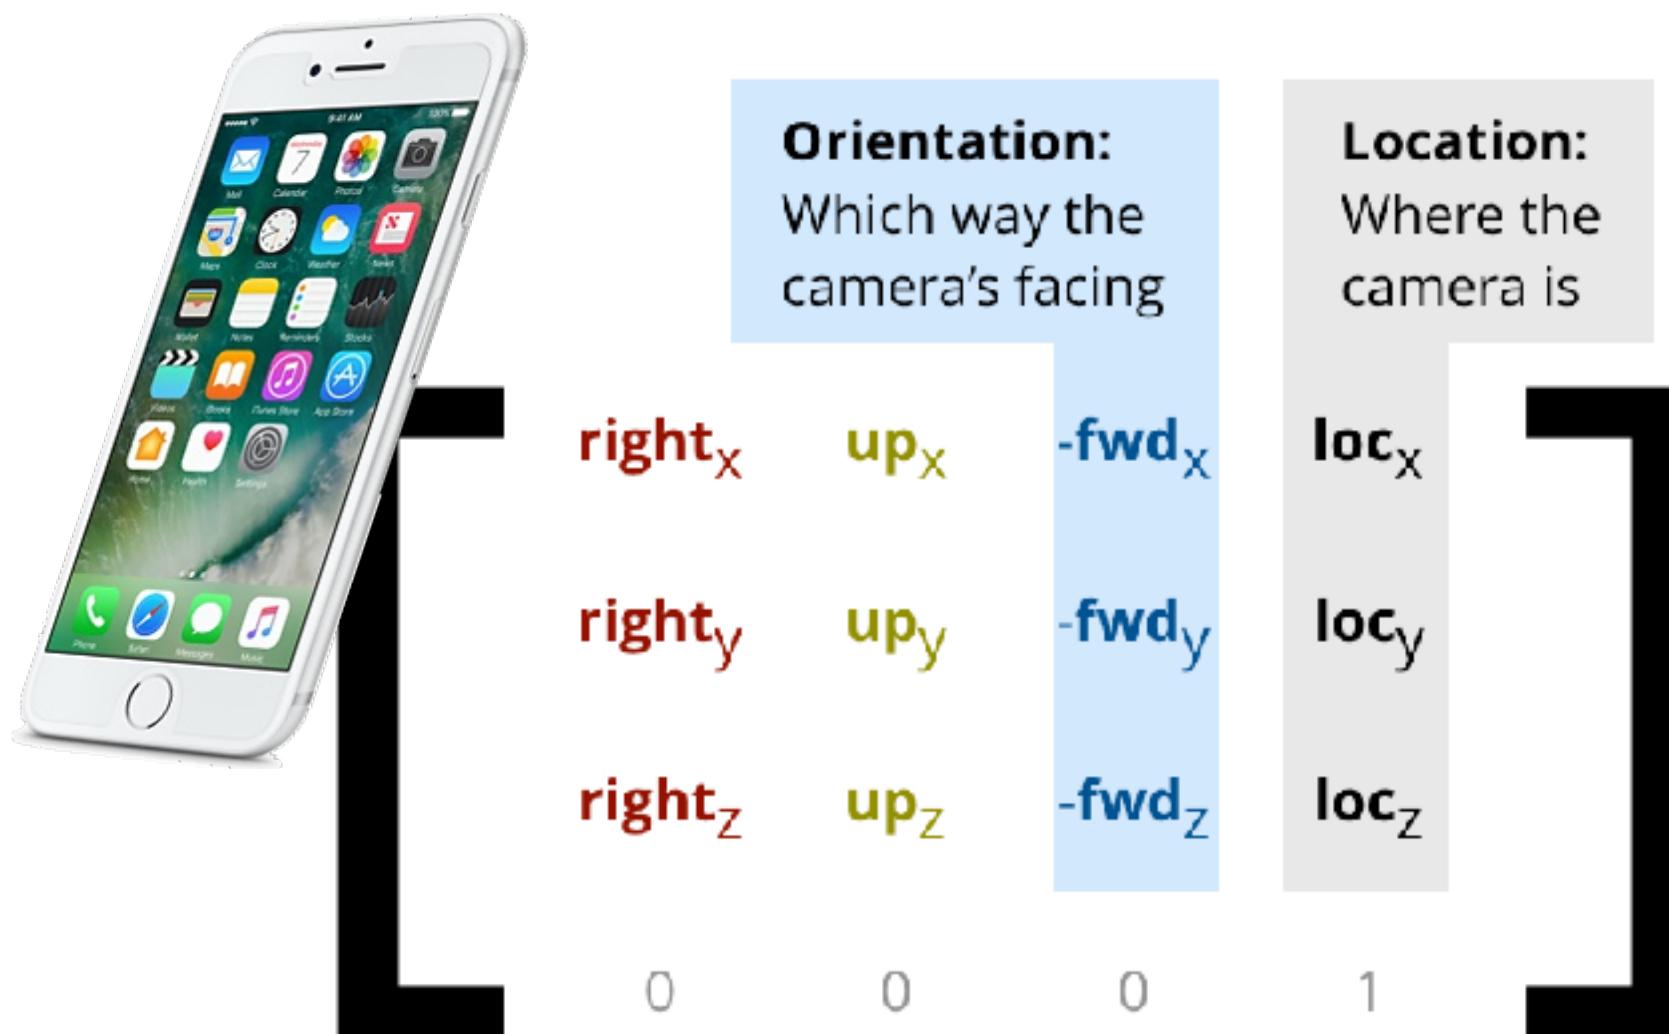


3. Assign the geometry  
to a node.

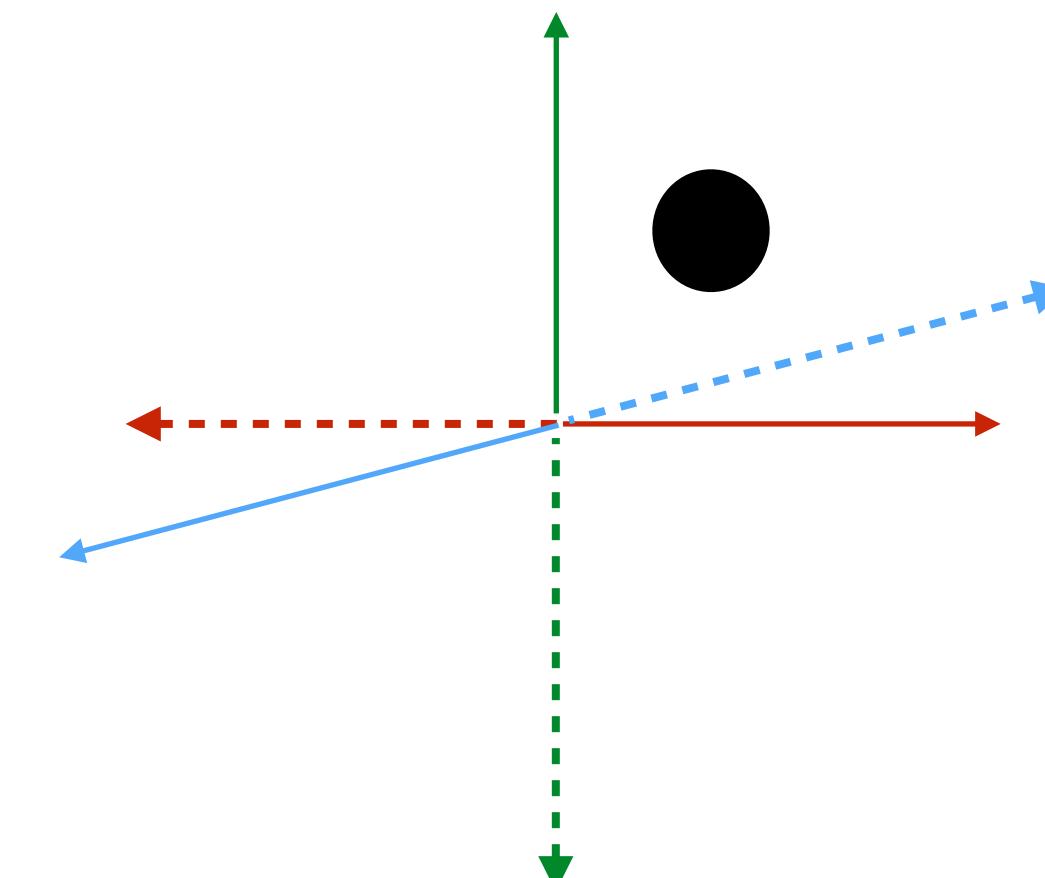


# How Happy AR Painter Works, Part 3

4. Get the device's orientation and position from the SceneView's transform matrix.



5. Set the node's orientation and position to that of the device's orientation and position.

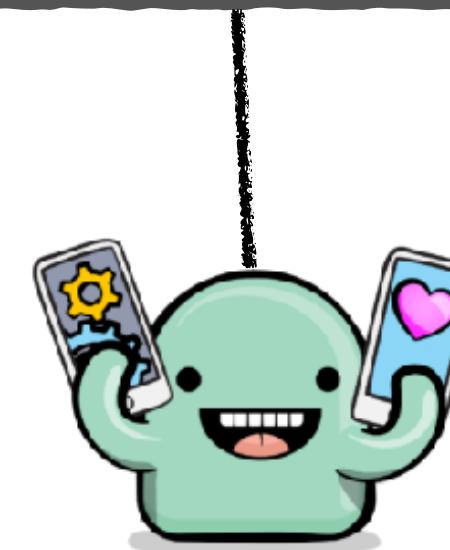


# HOW HAPPY AR PAINTER WORKS, PART 4

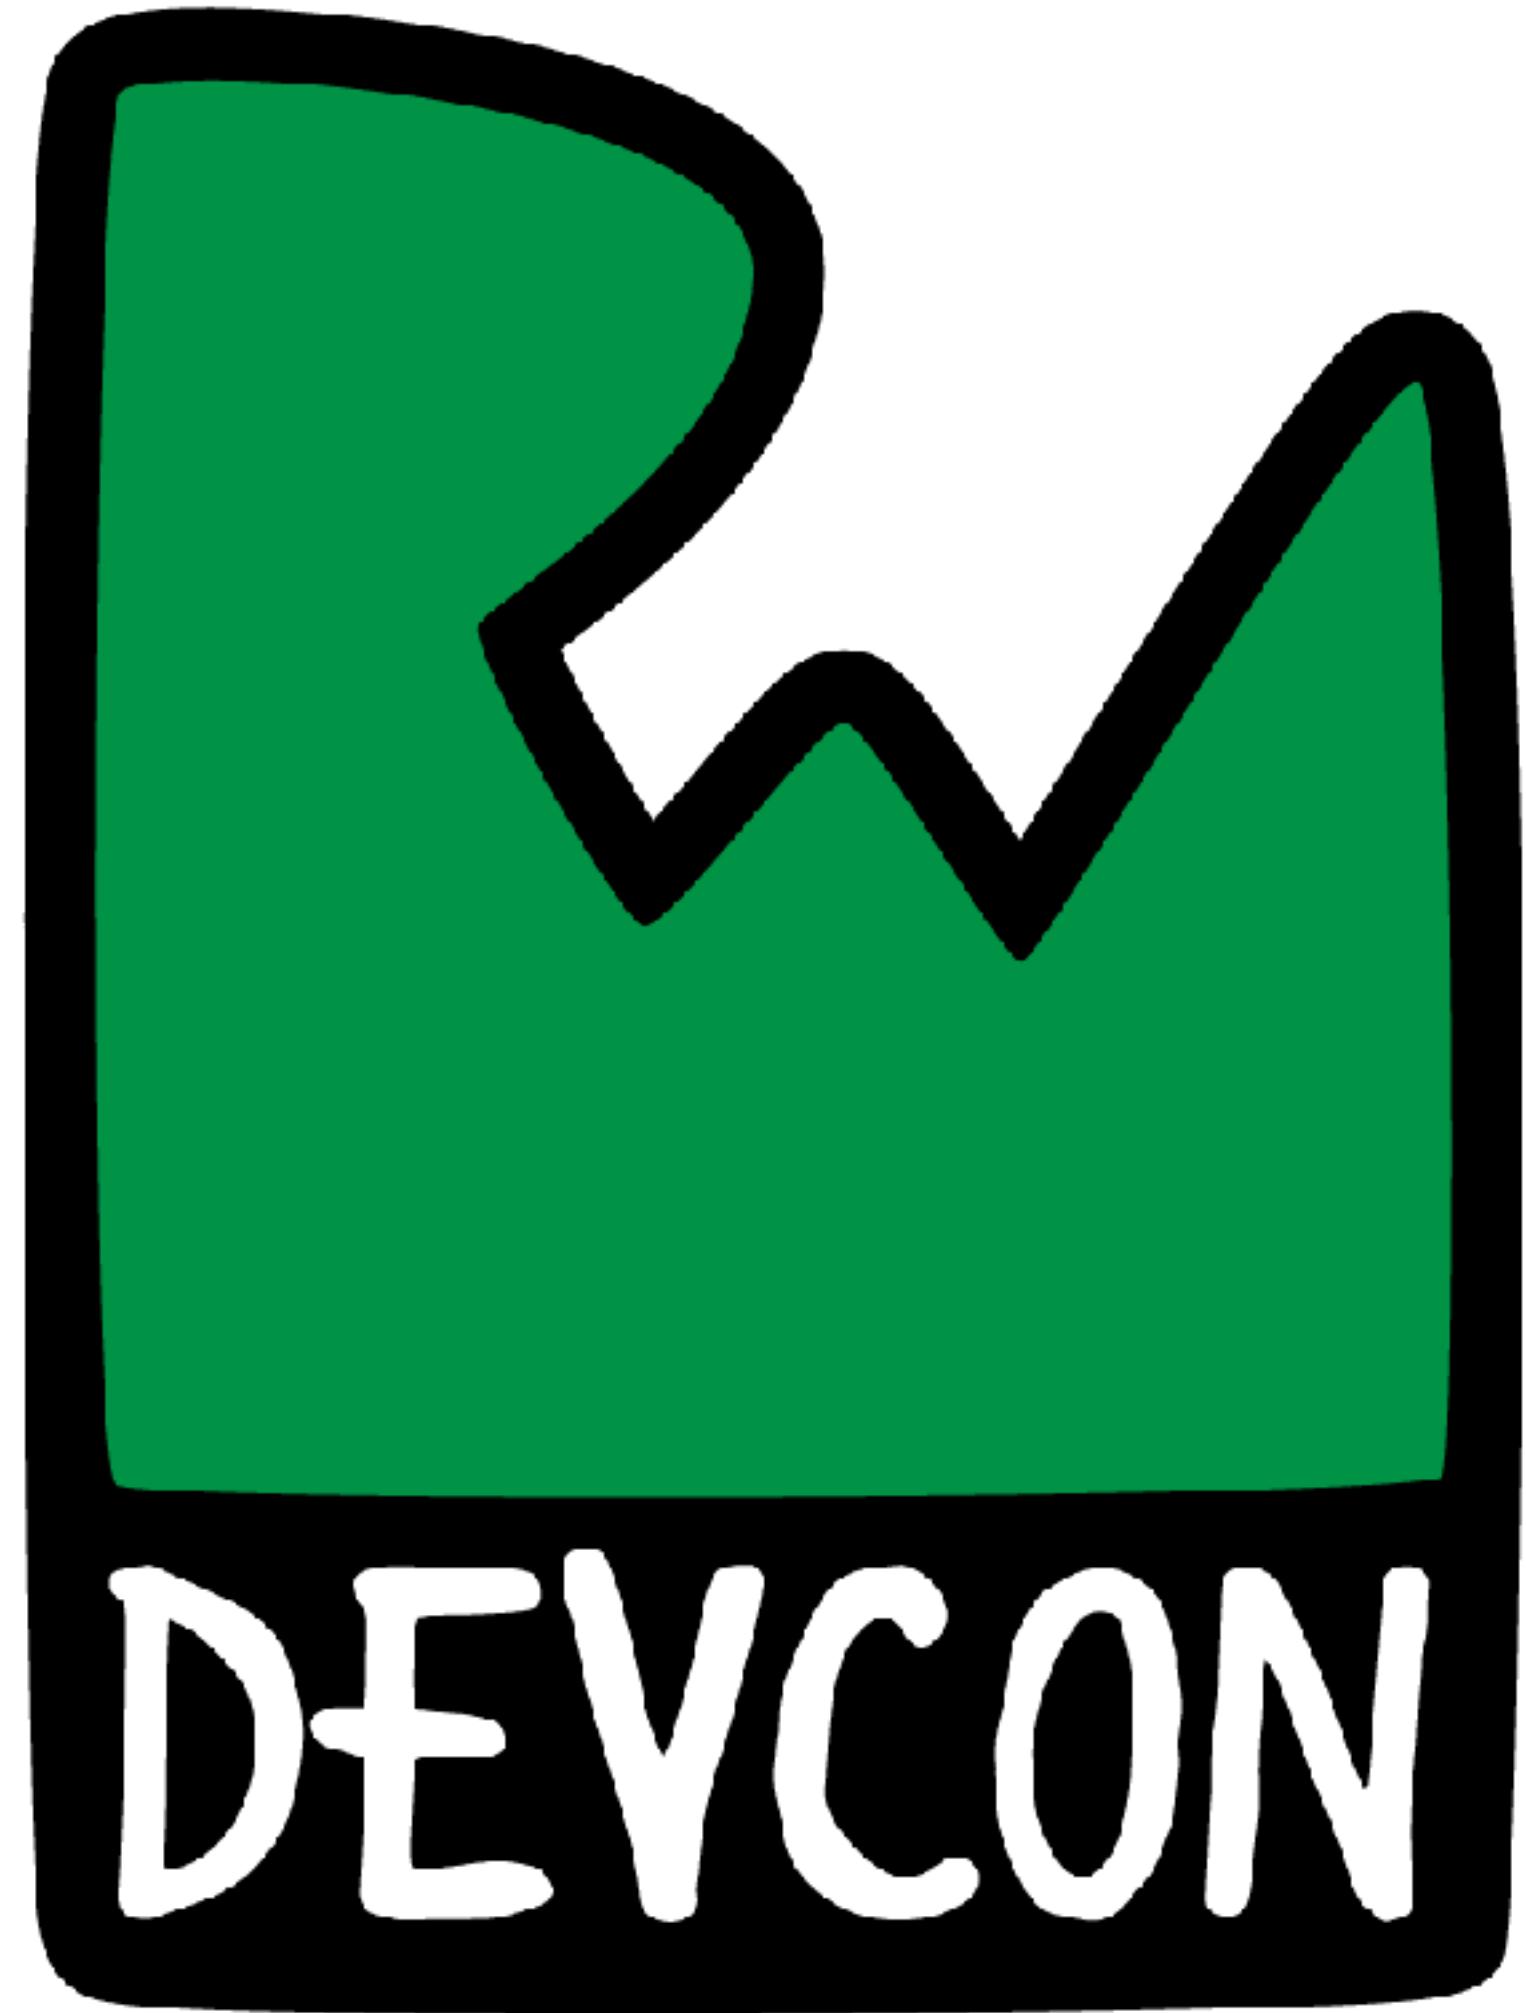
6. Add the node to the scene.



The  
root  
node



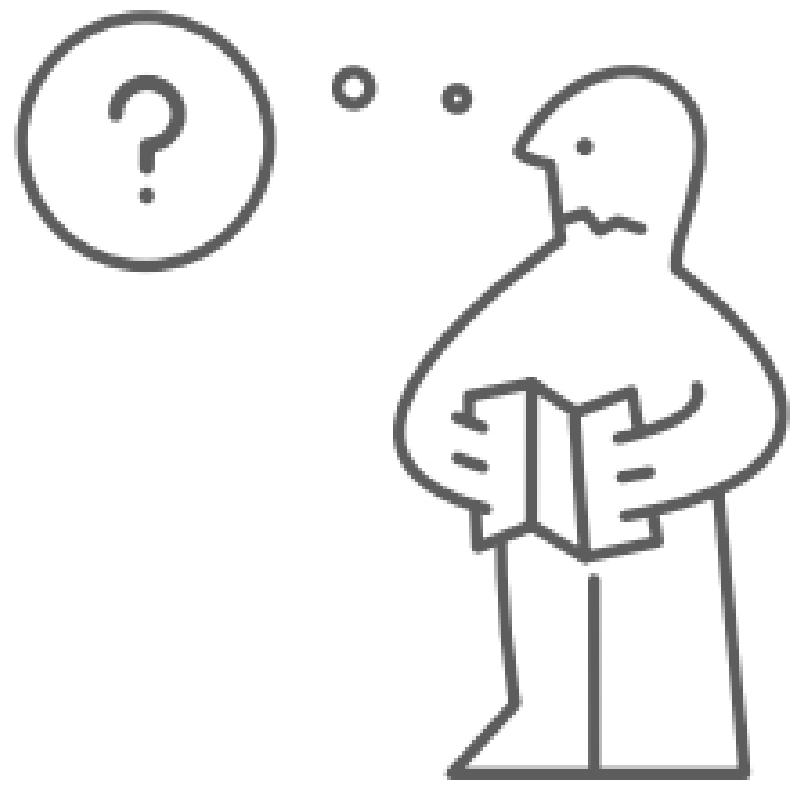
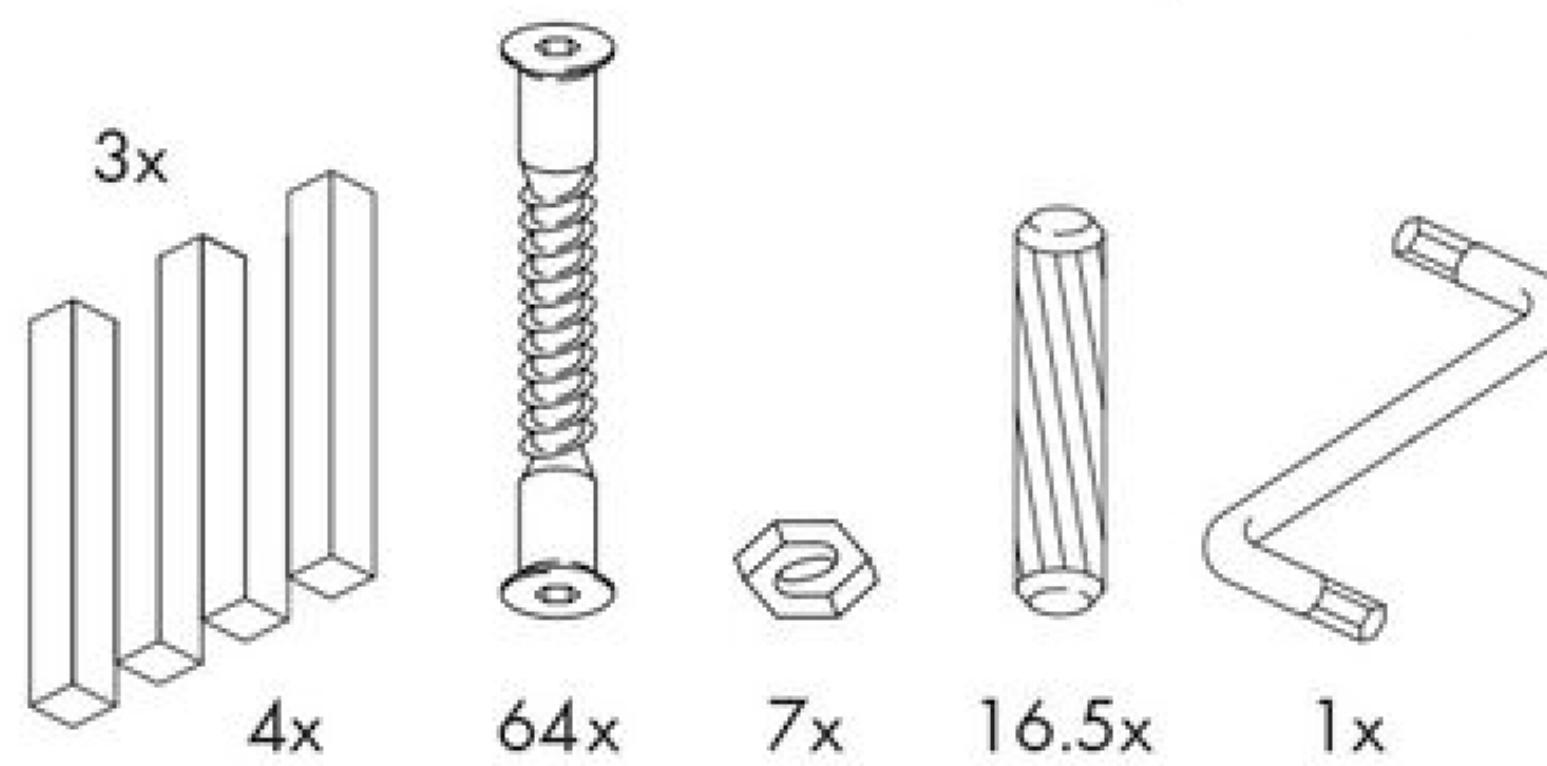
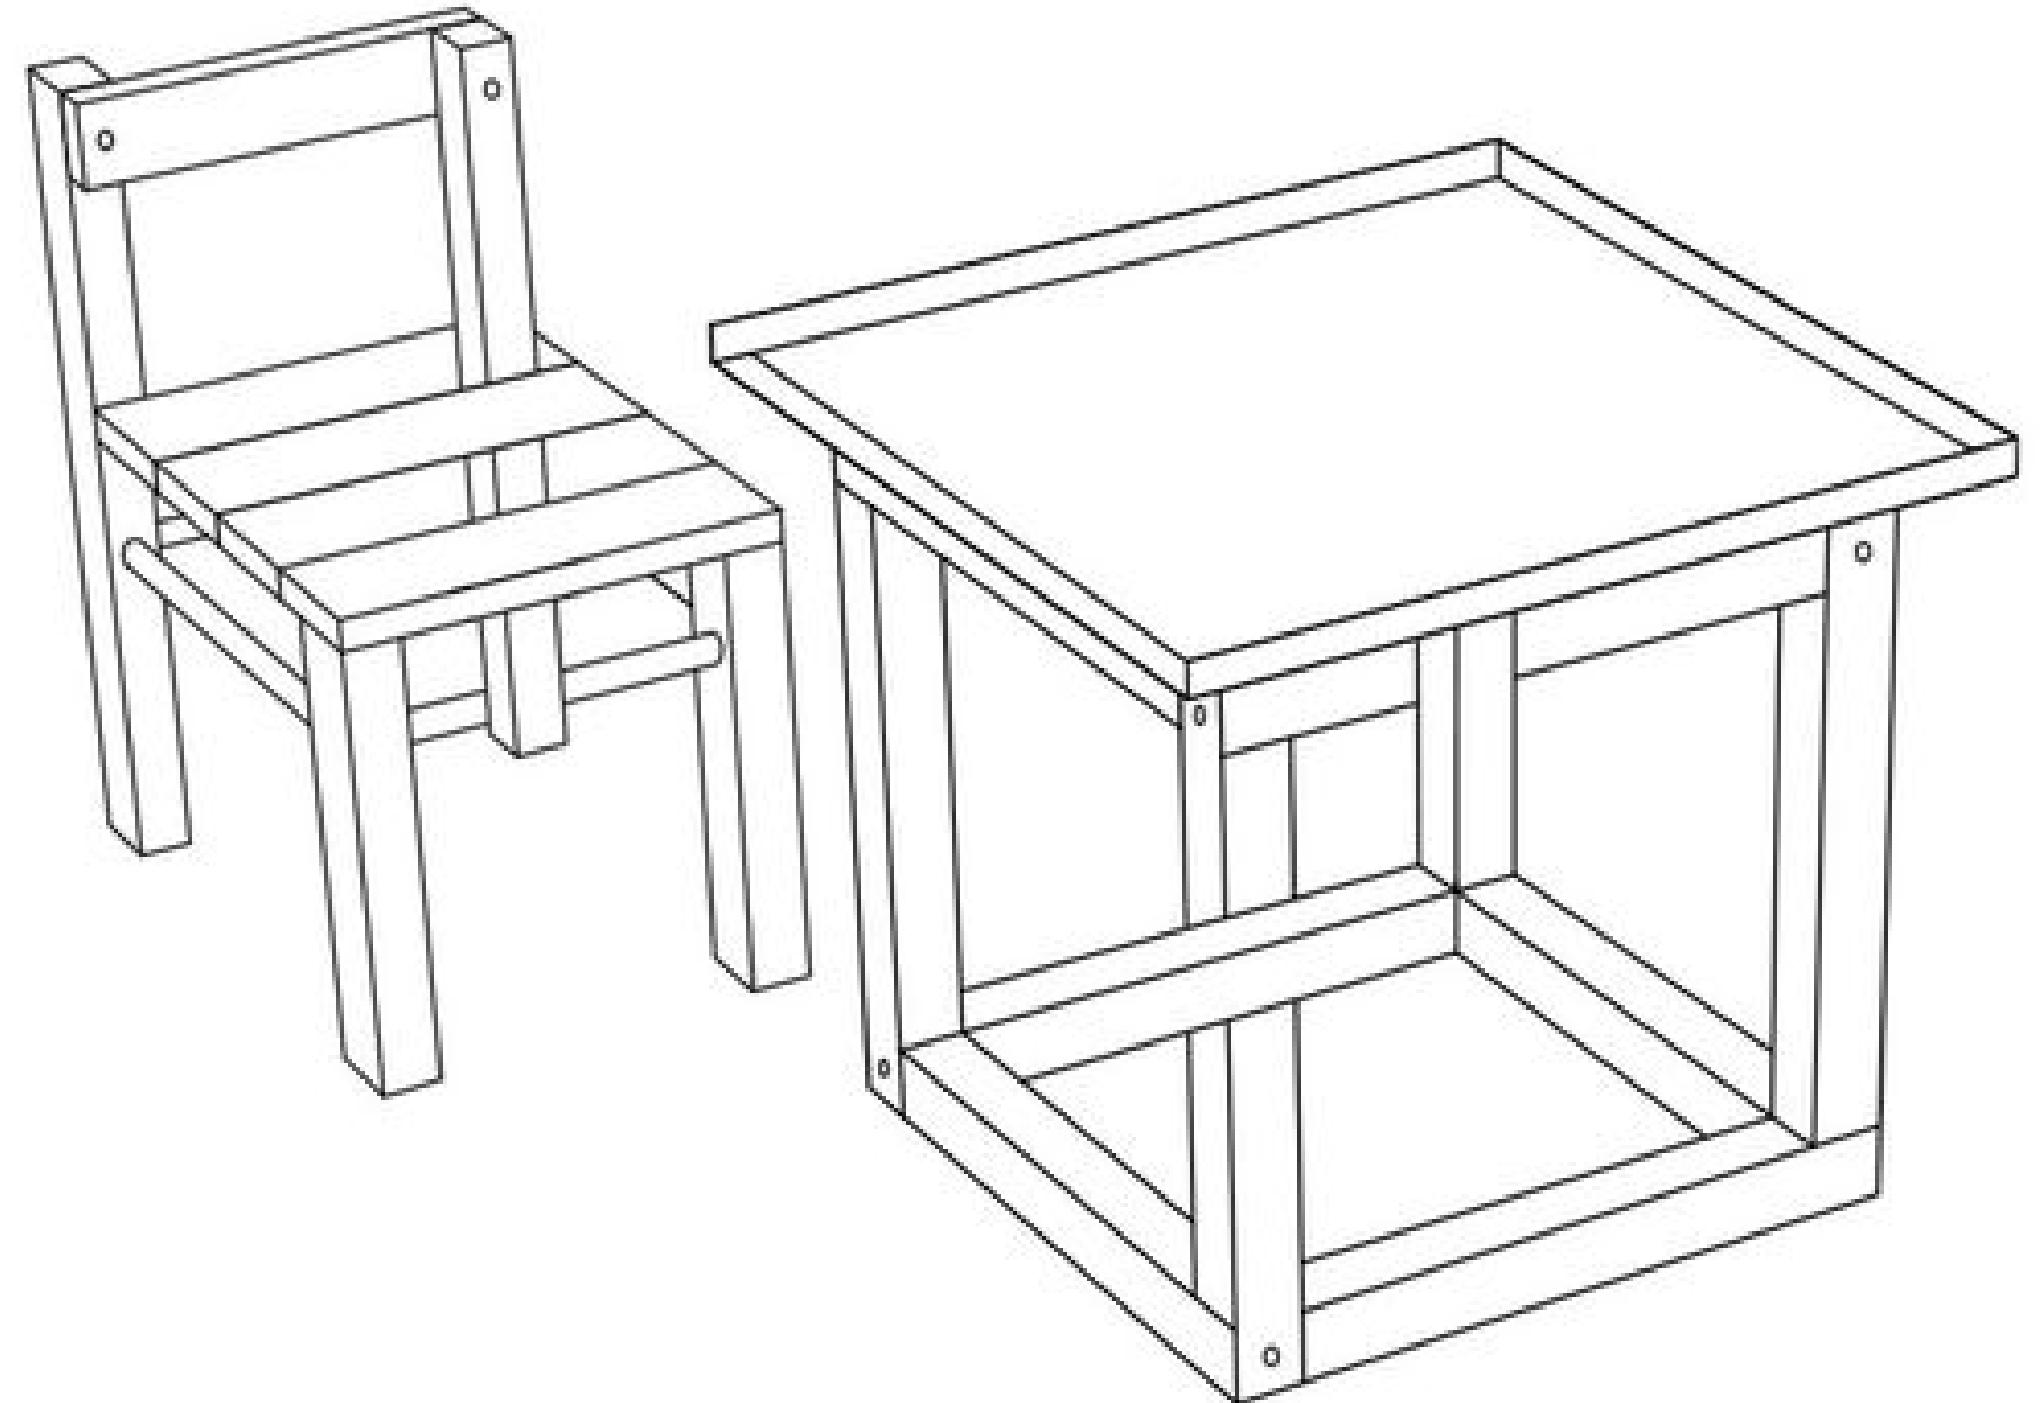
# Workshop: ARKit in Depth



RW  
RAYKEA!

# PARADÖX

play table and chair



# PLANE DETECTION



R  
W

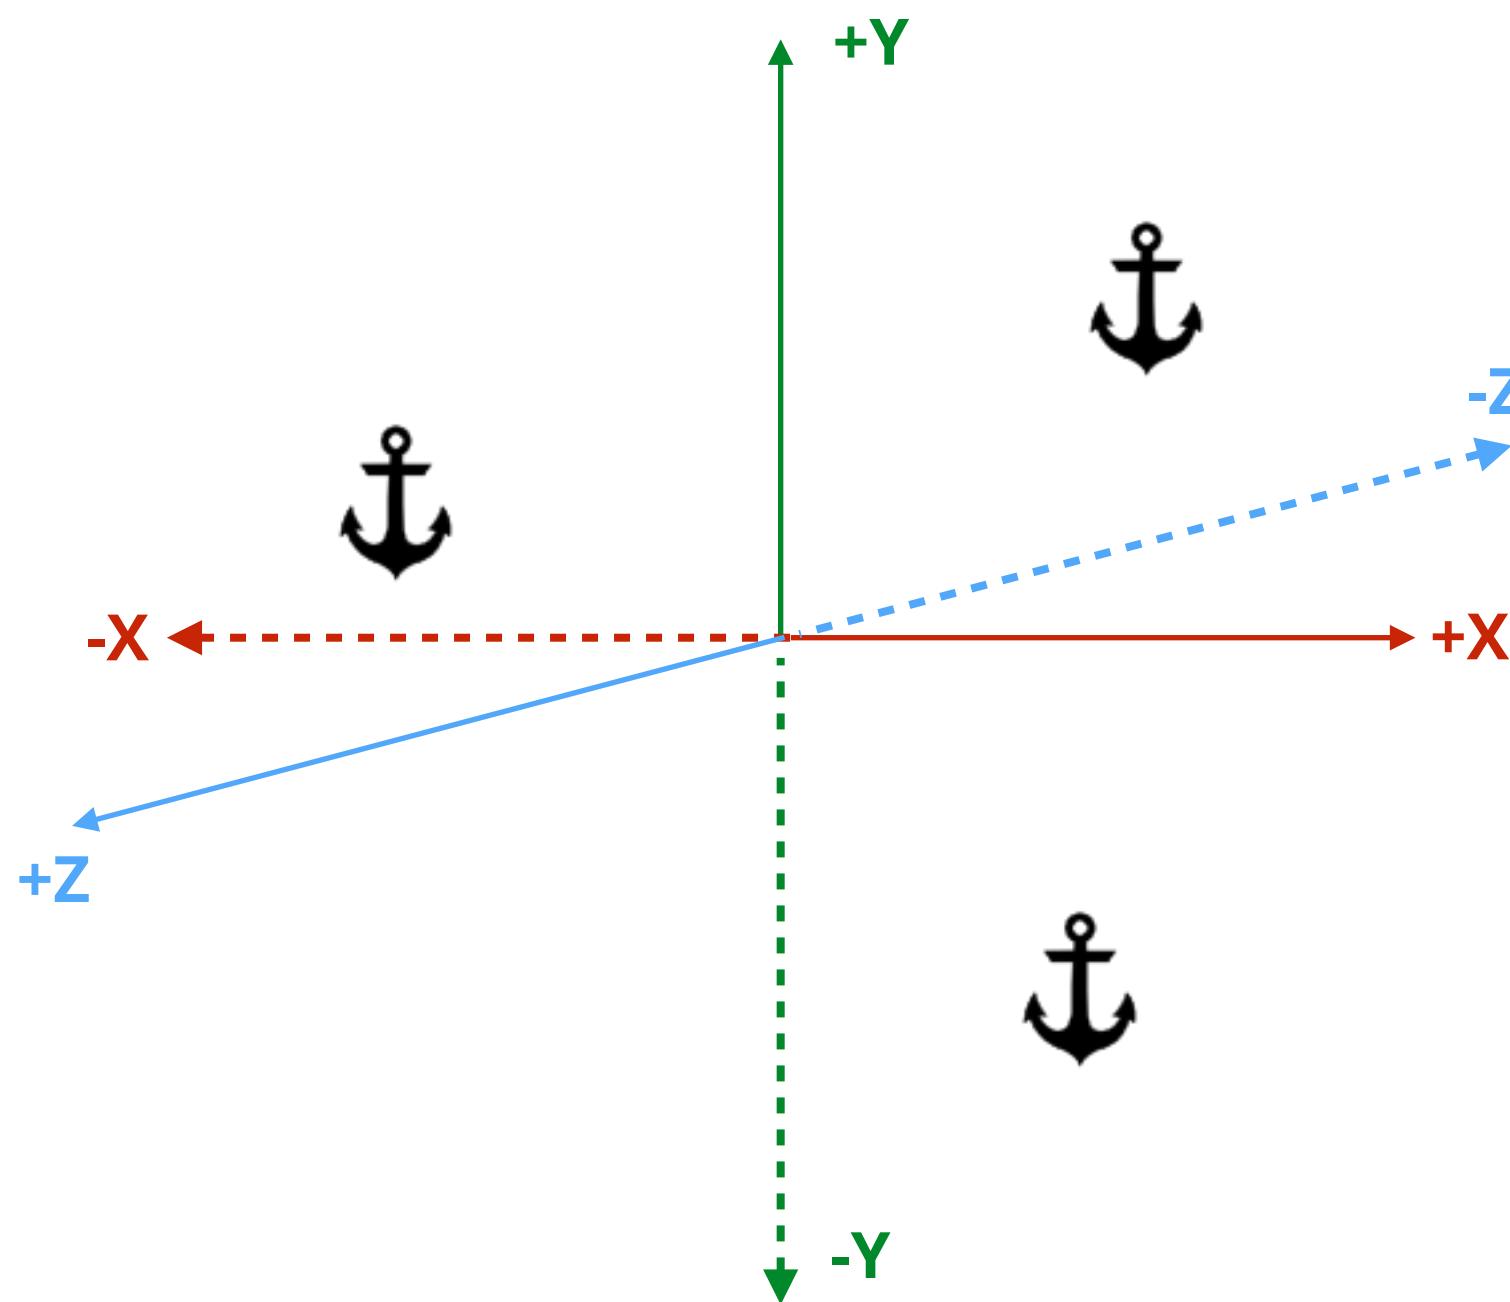
# ACTIVATING PLANE DETECTION



These two **ARWorldTrackingConfiguration** settings will enable horizontal plane detection:

```
config.worldAlignment = .gravity  
// or  
config.worldAlignment = .gravityAndHeading  
  
config.planeDetection = .horizontal  
// or  
config.planeDetection = .vertical  
// or  
config.planeDetection = [.horizontal,  
                        .vertical]
```

# AR ANCHORS



An **ARAnchor** is a reference point in 3D space, marking *position* and *orientation* in an AR scene.

They can be added to a scene in 2 ways:

- 1. Programmatically, by your code.**

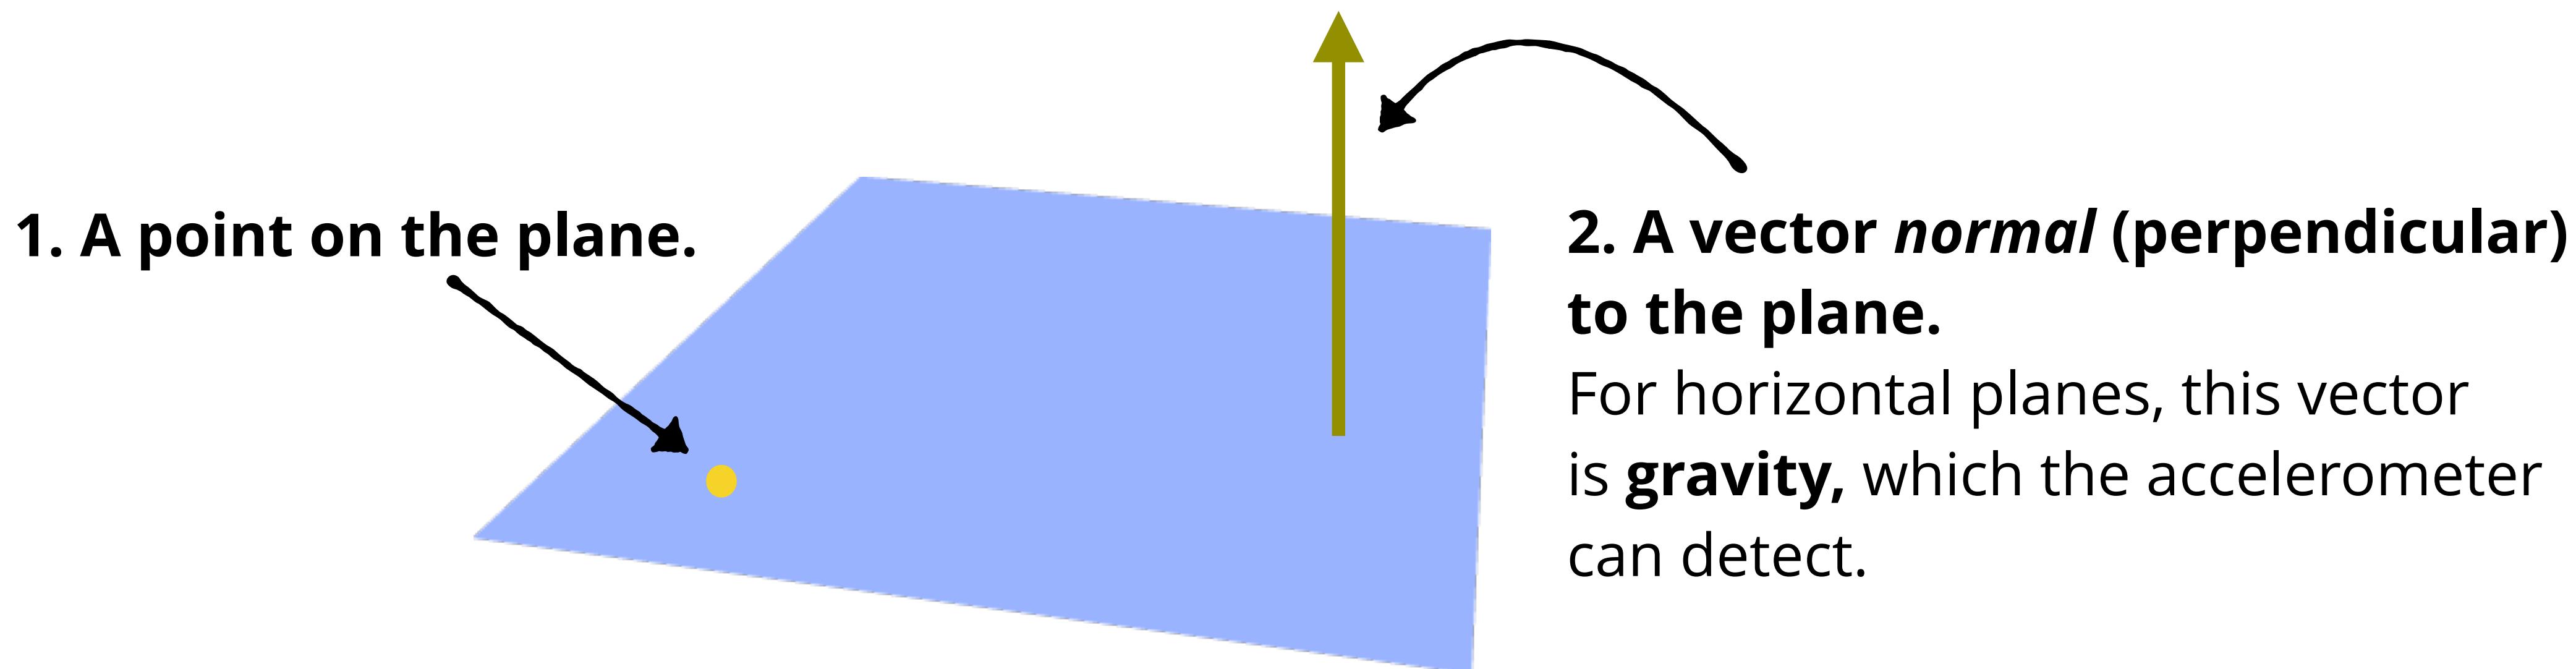
Do this to keep track of specific objects or locations.

- 2. Automatically, by ARKit.**

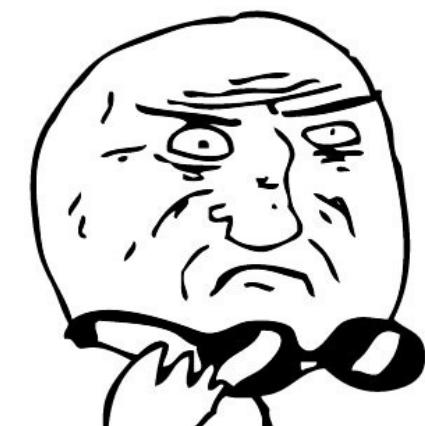
With plane detection turned on, ARKit adds **ARPlaneAnchors** (a subclass of **ARAnchor**) to the scene.

# How ARKIT DETECTS HORIZONTAL PLANES

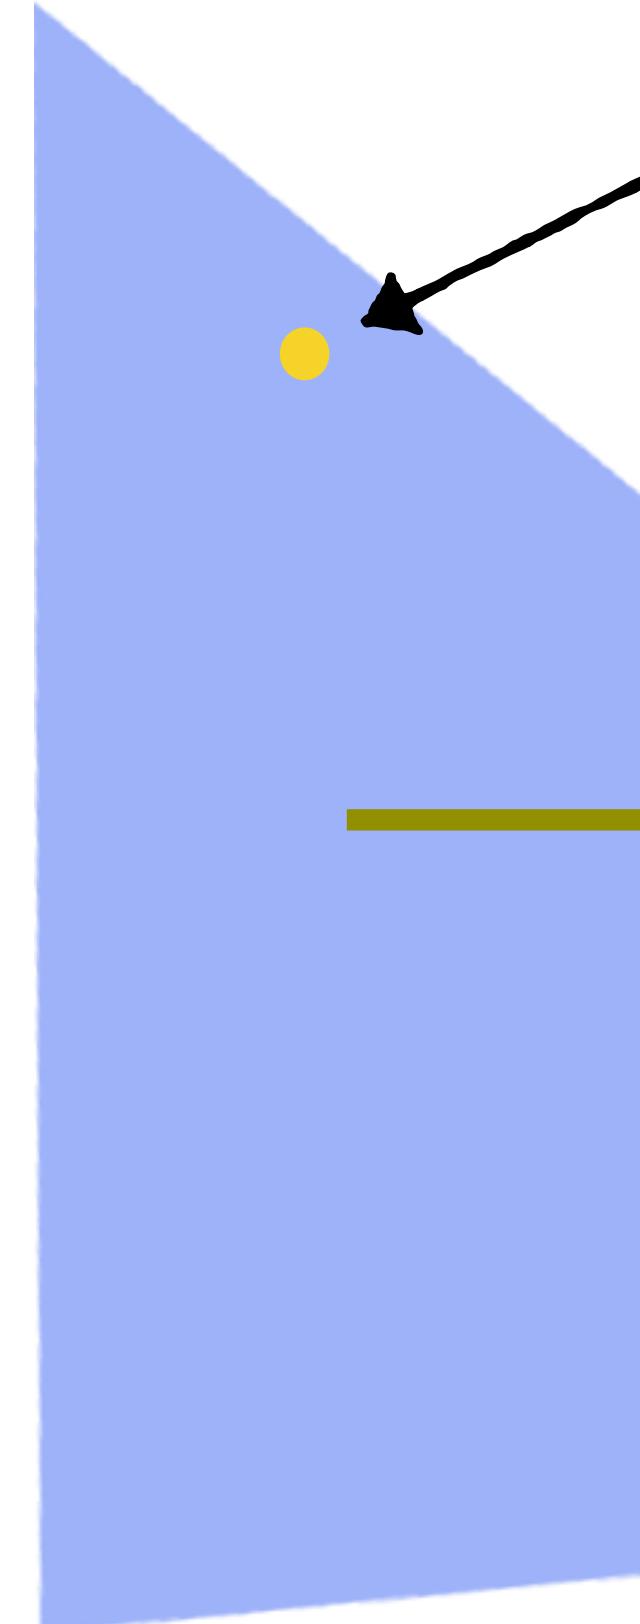
The equation for a plane is defined by 2 things...



(Don't believe this? Try plane detection with the AR configuration's **worldAlignment** property to **camera**, the one setting that doesn't align the Y-axis with gravity.)

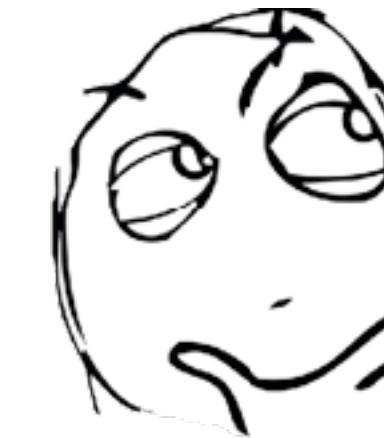


# DETECTING VERTICAL PLANES IS TRICKIER



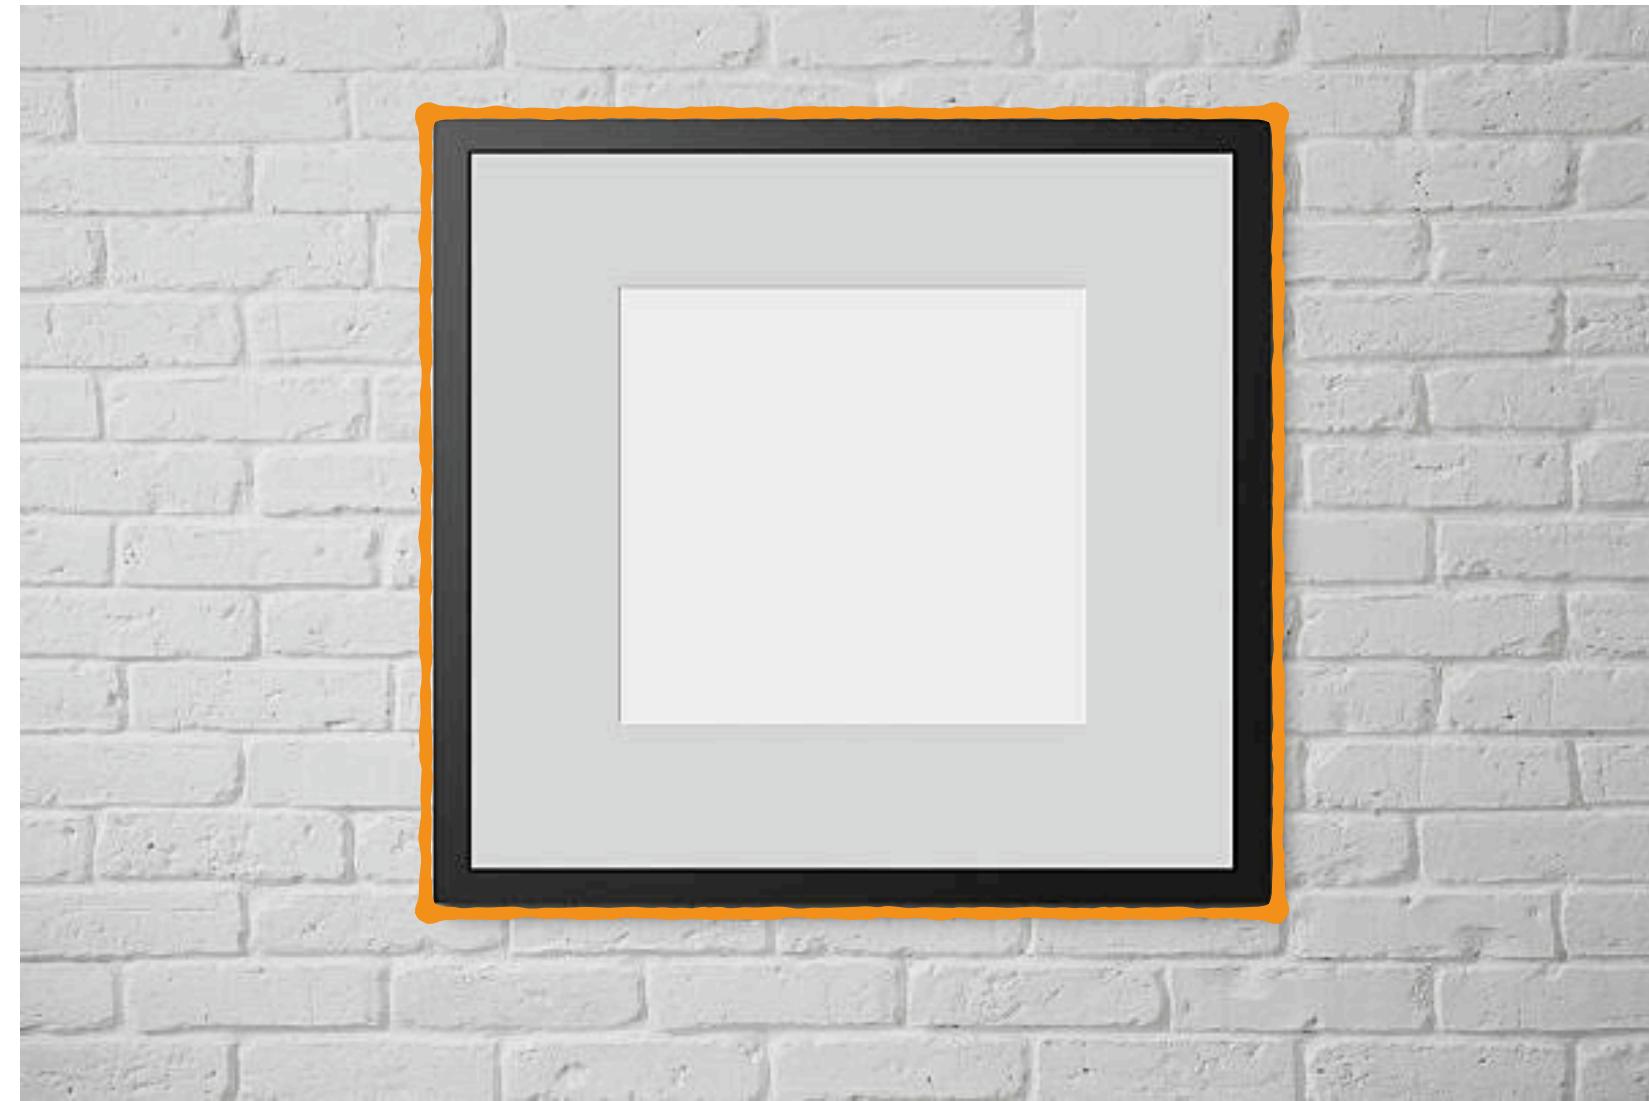
1. ARKit can find points on a vertical plane.  
It can find feature points on vertical surfaces easily.

2. The hard part is *finding the plane's normal*.  
With vertical planes, we don't have a handy force like gravity to rely on to find the normal.

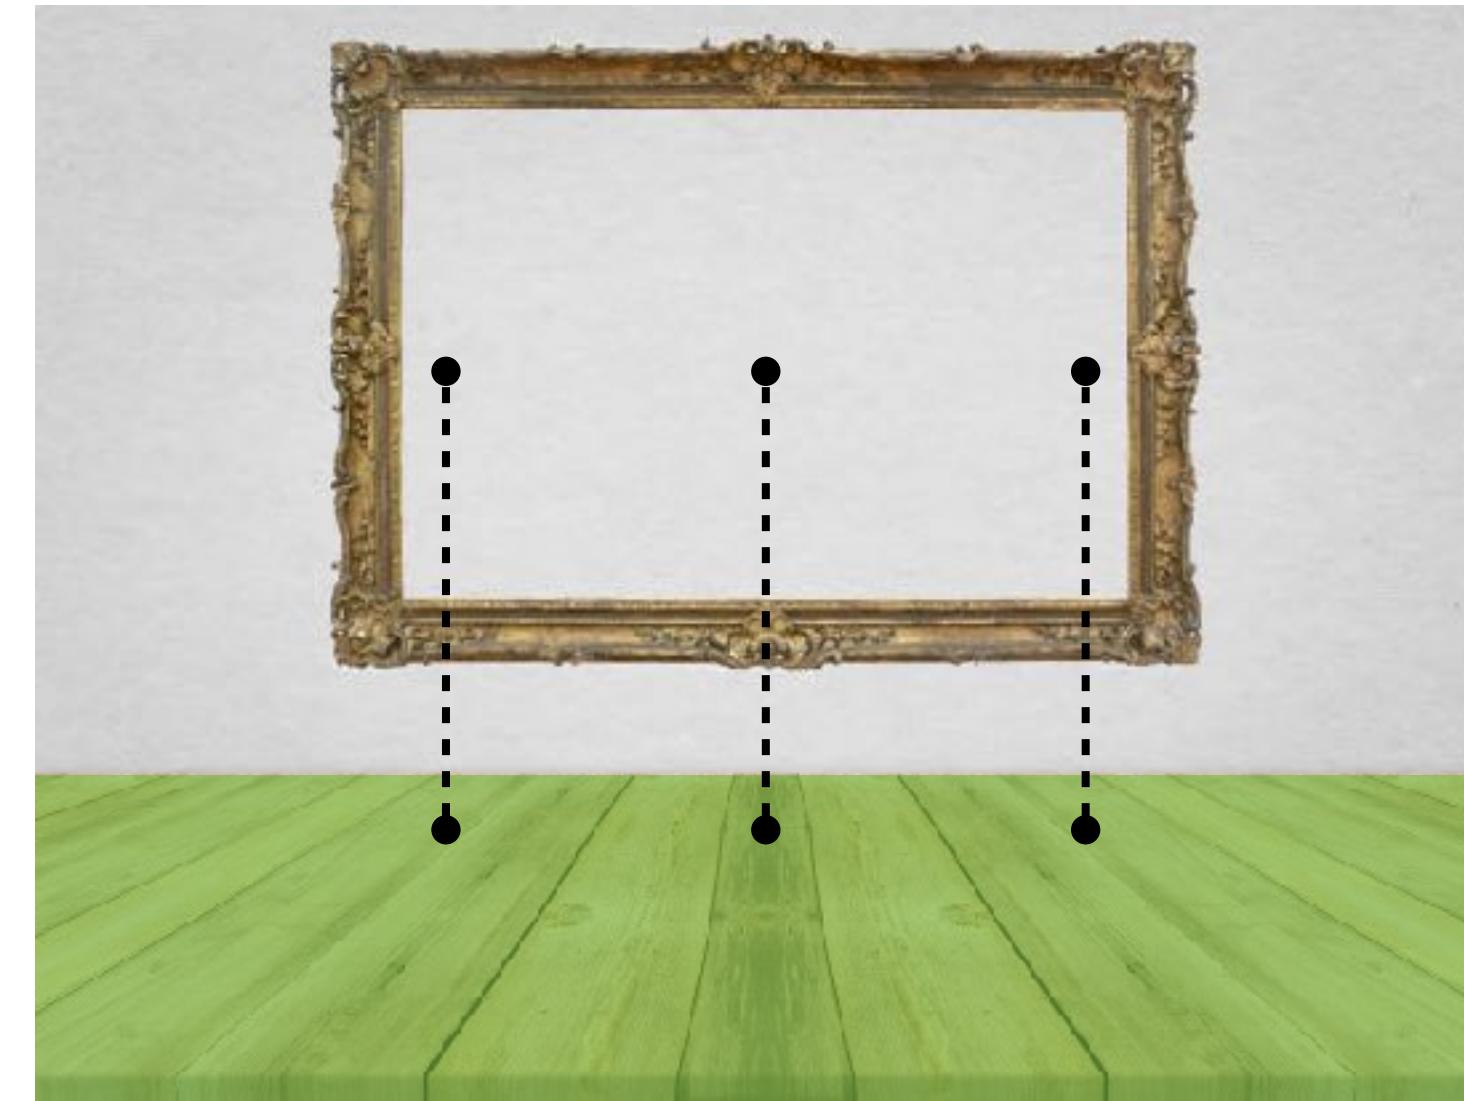


(Special bragging rights to the first person who can answer this question:  
“Why can’t we simply use a vector that’s *perpendicular* to gravity as the normal for vertical surfaces?”)

# SO HOW DOES ARKIT DETECT VERTICAL PLANES?



High-contrast borders



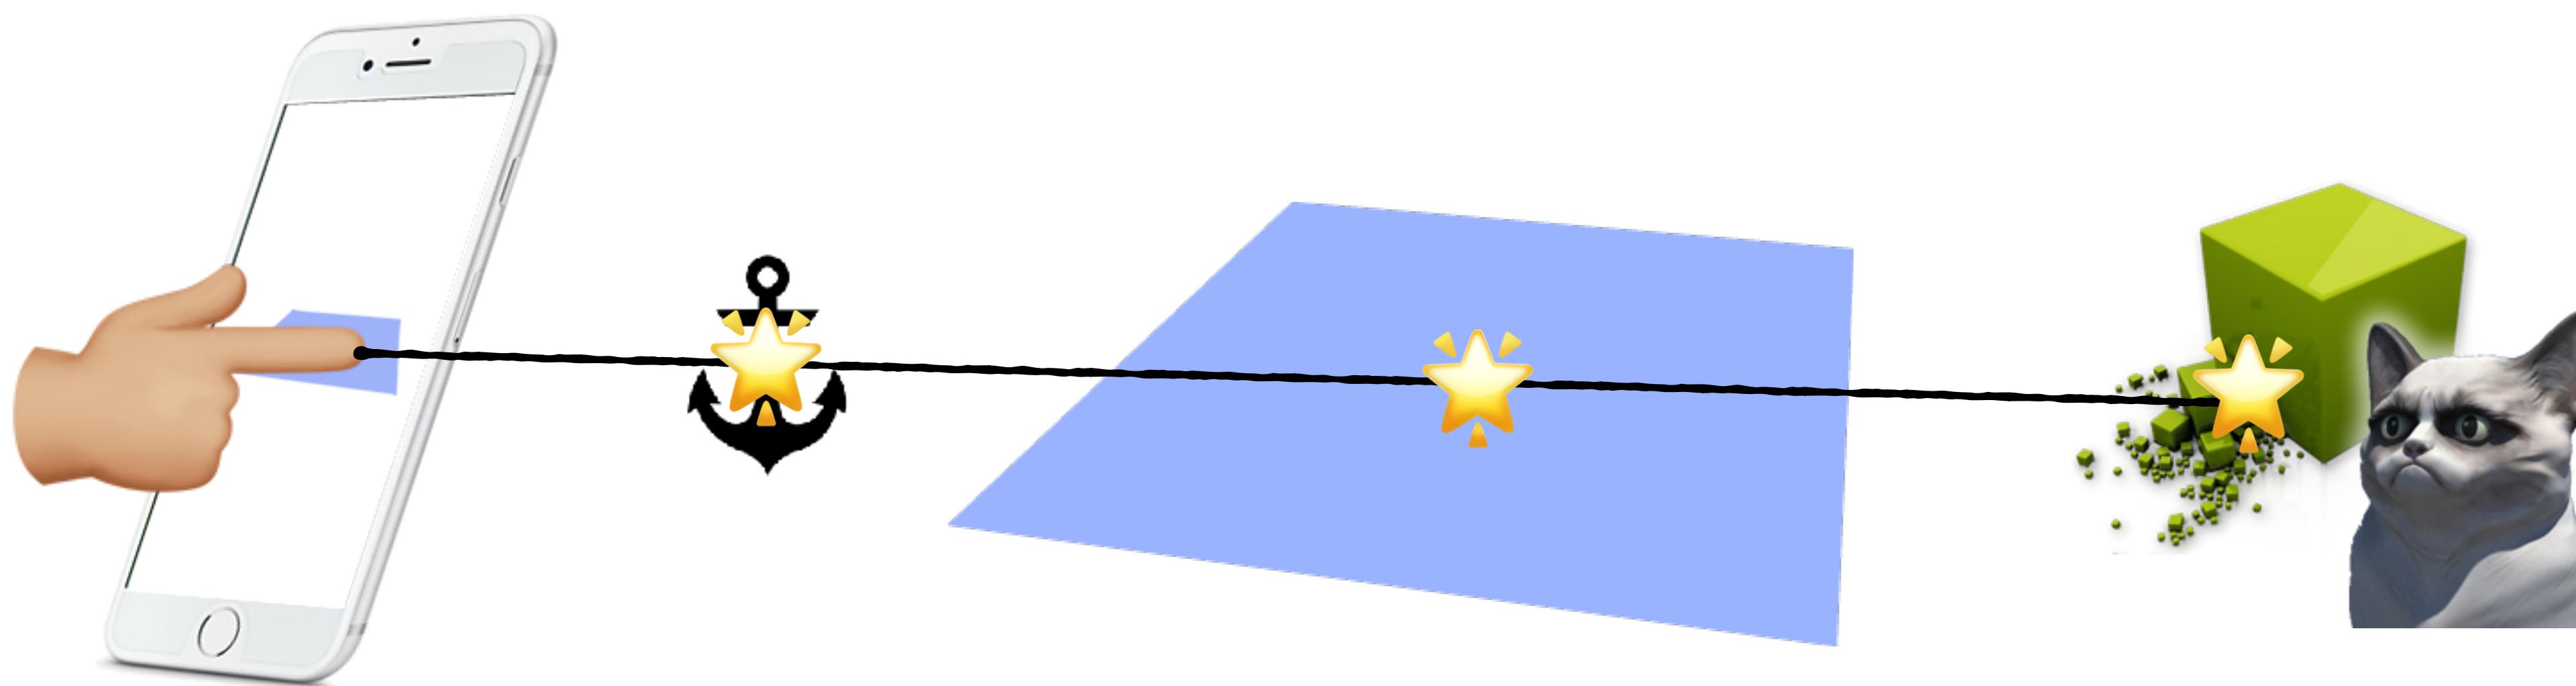
Points and  
a known horizontal plane

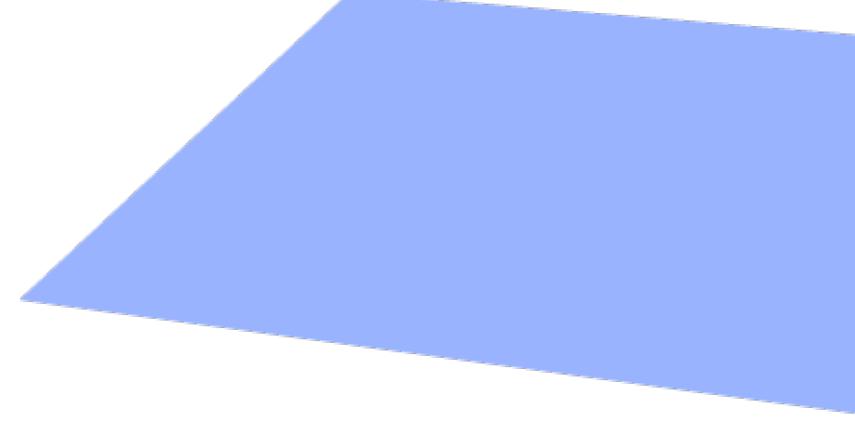


Raycasting

# HIT TESTING

Hit testing answers the question “Does this 2D screen coordinate correspond to the 3D coordinates on an AR anchor, real-world feature, or virtual object?



[  ,  ,  ]



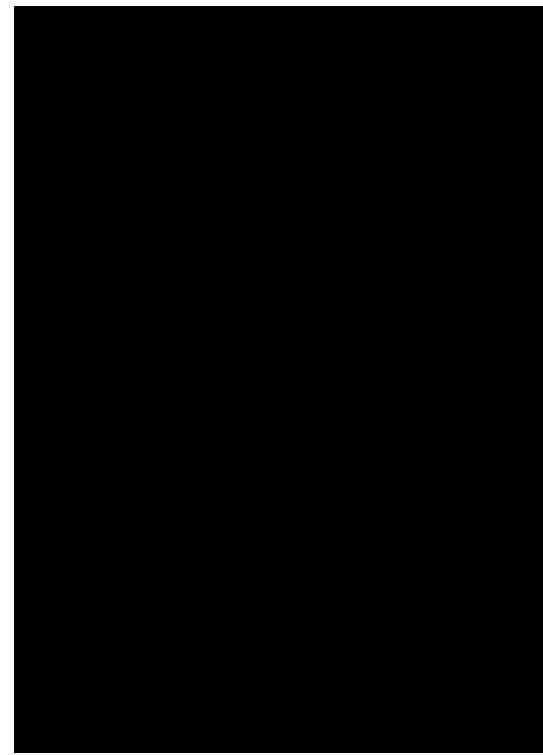
# WHAT CAUSES TRACKING TO FAIL?



Too few features



Too much motion



Too little light



Relocating

# HOW WELL IS THE CAMERA TRACKING?

This ARKit delegate method gets called whenever the camera's tracking state changes...

```
session(_ :cameraDidChangeTrackingState:)
```

...and it gives us this very handy property...

```
ARCamera.trackingState
```

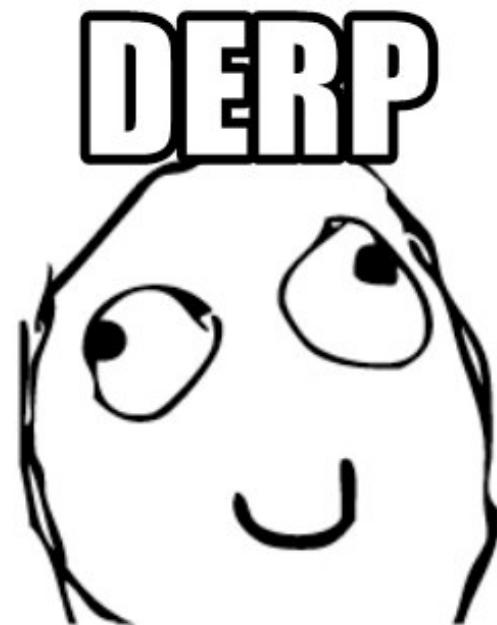
...which will contain one of three possible values:



```
.notAvailable
```

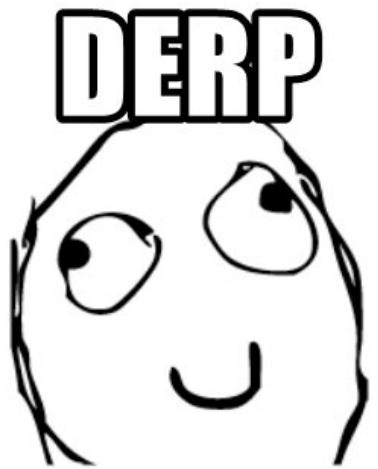


```
.normal
```



```
.limited
```

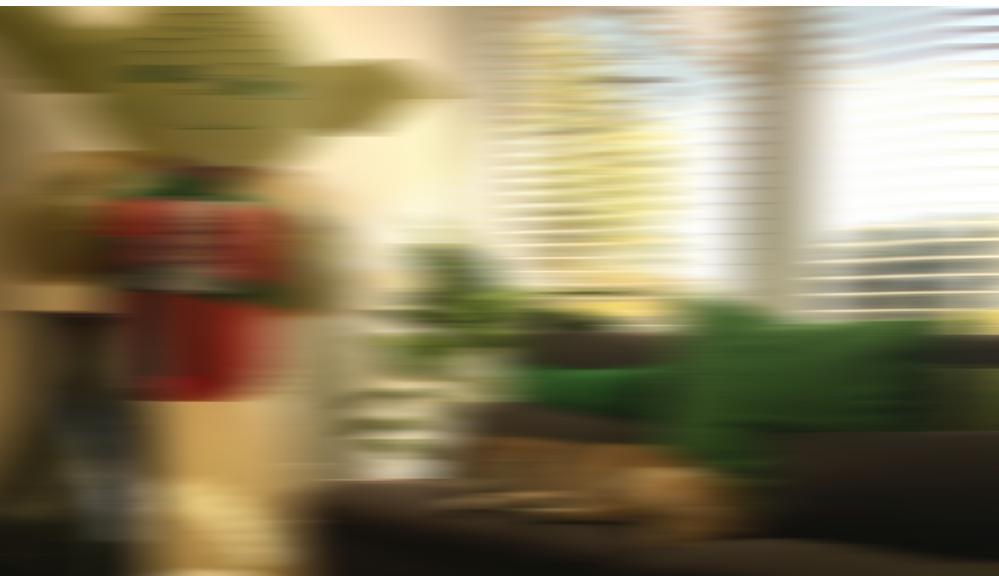
# REASONS WHY TRACKING IS LIMITED



If tracking is limited, you can check the **reason** property to see why.



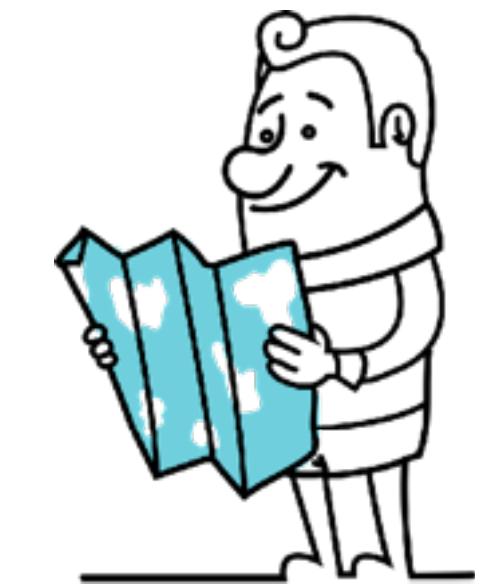
.initializing



.excessiveMotion



.insufficientFeatures



.relocalizing



# DEMO 2: RAYKEA



# How RAYKEA WORKS, PART 1

1. Continuously seek horizontal and vertical planes.



2. Cover any detected vertical planes with a poster.

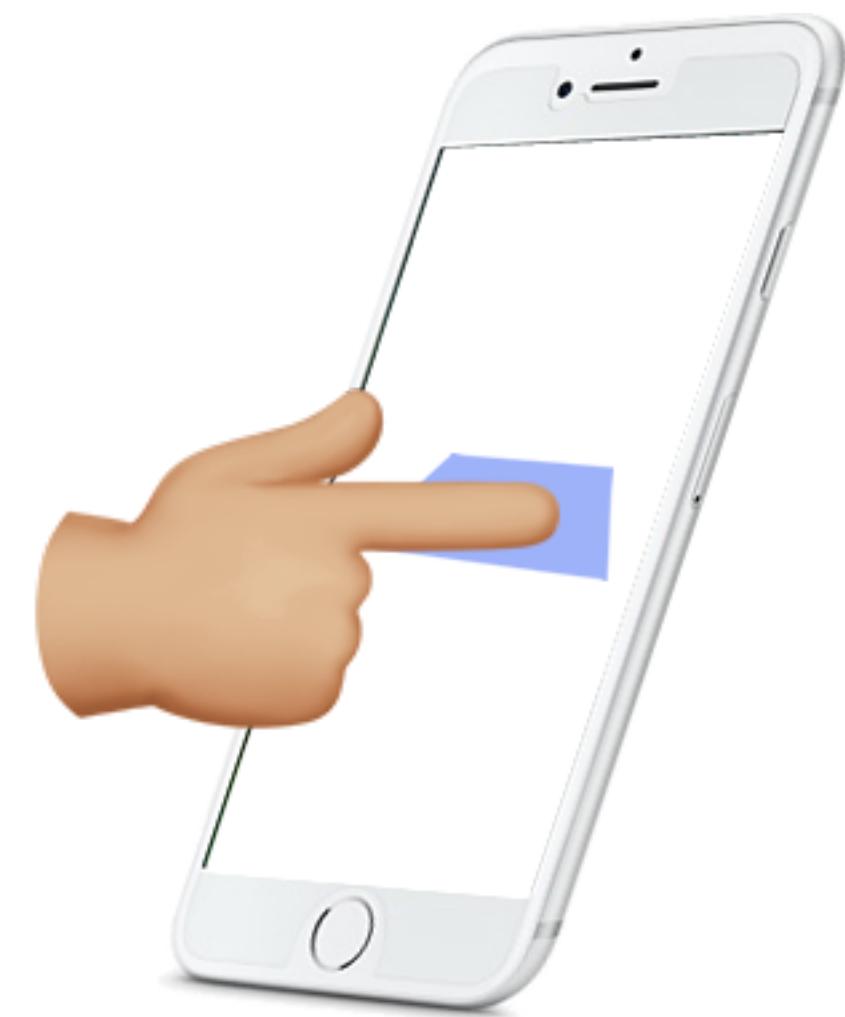


# How RAYKEA WORKS, PART 2

3. Cover any detected horizontal planes with a “place furniture here” grid.



4. When the user taps the screen, perform a hit test to see if that tap corresponds to a detected horizontal plane.

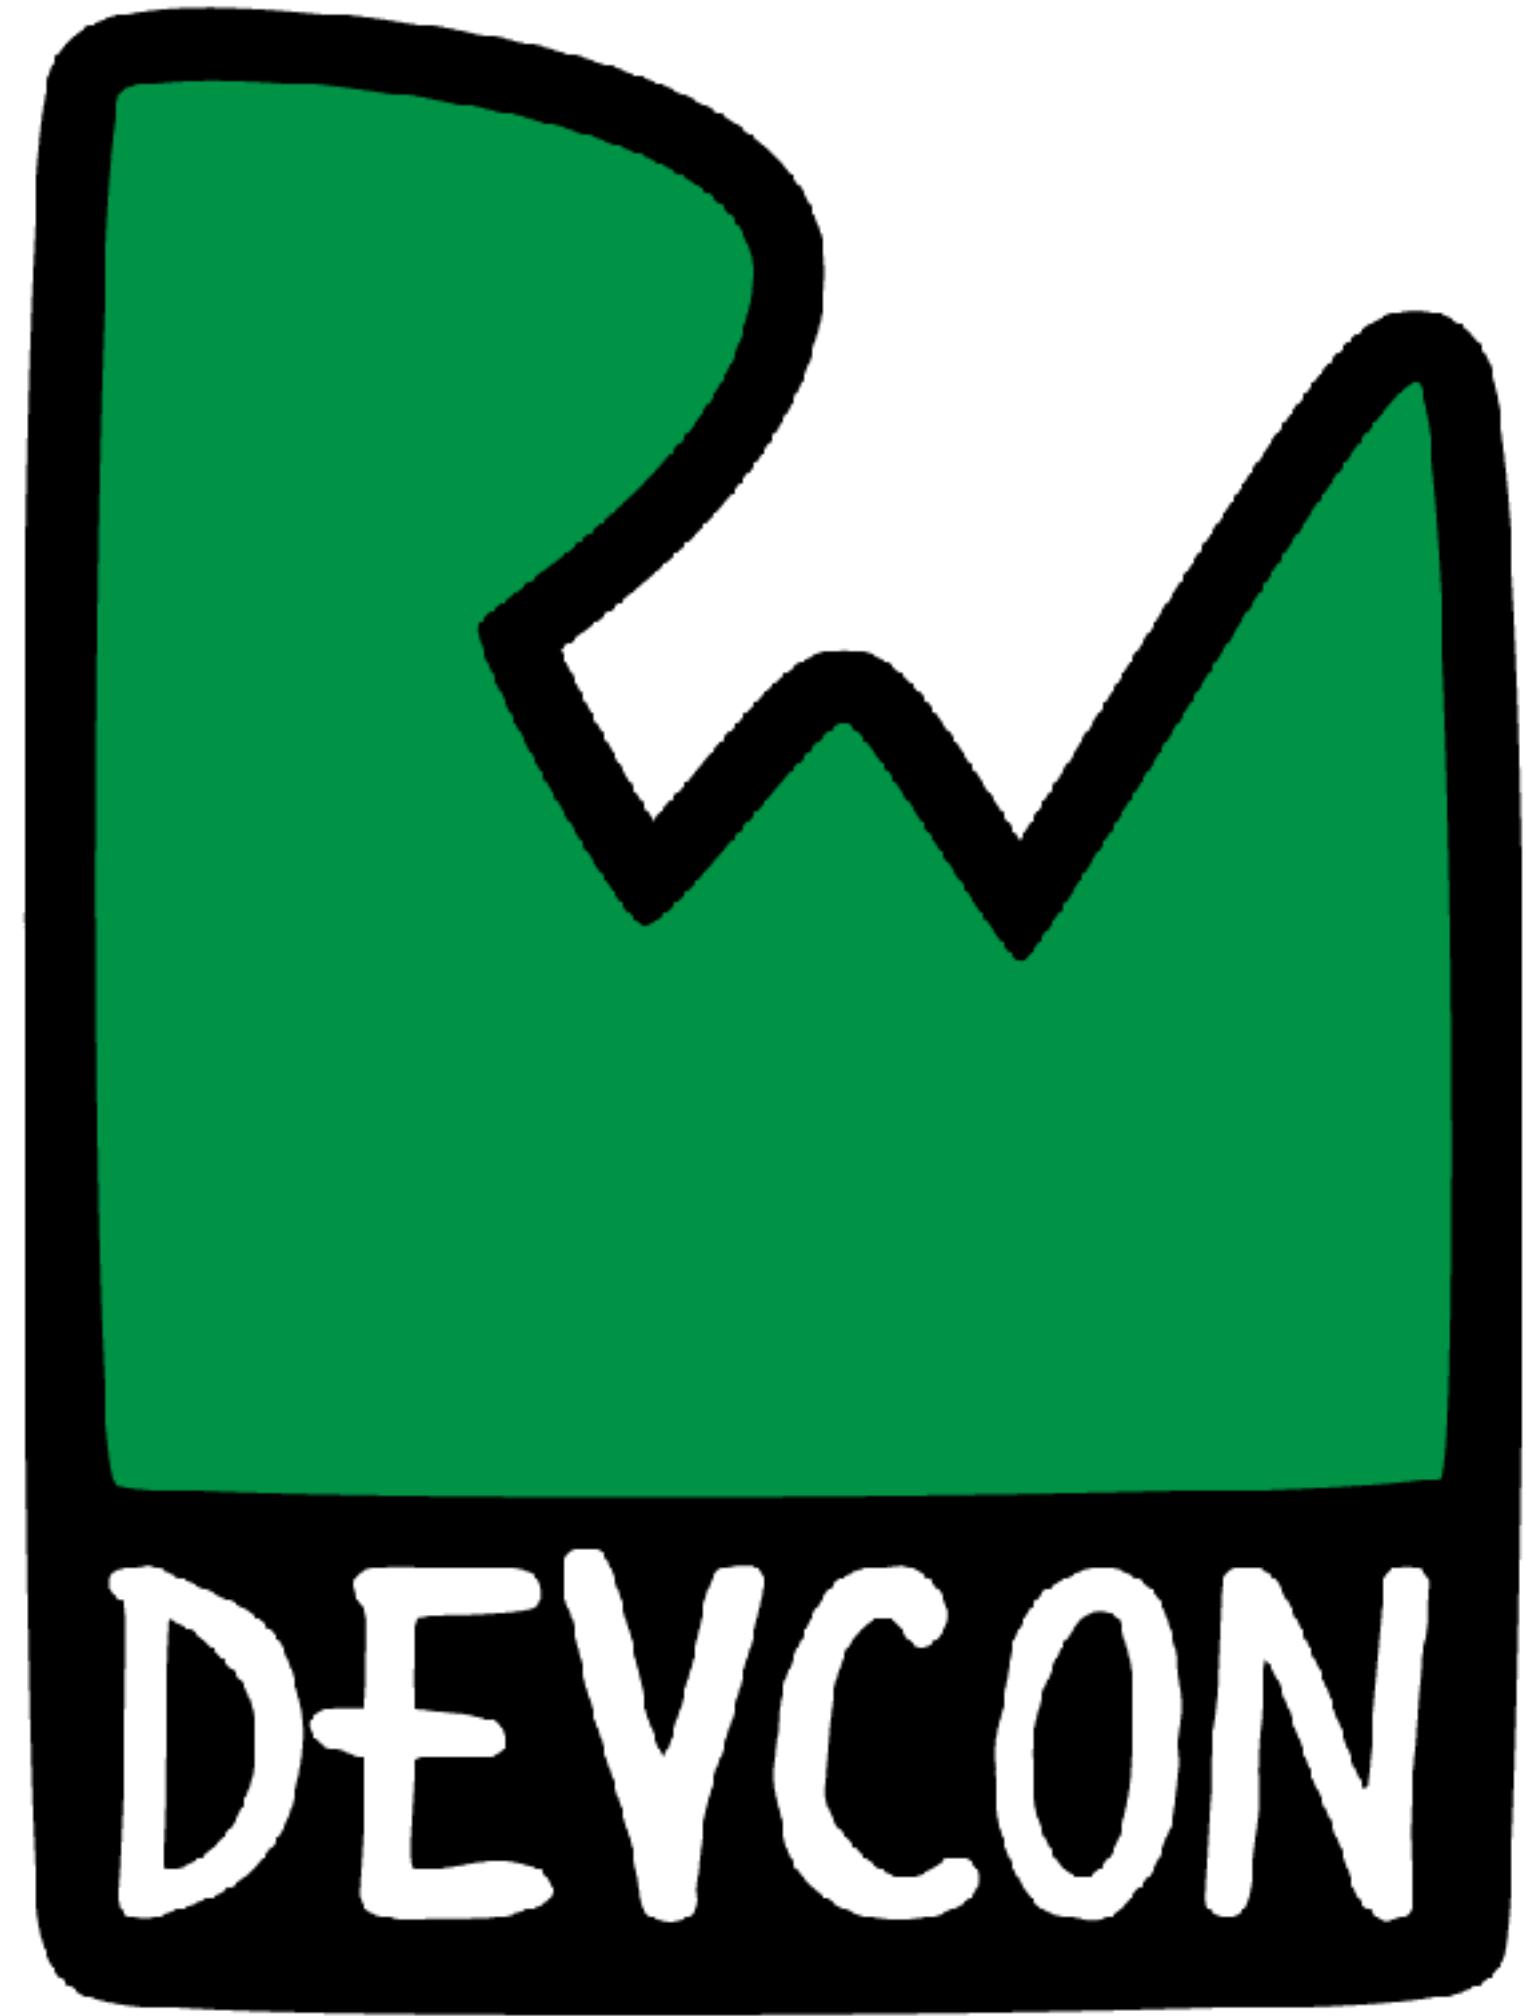


# HOW RAYKEA WORKS, PART 3

5. If the tap corresponds to a detected horizontal plane, find the real-world coordinates that correspond to that tap, and draw furniture at those coordinates.



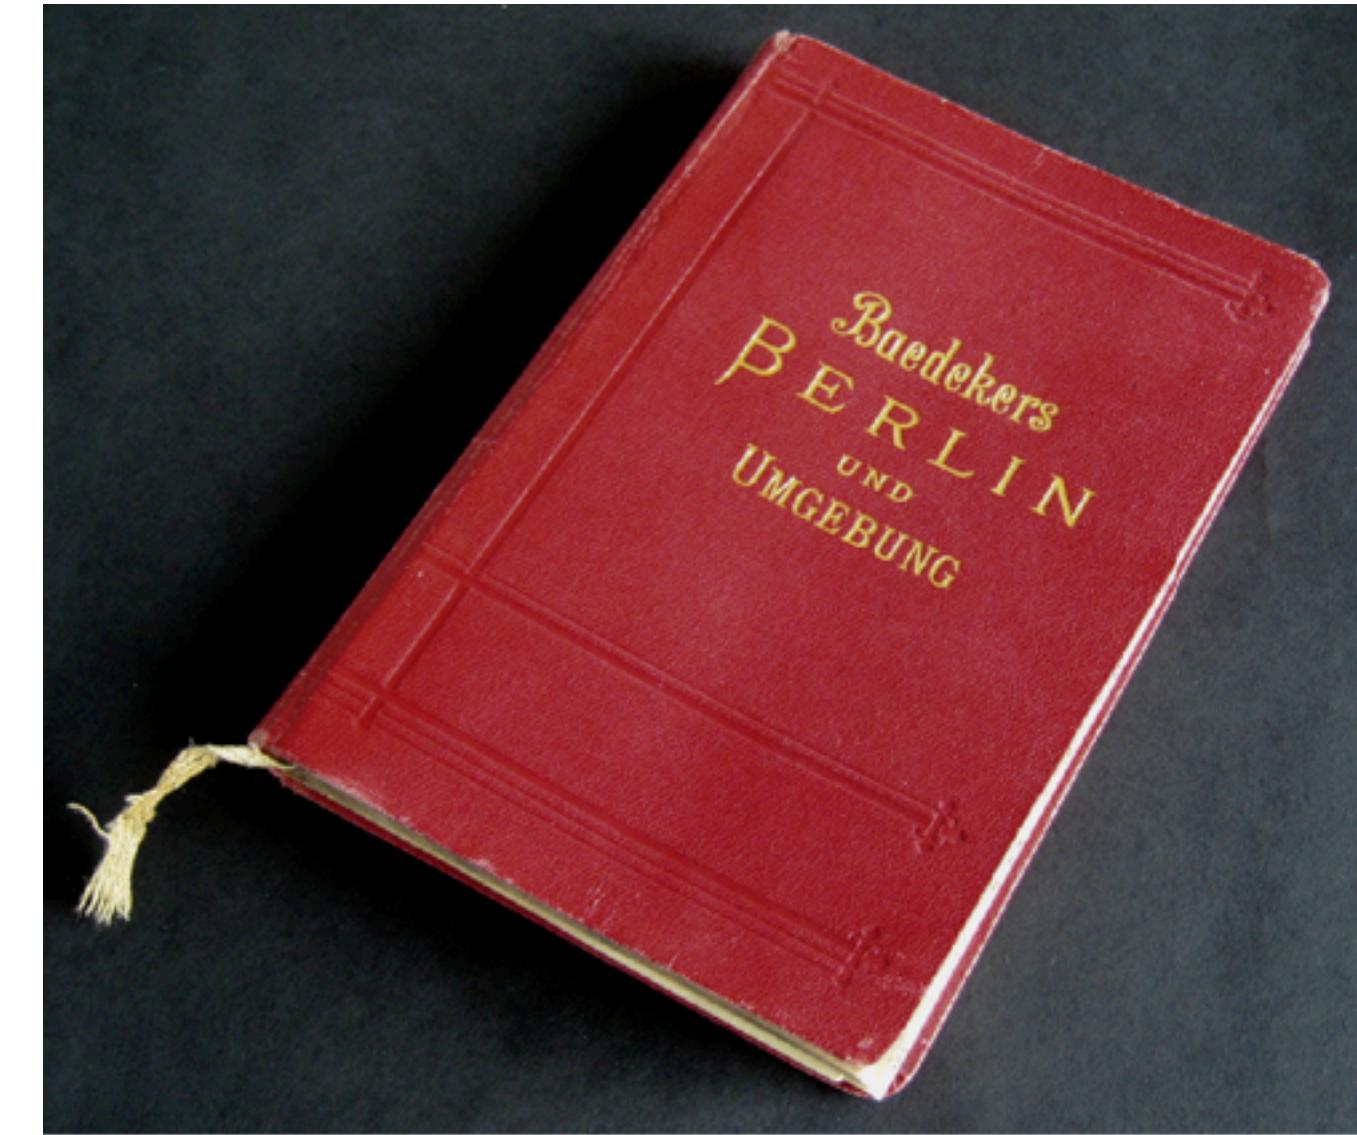
# Workshop: ARKit in Depth



RW  
BAEDEKAR

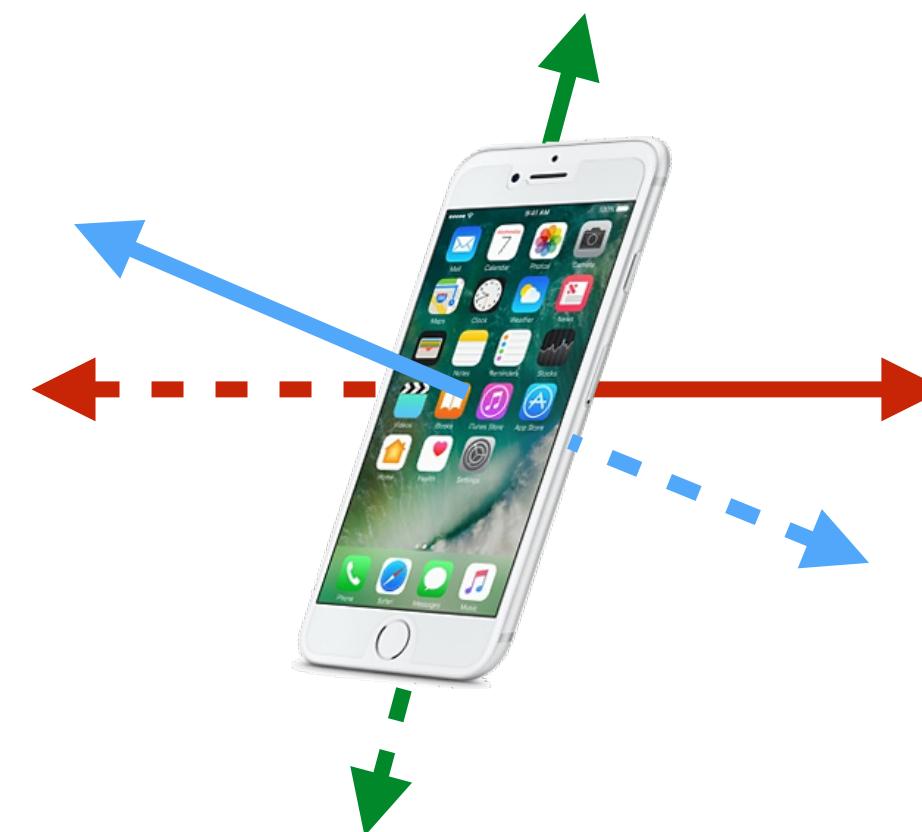
# “BAEDEKAR”?

**Baedeker** is an old colloquial term for “travel guidebook”. The term comes from Karl Baedeker, whose publishing house started publishing world travel guides under their name in 1827.



**Fun fact:** BaedekAR is the basis of an app that I'm actually writing for cruise lines and other tourism-based businesses. You just might run into it on your next vacation!

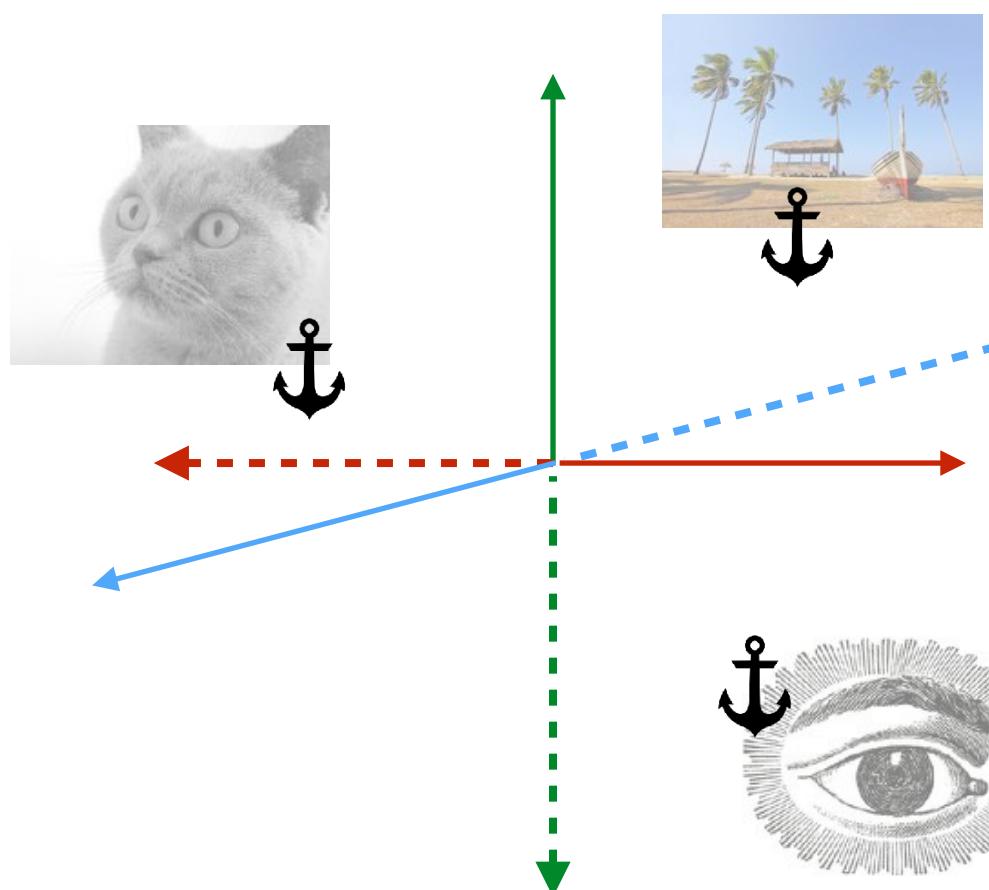
# ARKit 1.5 AND 2D IMAGE RECOGNITION



In ARKit 1.5, you can take  
an **ARWorldTrackingConfiguration**...



...and pass a set of reference images to its  
**detectionImages** property...



...then ARKit will place **ARImageAnchors** (a subclass of **ARAnchor**)  
in the scene wherever it finds any of those reference images.

# STORING 2D REFERENCE IMAGES



There's a new asset type called an **AR Resource Group** that is used to store 2D reference images for image recognition with ARKit.

It's easy to load them into a **Set of ARReferenceImages**, which you can then pass to the session's configuration via the **detectionImages** property:

```
let referenceImages = ARReferenceImage.referenceImages  
  (inGroupNamed: "AR Resources",  
   bundle: nil)  
config.detectionImages = referenceImages  
  
// Start the AR session.  
sceneView.session.run(config, options: [])
```

# THE REAL-WORLD SIZE OF REFERENCE IMAGES

---

In addition to providing a set of reference images, you also have to provide the real-world size of those images. Why?



It's because objects that are closer appear to be larger...

...than objects that are farther away.

**Providing the real-world size of reference images is mandatory.** If you don't provide the size for a reference image, ARKit doesn't report its presence in a scene.



# RESPONDING TO RECOGNIZED IMAGES



```
func renderer(_ renderer: SCNSceneRenderer, didAdd node: SCNNNode, for anchor: ARAnchor) {  
    // What kind of anchor are we dealing with?  
    if let imageAnchor = anchor as? ARImageAnchor {  
        // Code to handle recognized 2D images goes here  
    }  
}
```

This is the same method that we used when detecting planes in *Raykea!*

# 2D IMAGE RECOGNITION TIPS AND TRICKS, PART 1



Partially obscured images

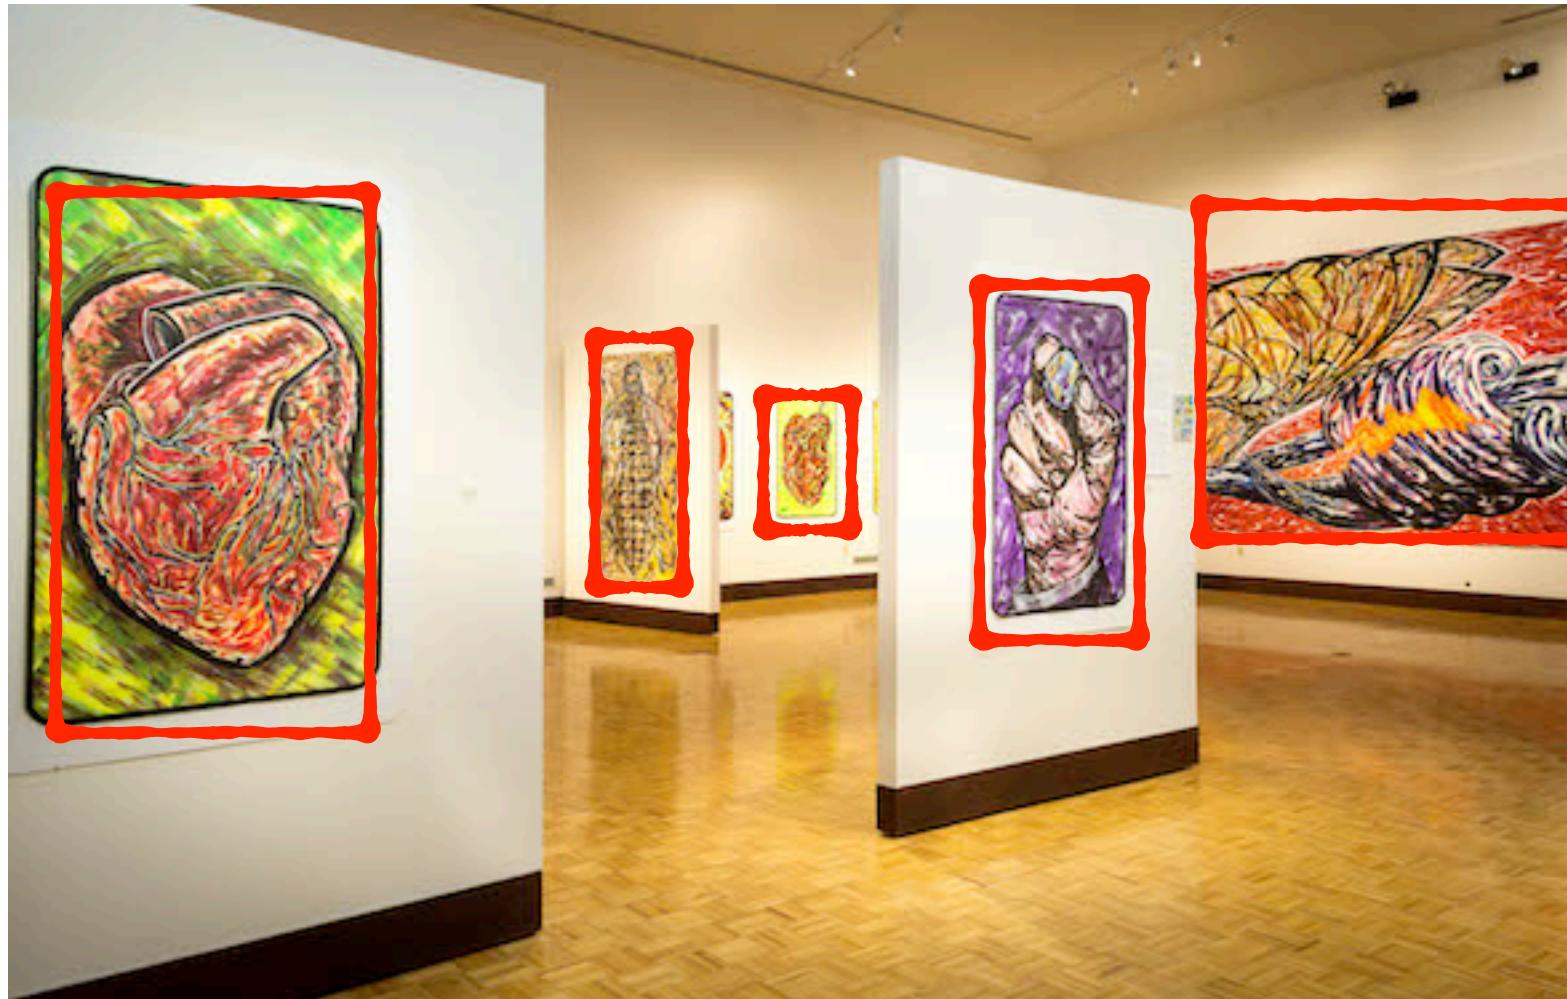


Works best with flat surfaces



Poor lighting conditions  
and effects

# 2D IMAGE RECOGNITION TIPS AND TRICKS, PART 2

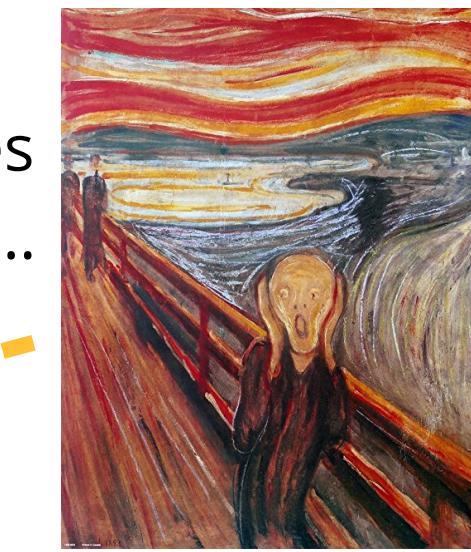


“Visual iBeacons”

This  
reference image...



...matches  
with this...



...and this!



Experiment with reference images



# DRAWING 3D TEXT IN SCENEKIT

Define the text content by creating an SCNText instance:

```
let myText = SCNText(string: "Hello, AR world!", extrusionDepth: 0.01)
```

Define the text's appearance by setting its properties:

```
myText.font = markerFont  
myText.alignmentMode = kCAAlignmentCenter  
myText.firstMaterial?.diffuse.contents = UIColor.yellow
```

Create a SceneKit node using the text as its geometry, set its position, and add it to the scene:

```
let myTextNode = SCNNode(geometry: myText)  
myTextNode.position = SCNVector3(0.0, 0.0, 0.0)  
sceneView.scene.rootNode.addChildNode(myTextNode)
```



# THE PROBLEM WITH 3D TEXT IN AN AR WORLD

---



**THIS IS  
HARD TO READ**

If you walk past the position  
of an AR text node and  
look back to read it,  
you'll see something like this.

**THIS IS  
EASY TO READ**



**THIS IS  
EASY TO READ**

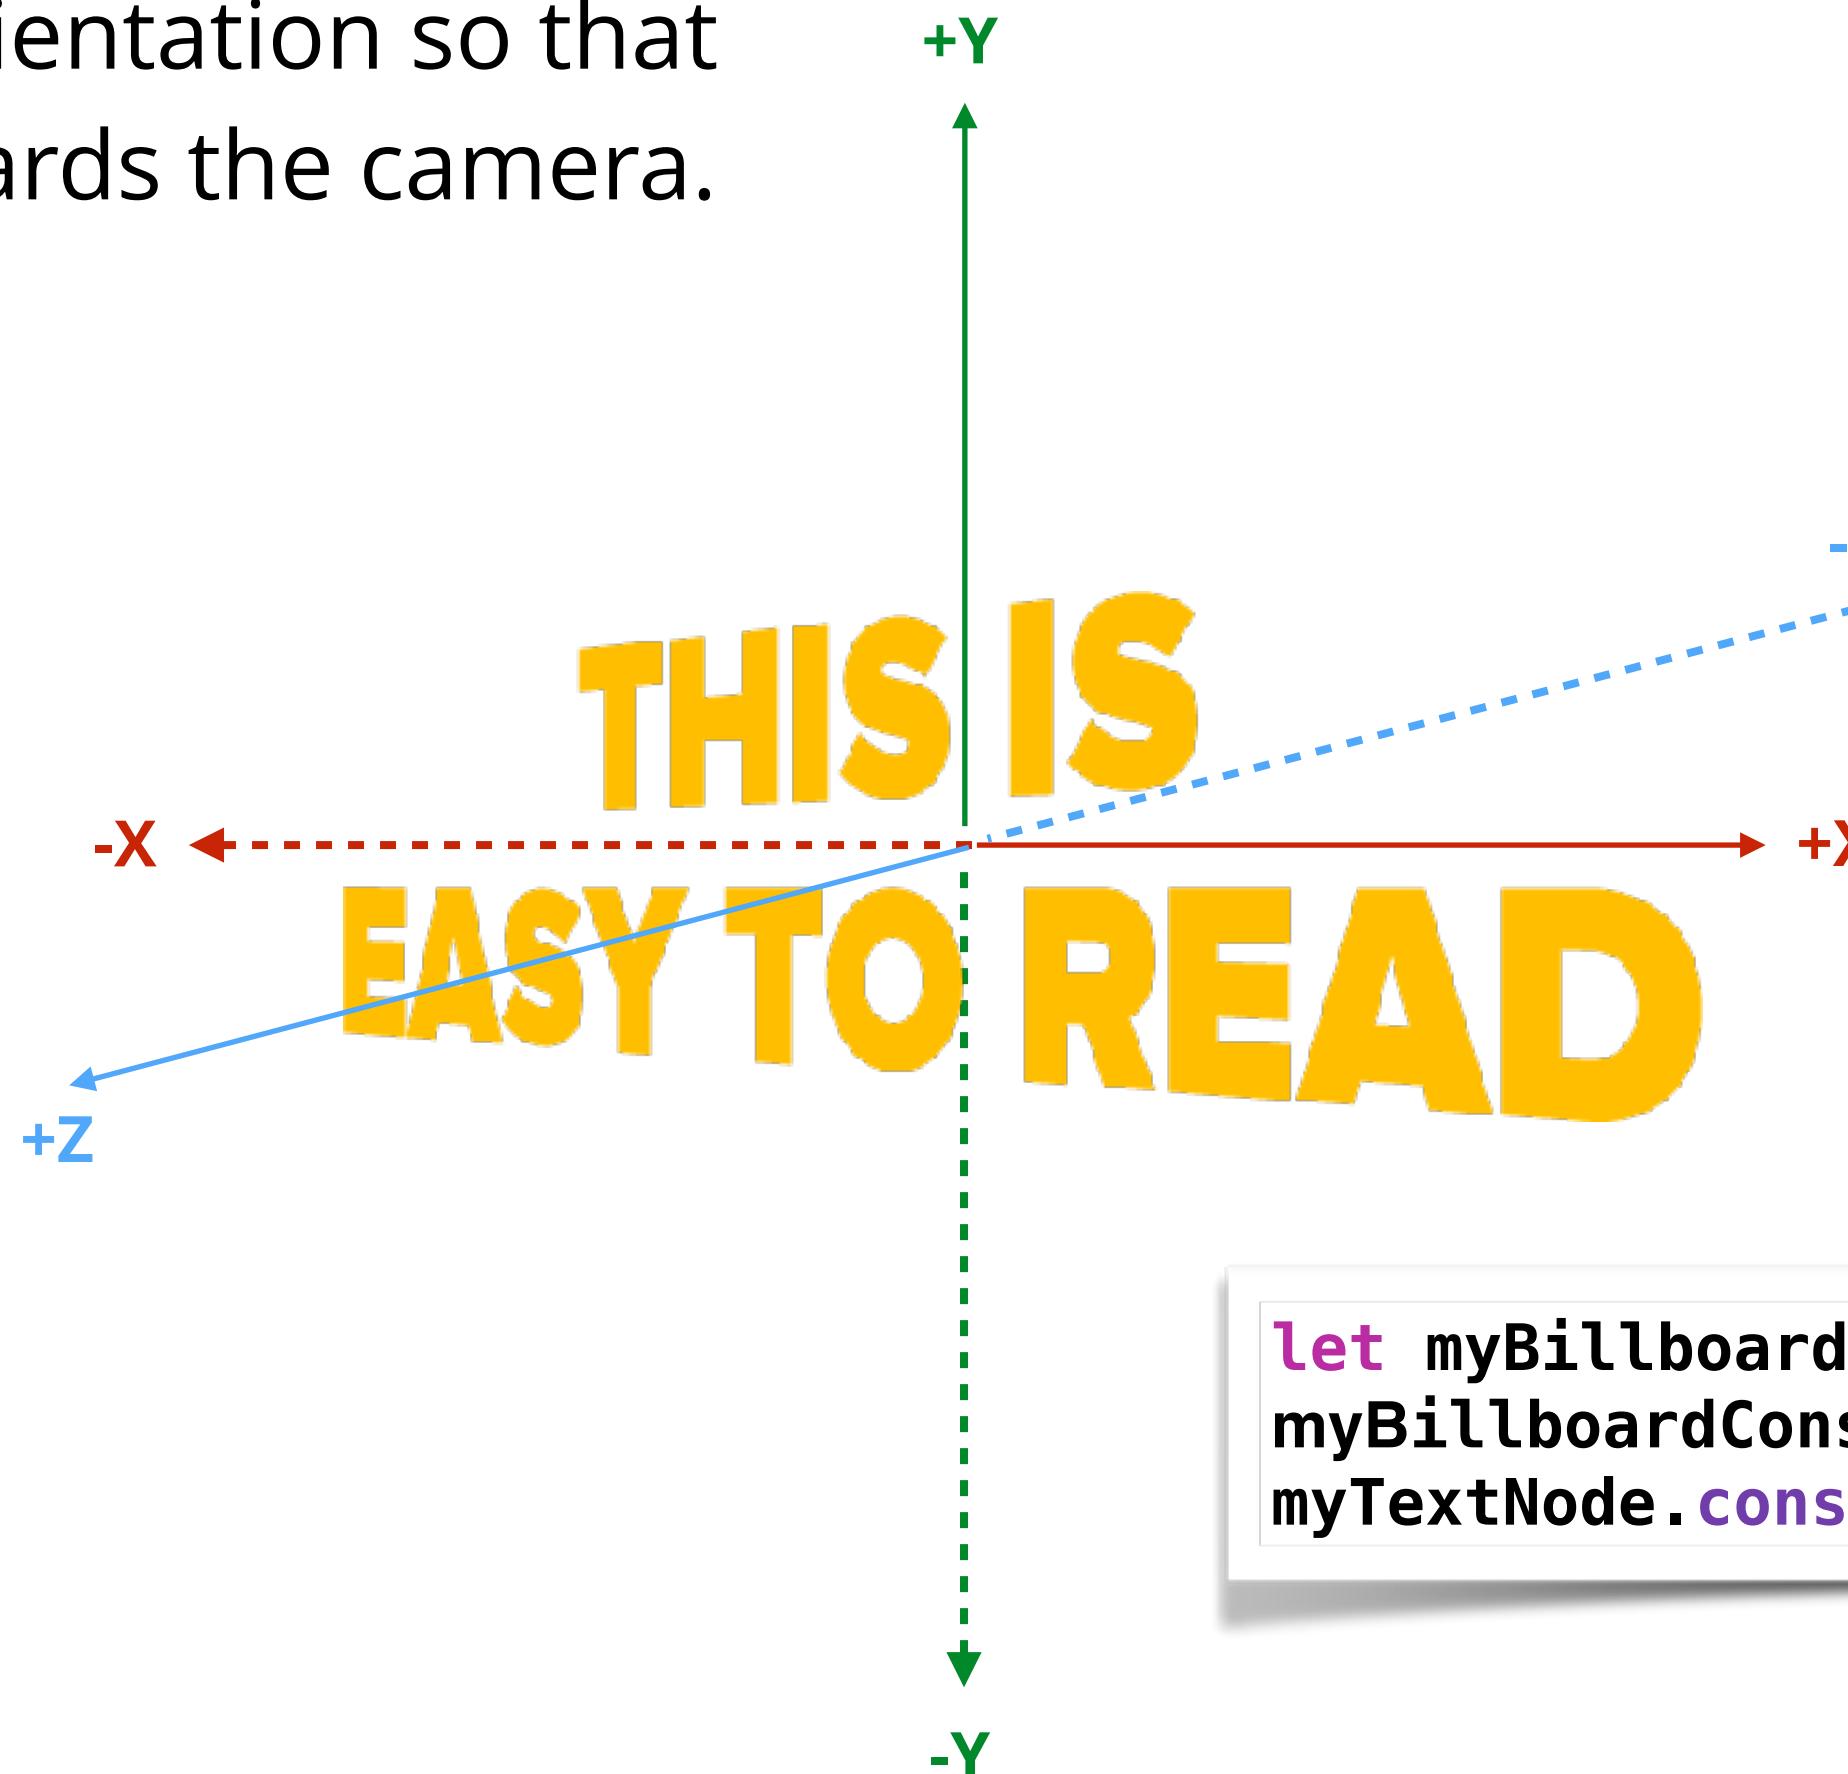
**THIS IS  
EASY TO READ**

We'd like it if 3D text would swing around  
to face the user wherever they went  
so that it's always easy to read.



# BILLBOARD CONSTRAINTS

A **billboard constraint** automatically adjusts a SceneKit node's orientation so that it always points towards the camera.



```
let myBillboardConstraint = SCNBillboardConstraint()  
myBillboardConstraint.freeAxes = SCNBillboardAxis.Y  
myTextNode.constraints = [myBillboardConstraint]
```

# DEMO 3: BAEDEKAR



# How BAEDEKAR WORKS, PART 1

1. Provide the app with reference 2D images.



2. Continuously seek instances of the reference 2D images in the real world.



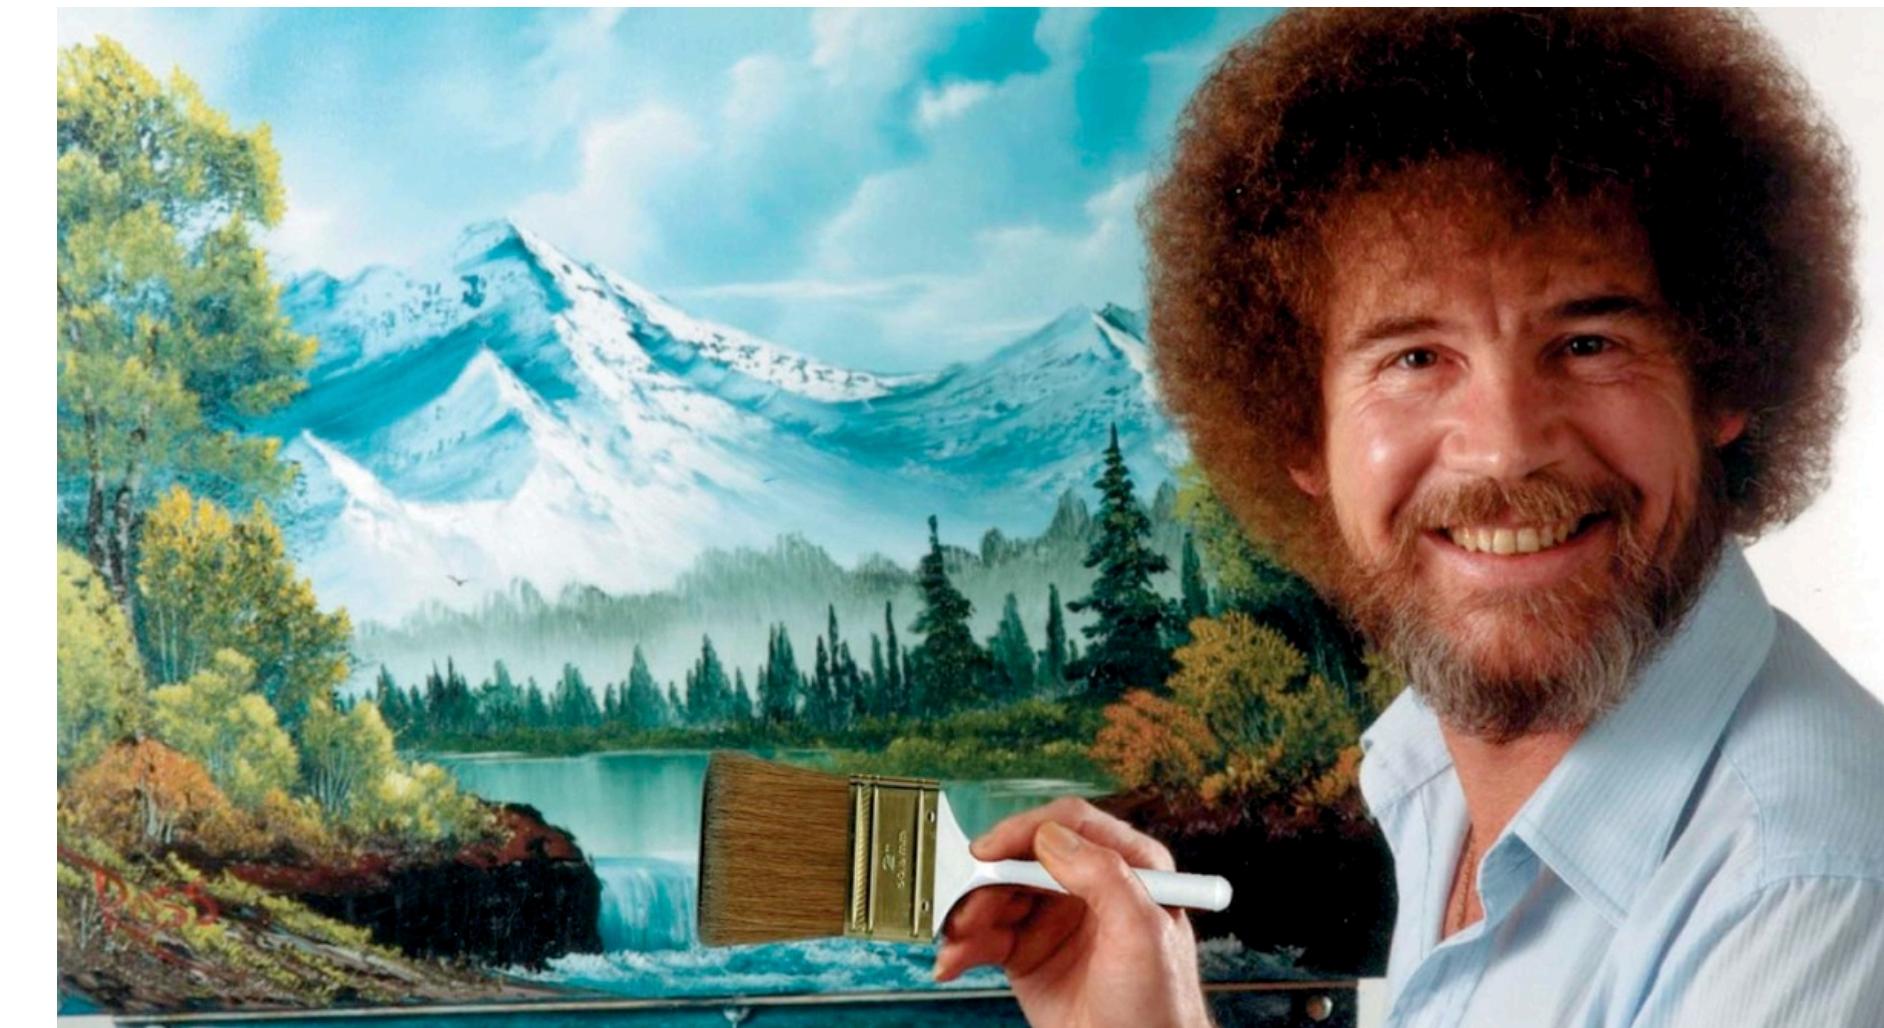
# HOW RAYKEA WORKS, PART 2

3. If the detected 2D image's reference name doesn't begin with "block\_this", overlay the image with a tappable hotspot and AR text displaying its name.

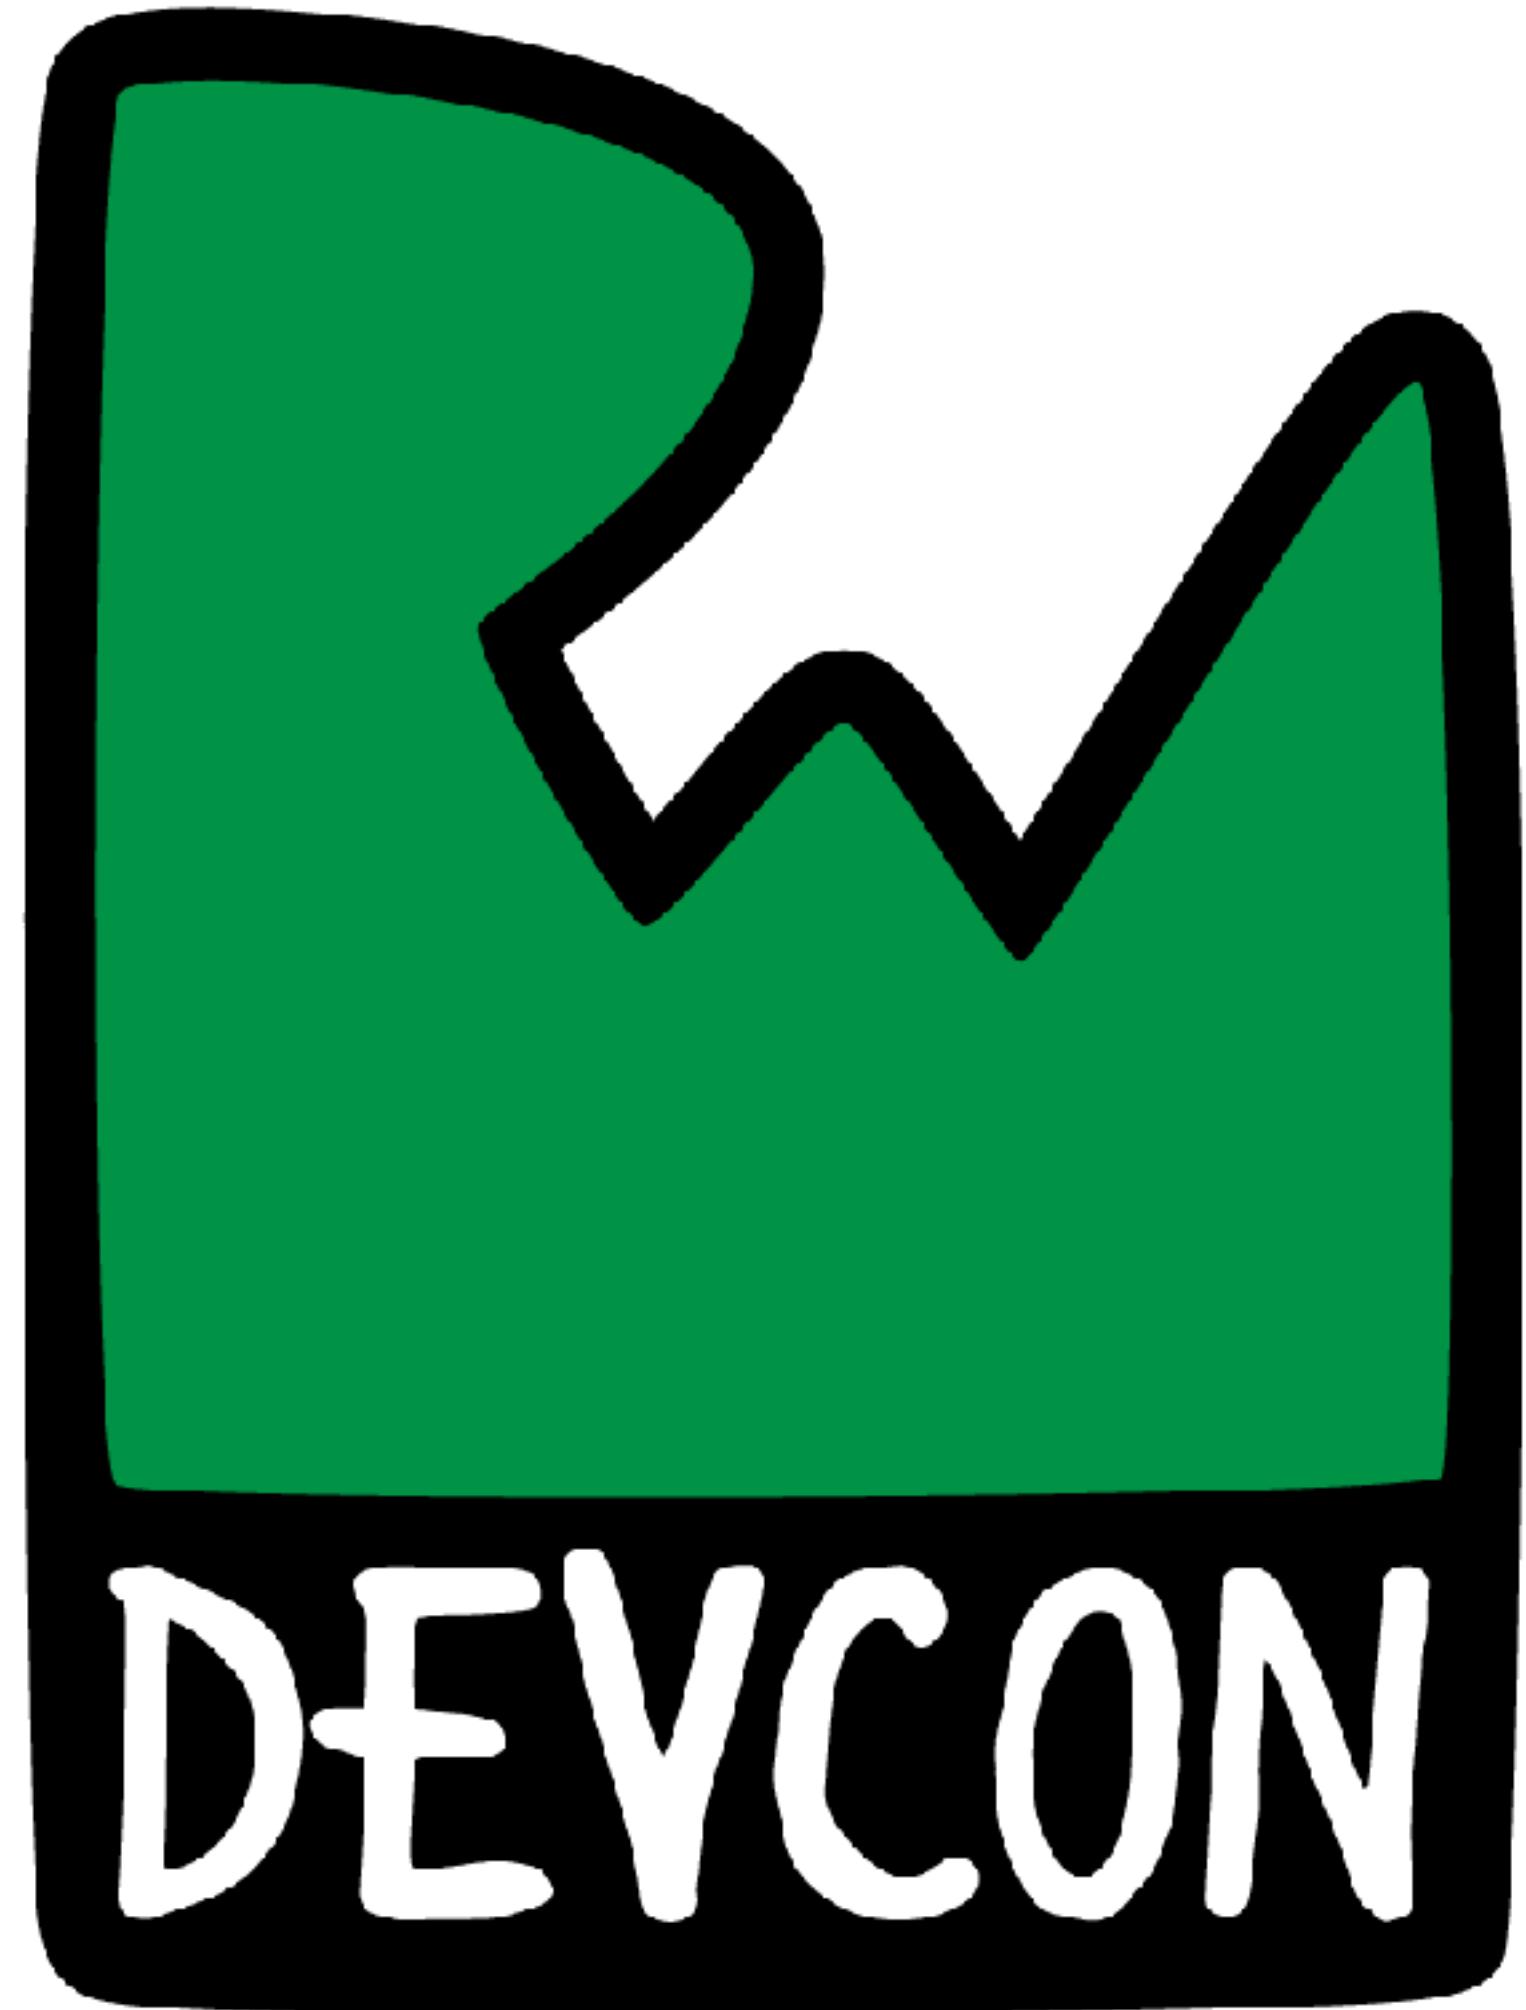


**Starry night**

4. If the detected 2D image's reference name begins with "block\_this", cover up the image with a more comforting one.



# Workshop: ARKit in Depth



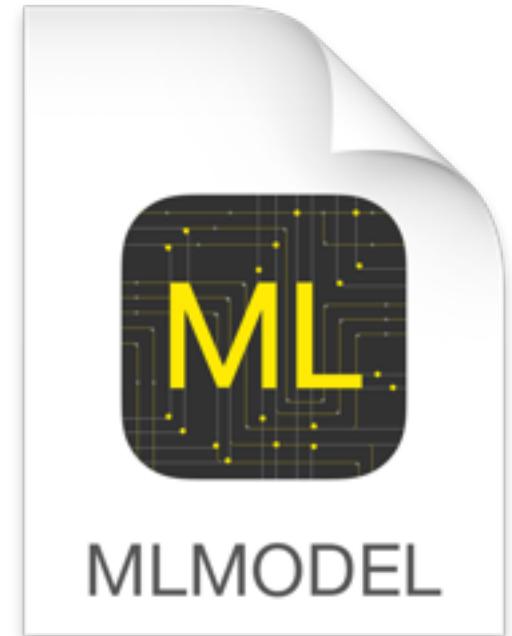
RV  
VISION QUEST

# CATEGORIZATION AND MACHINE LEARNING

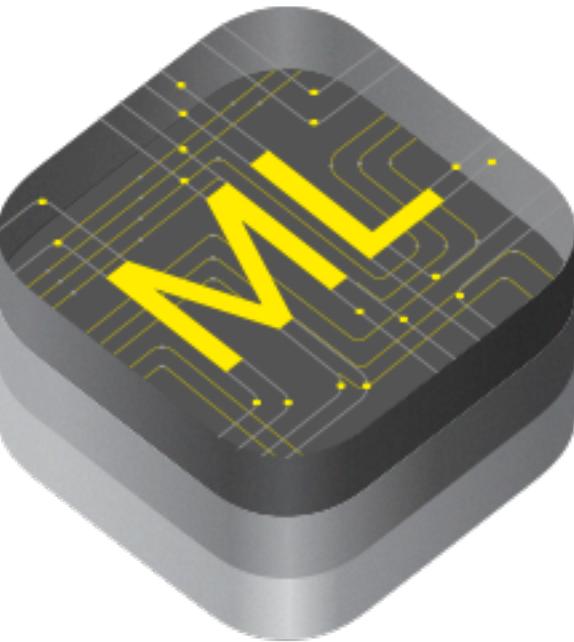


R  
W

# MODELS AND CORE ML



Think of a machine learning model  
as software...



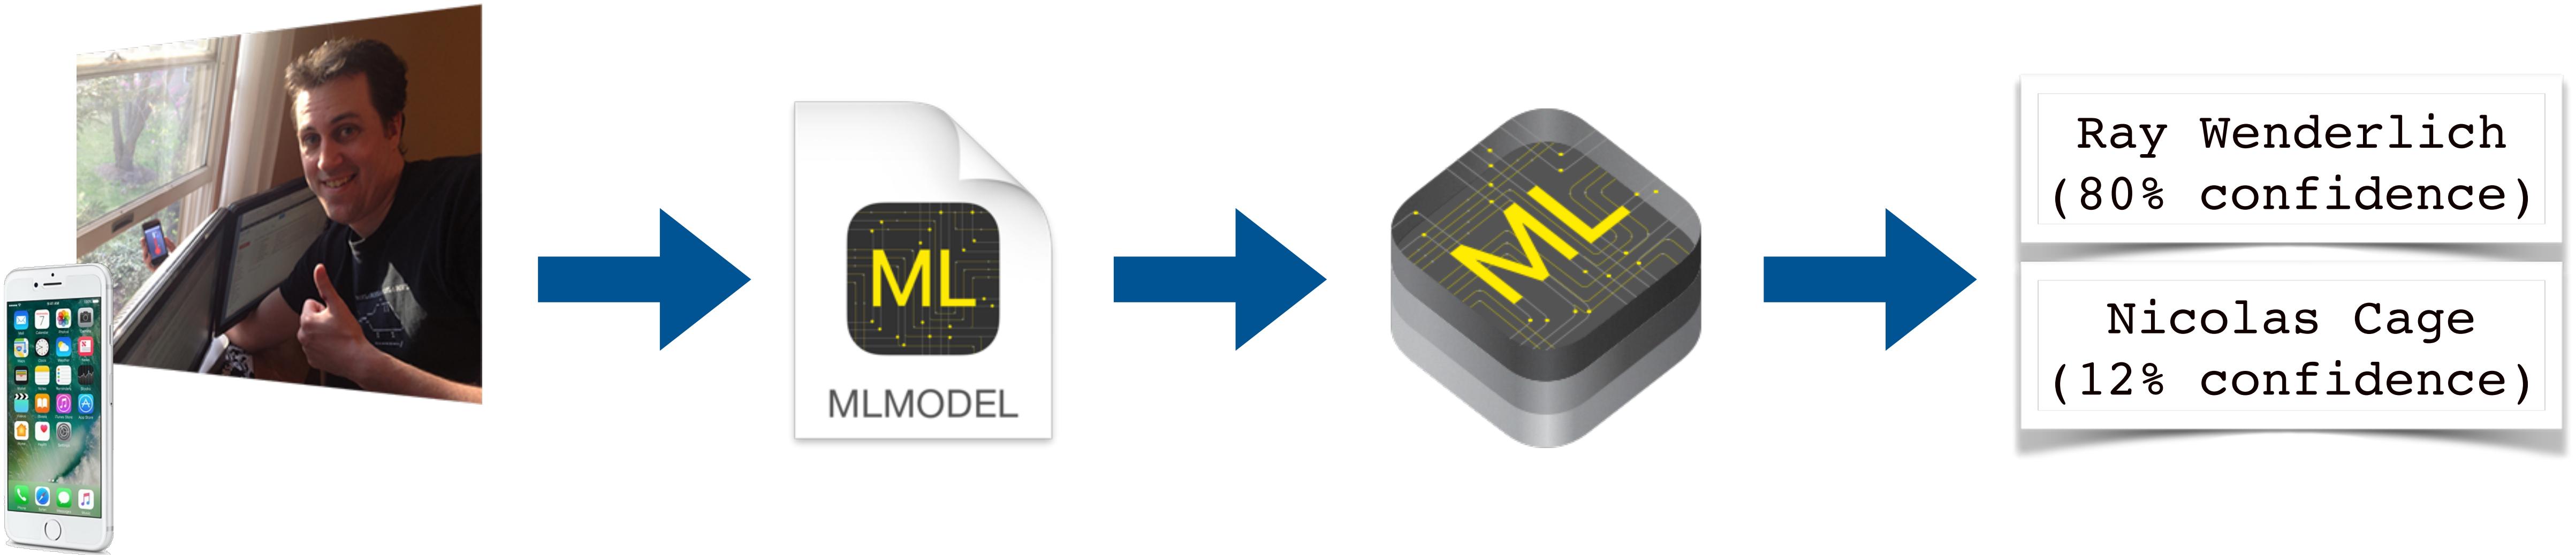
...and think of Core ML as the hardware  
on which that software runs.

When you “run” a model on Core ML,  
you’re harnessing the model’s training  
to take various inputs and **categorize** them.



A collage of images including a black bear playing a guitar, a grumpy-looking cat, and two men in dark sunglasses. <p>Metal</p>	A collage of four young men posing together. <p>Pop</p>	A portrait of a man with blonde hair and a blue suit, with the text "Goin' back to Miami" curved above him. <p>???</p>
---	---	--

# IMAGE RECOGNITION WITH CORE ML AND VISION



Take some **image data...**

...and a  
**machine learning model**  
for **categorizing images...**

...then use **Core ML**  
to categorize the image  
using the model...

...and get the results,  
along with Core ML's  
confidence in them.



Sometimes the results  
are **hilariously wrong**.



Actual sea urchin



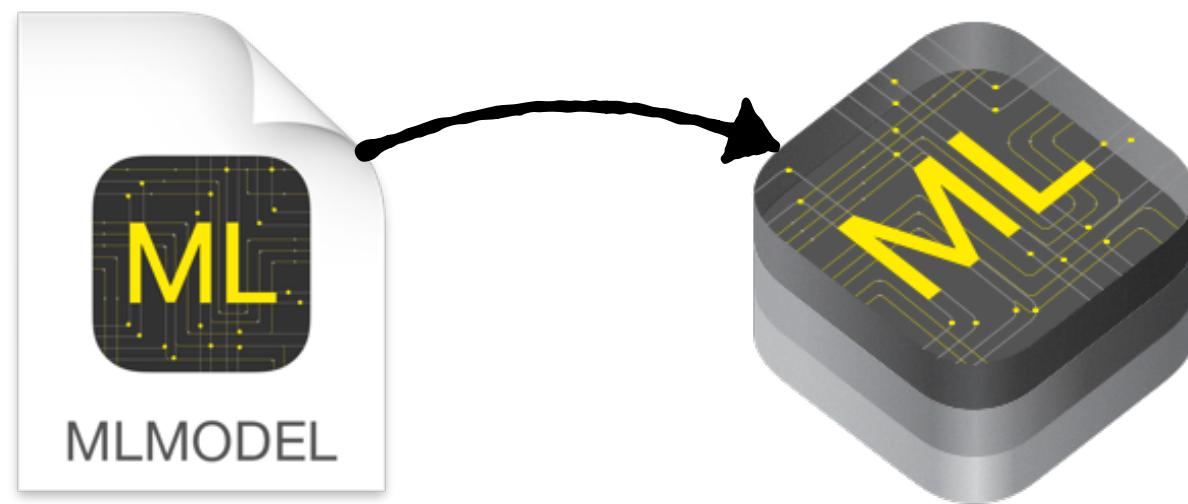
(Hey, it's pretty close...)

# DEMO 4: VISION QUEST



# How VISION QUEST WORKS, PART 1

1. Load the appropriate model into Core ML and request that Core ML use the model to classify images.



2. Continuously capture the camera image and feed it to Core ML.



# How VISION QUEST WORKS, PART 2

3. Continuously display Core ML's top 3 guesses of what the dominant item in the image is at the top of the screen.

**Beer mug - 33%**

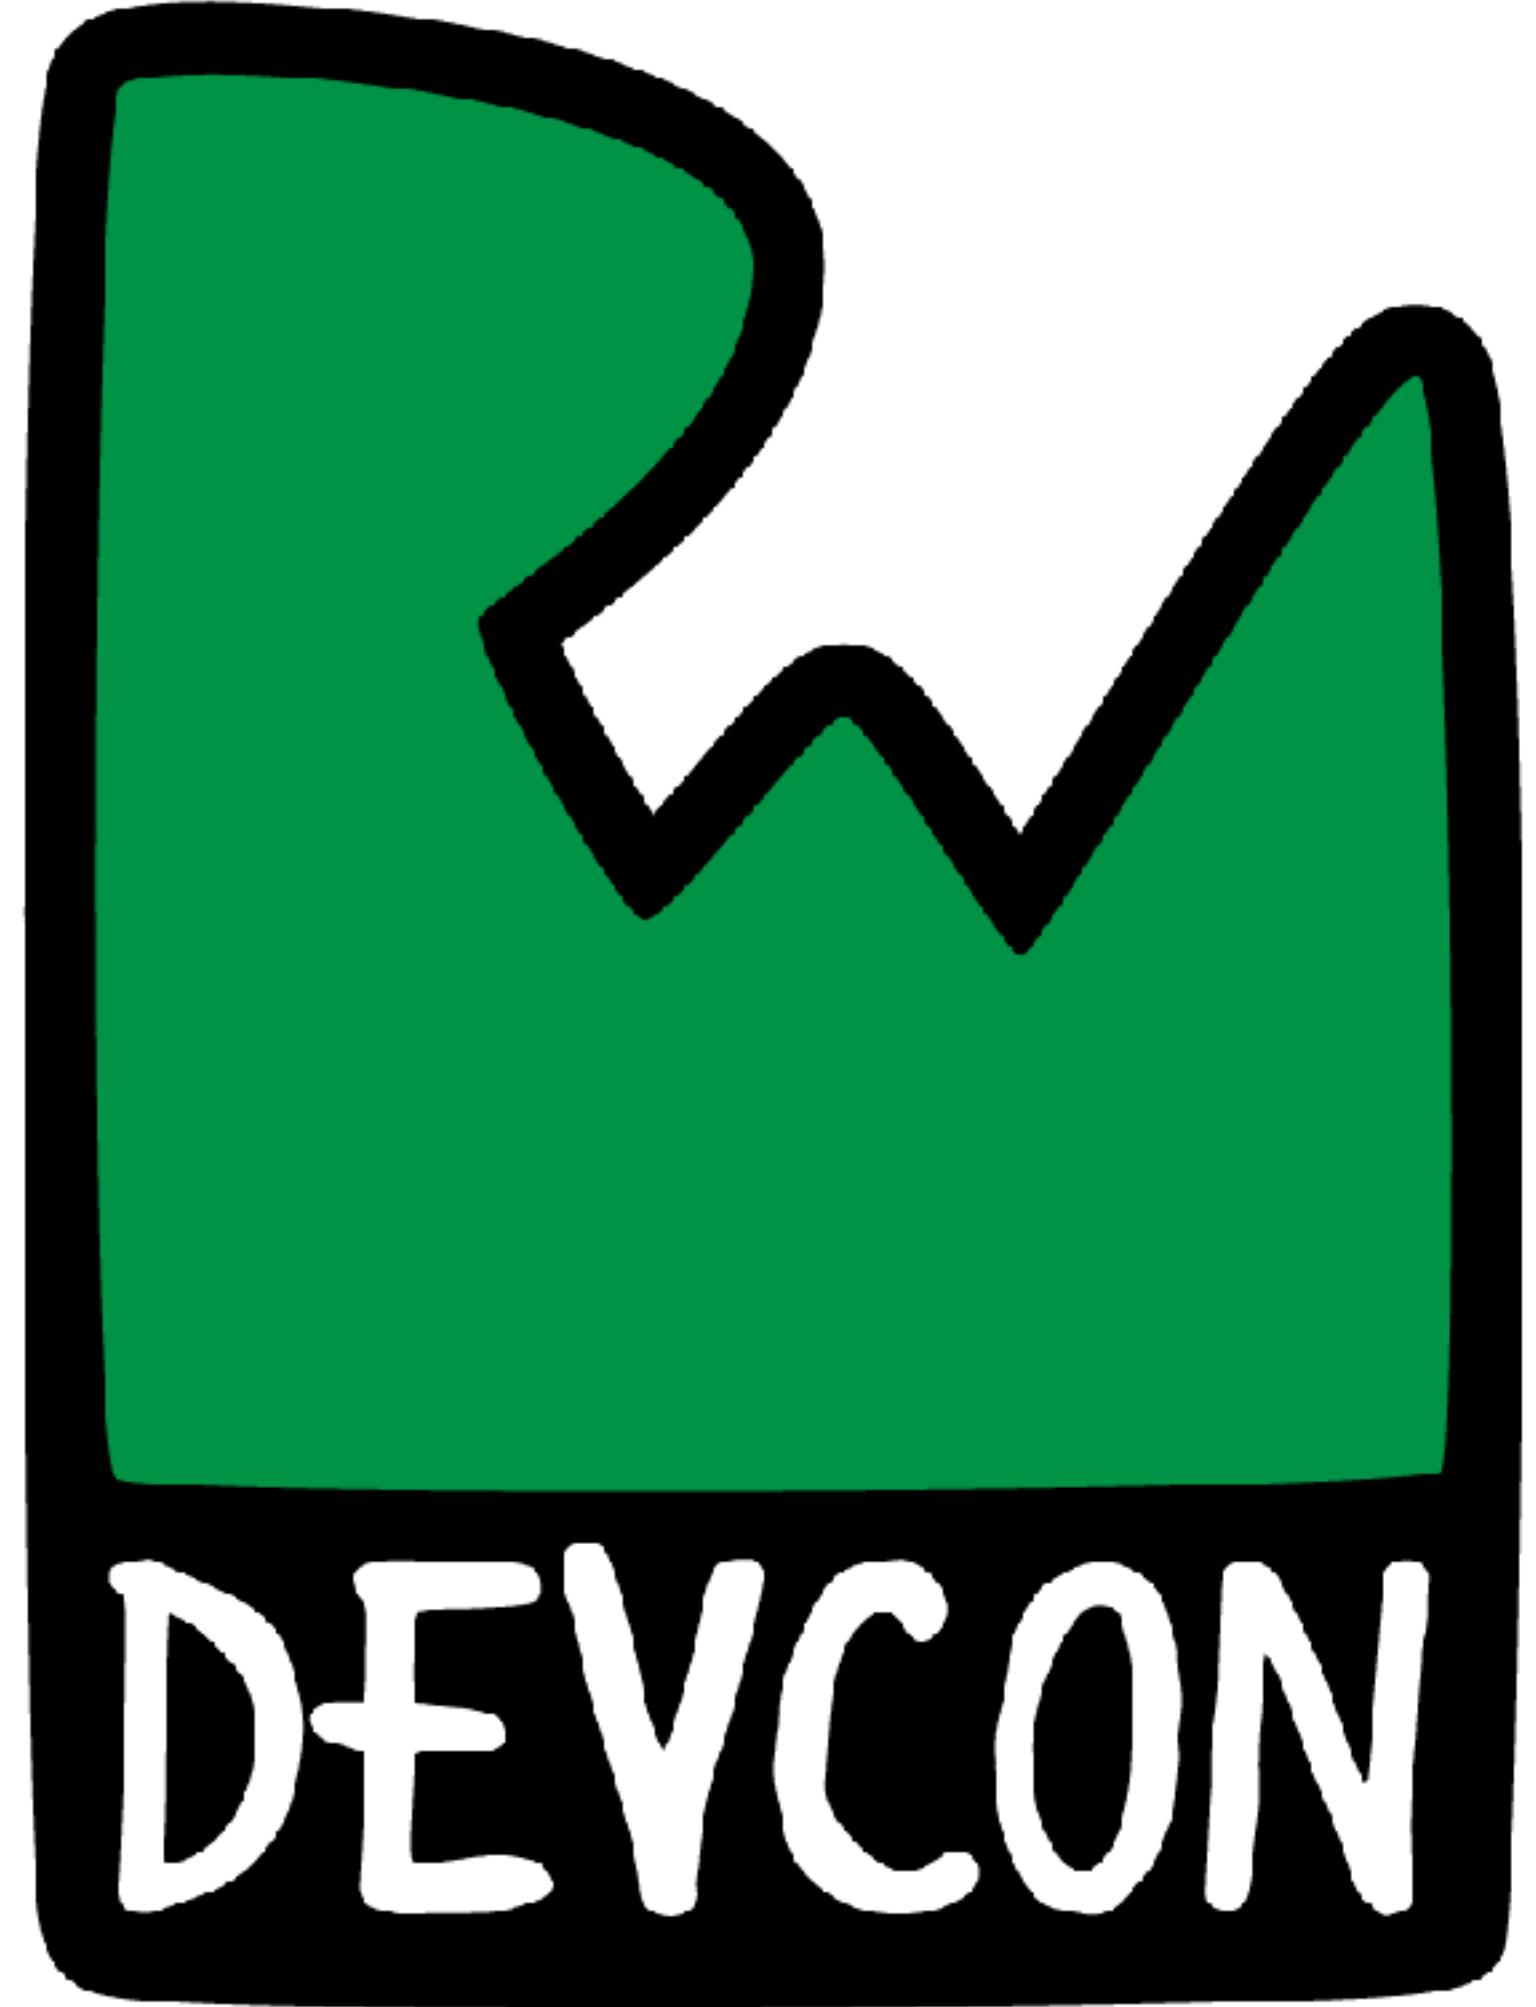
**Beaker - 10%**

**Glass jar - 5%**

4. When the user taps on a feature point, draw the top guess and its confidence as AR text.



Session #:  
Getting started  
with ARKit



CONCLUSION

# WHAT YOU LEARNED



## Demo 1: *Happy AR Painter*

Your first AR scene, where you added simple SceneKit geometric shapes to the scene and used the device's position and orientation.



## Demo 2: *Raykea*

You took Demo 1's lessons, added 3D models, hit tests, and plane detection, and started interacting with the real world.



## Demo 3: *BaedekAR*

You used a new feature in ARKit 1.5 to detect, add hotspots to, and block known 2D images, and tag them with easily-read AR text.



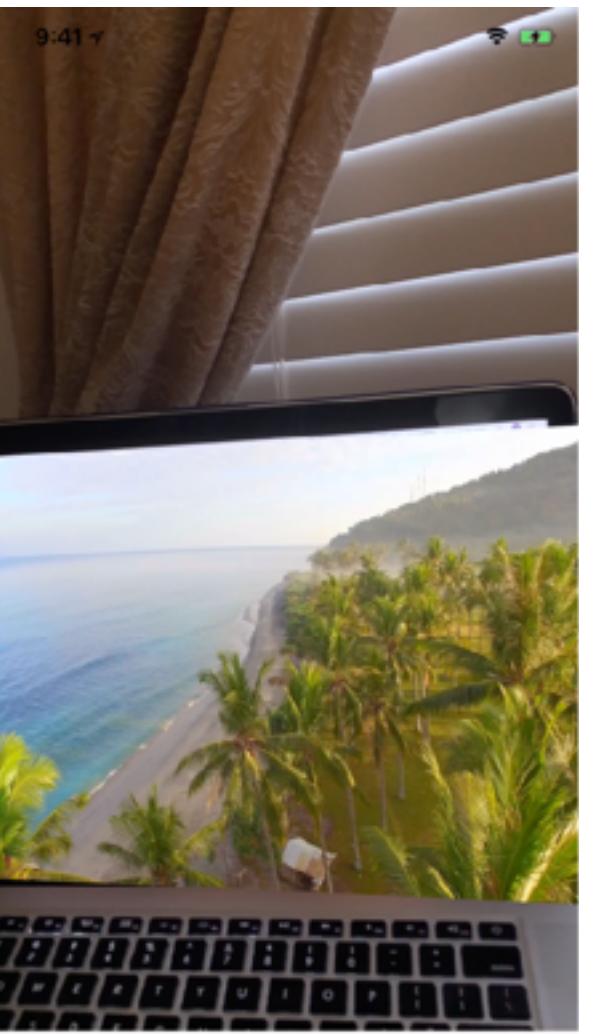
## Demo 4: *Vision Quest*

Harnessing the true potential of AR by taking what you learned about ARKit and combining it with machine learning.



# WHAT You DIDN'T LEARN

Even at less than a year old, ARKit already offers so much ground to cover that a single workshop can't cover it all...



Embedding video in AR scenes



SceneKit physics



SpriteKit and  
ARKit



ARKit and Core Location

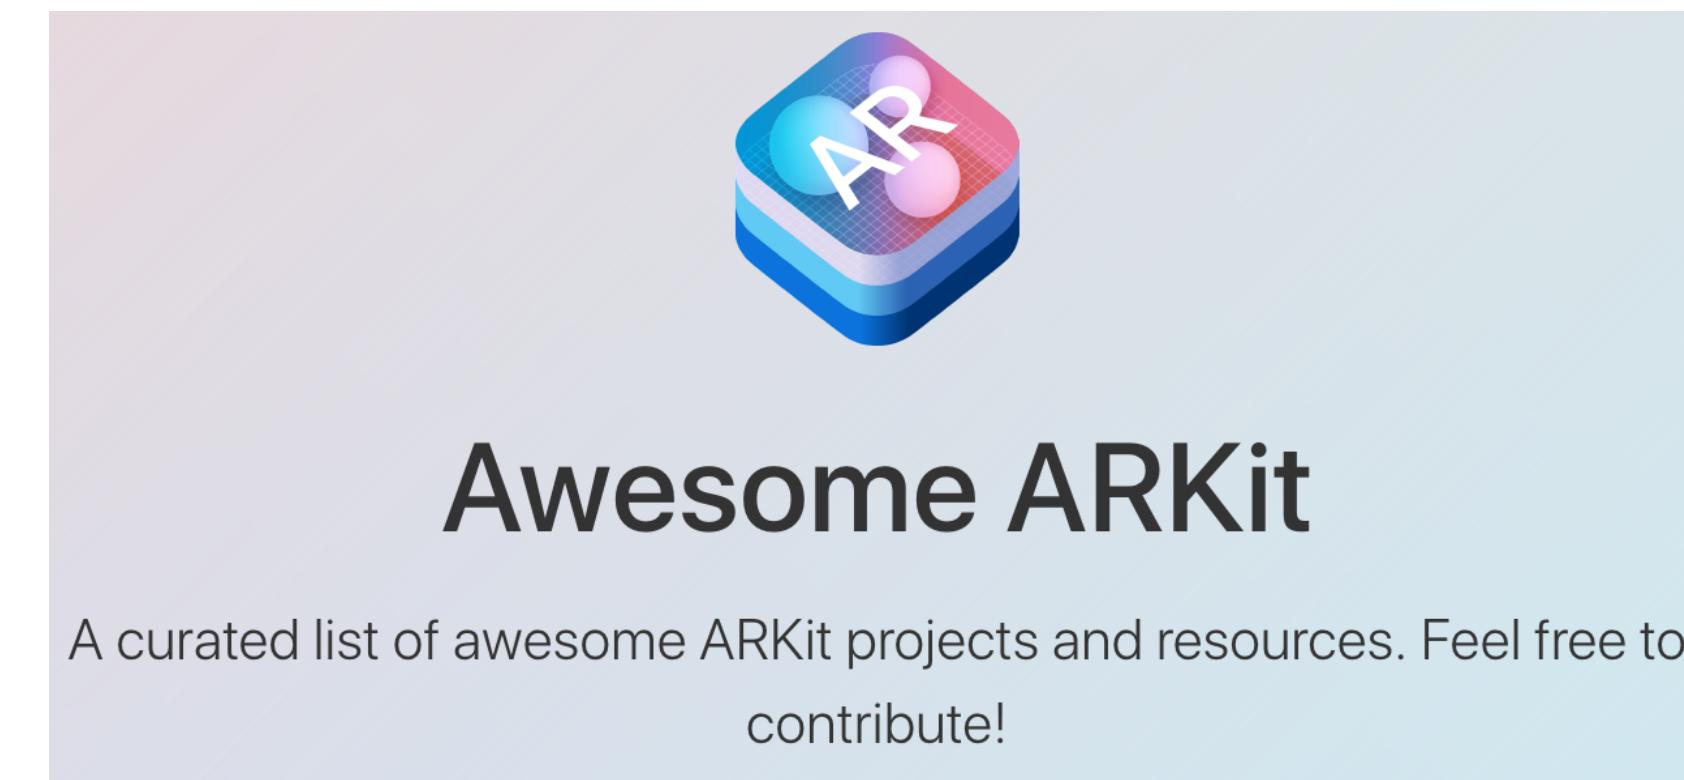


# WHERE To Go FROM HERE?



## Apple's ARKit documentation

- ⚙️ Apple developer ([developer.apple.com/arkit](https://developer.apple.com/arkit))
- ⚙️ Apple's human interface guidelines  
([developer.apple.com/ios/human-interface-guidelines](https://developer.apple.com/ios/human-interface-guidelines))



**“Awesome ARKit” repo on GitHub**  
[github.com/olucurious/Awesome-ARKit](https://github.com/olucurious/Awesome-ARKit)



# RAYWENDERLICH.COM AND ARKIT



## *ARKit by Tutorials*

Coming soon!

**Augmented Reality in Android with Google's Face API**

 [joey devilla](#) on July 12, 2017

If you've ever used [Snapchat's "Lenses" feature](#), you've used a combination of **augmented reality** and **face detection**.

Augmented reality — *AR* for short — is technical and an impressive-sounding term that simply describes real-world images overlaid with computer-generated ones. As for face detection, it's nothing new for humans, but finding faces in images is still a new trick for computers, especially *handheld* ones.



## Articles

Coming soon!

**raywenderlich.com Video Tutorials**

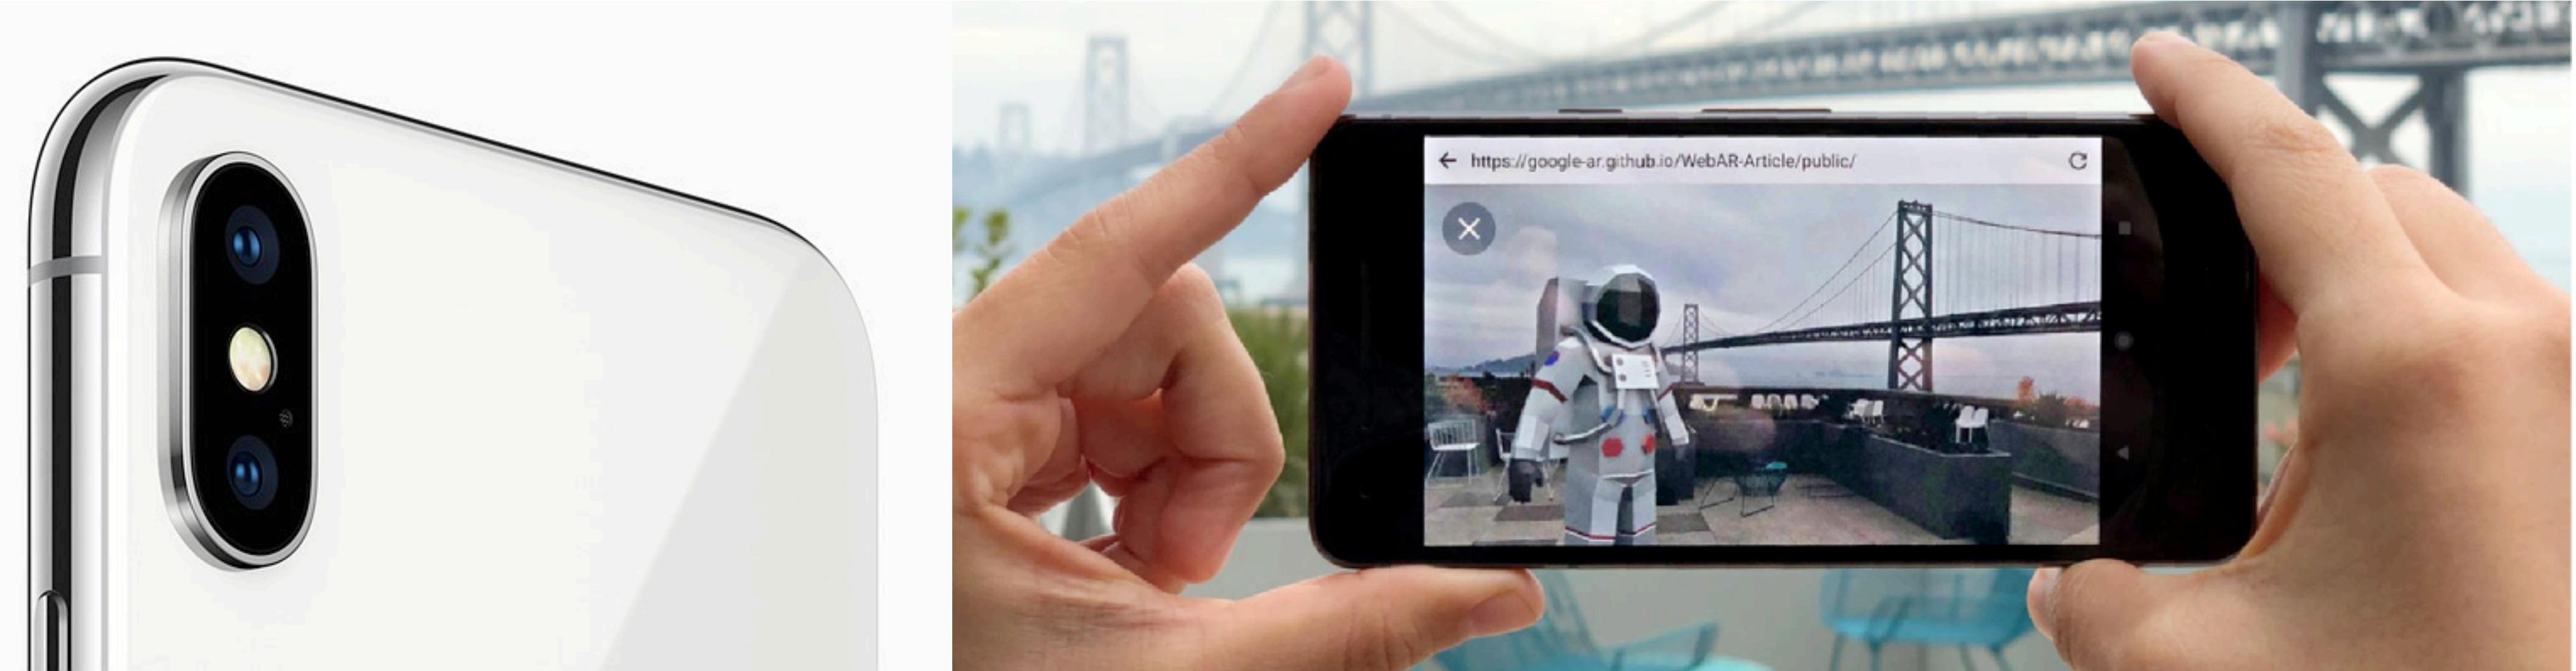
  


## Videos

Coming soon!



# THESE ARE THE EARLY DAYS OF MOBILE AR



R  
W

# YOU HAVE A RESPONSIBILITY



# HOW TO FIND ME ONLINE



***Global Nerdy, my tech blog***

[globalnerdy.com](http://globalnerdy.com)



**Twitter**

@AccordionGuy



**LinkedIn**

[linkedin.com/in/joeydevilla](https://linkedin.com/in/joeydevilla)